

Programmazione concorrente

Laurea Magistrale in Ingegneria Informatica

Università Tor Vergata

Docente: Romolo Marotta

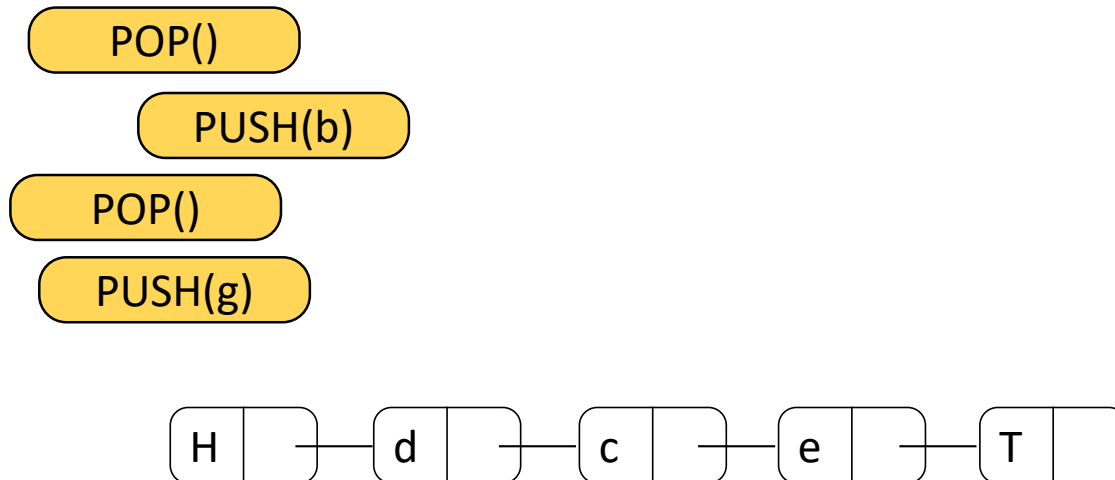
Concurrent data structures

1. Stack
2. Set

Concurrent Data Structures: **Stacks**

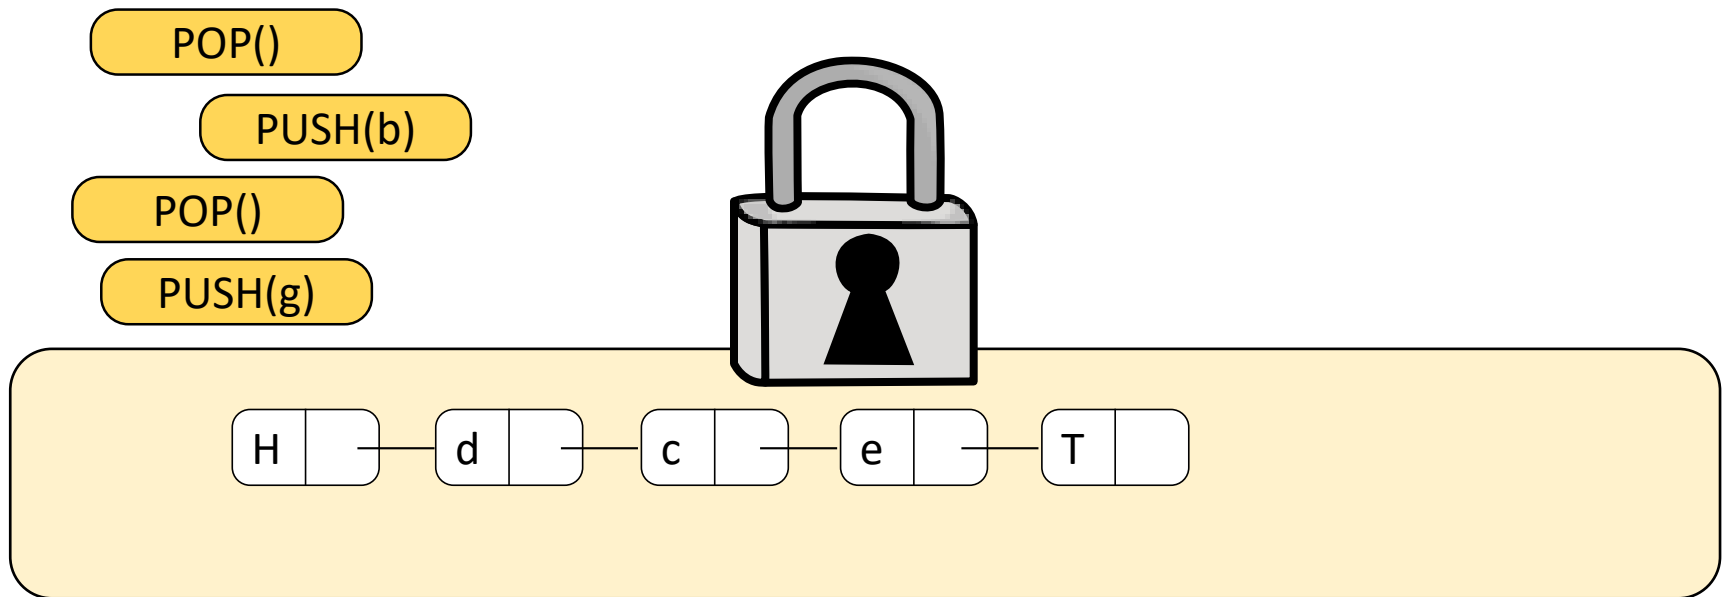
Stack implementation

- Stack methods:
 - `push(v)`
 - `pop()`
- Implemented as a linked list



Concurrent stack implementations

- Resort to a global lock



Read-Modify-Write

- RMW instructions allow to read memory and modify its content in an apparently instantaneous fashion.

```
1. RMW(MRegister *r, Function f){  
2.   atomic{  
3.     old = r;  
4.     *r = f(r);  
5.     return old;  
6.   }  
7. }
```

- Even conventional atomic Load and Store can be seen as RMW operations

Compare-And-Swap

- Compare-and-Swap (CAS) is an atomic instruction used in multithreading to achieve synchronization
 - It compares the contents of a memory area with a supplied value
 - If and only if they are the same
 - The contents of the memory area are updated with the new provided value
- Atomicity guarantees that the new value is computed based on up-to-date information
- If, in the meanwhile, the value has been updated by another thread, the update fails
- This instruction has been introduced in 1970 in the IBM 370 trying to limit as much as possible the use of spinlocks

Compare-And-Swap

- RMW instructions allow to read memory and modify its content in an apparently instantaneous fashion.

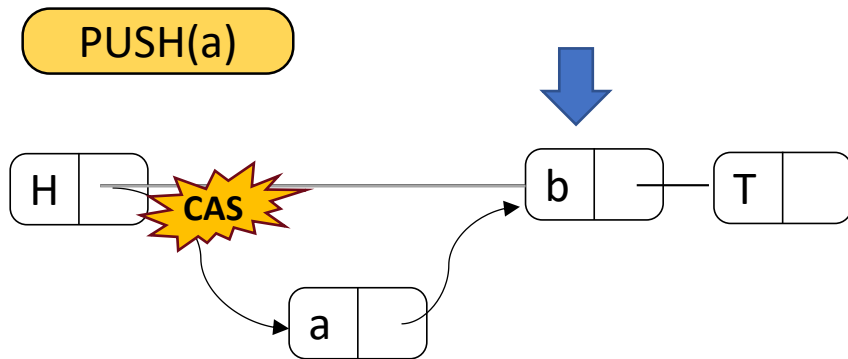
```
1. CAS(Mregister *r, Value old_value, Value new_value f){  
2.   atomic{  
3.     Value res = *r;  
4.     if(*r == old_value) *r = new_value;  
5.     return res;  
6.   }  
7. }
```

- CAS is implemented by x86 architectures (see CMPXCHG)
- gcc offers the `__sync_val_compare_and_swap` builtin

Attempt 1

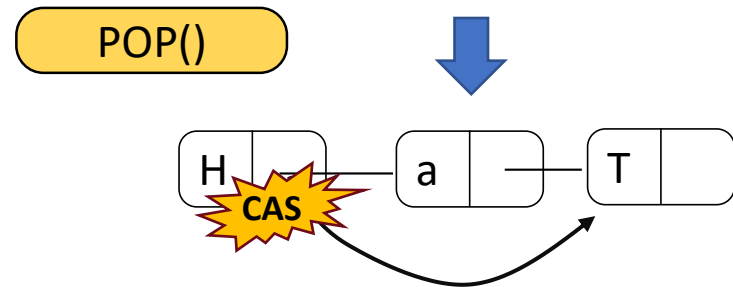
Push:

1. Get head next
2. Insert the new item with a CAS
3. If CAS fails, restart



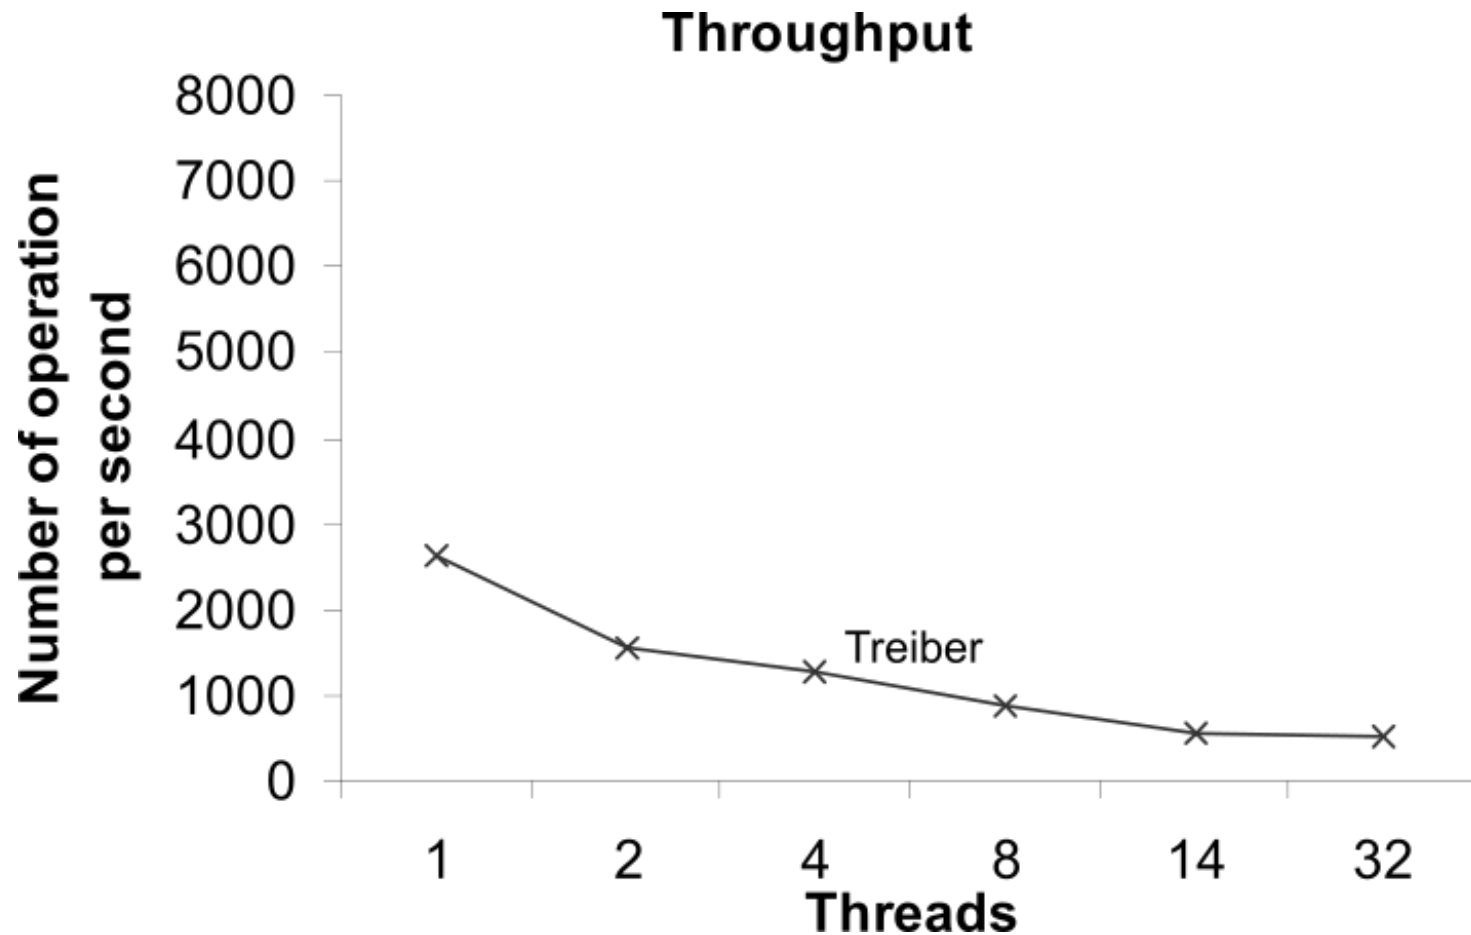
Delete:

1. Get head next
2. Disconnect the item with a CAS
3. If CAS fails, restart



- Is it scalable?

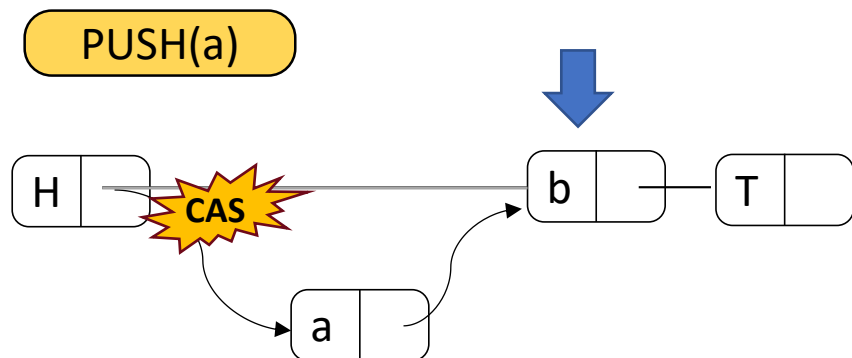
Non-blocking stack – Attempt 2 [Treiber+BO]



Non-blocking stack – Attempt 2 [Treiber+BO]

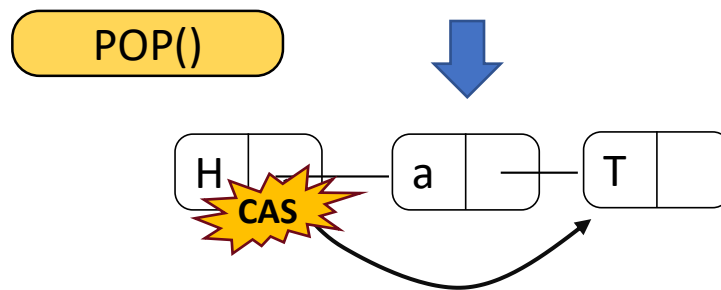
Push:

1. Get head next
2. Insert the new item with a CAS
3. If CAS fails, ~~restart~~ backoff and restart



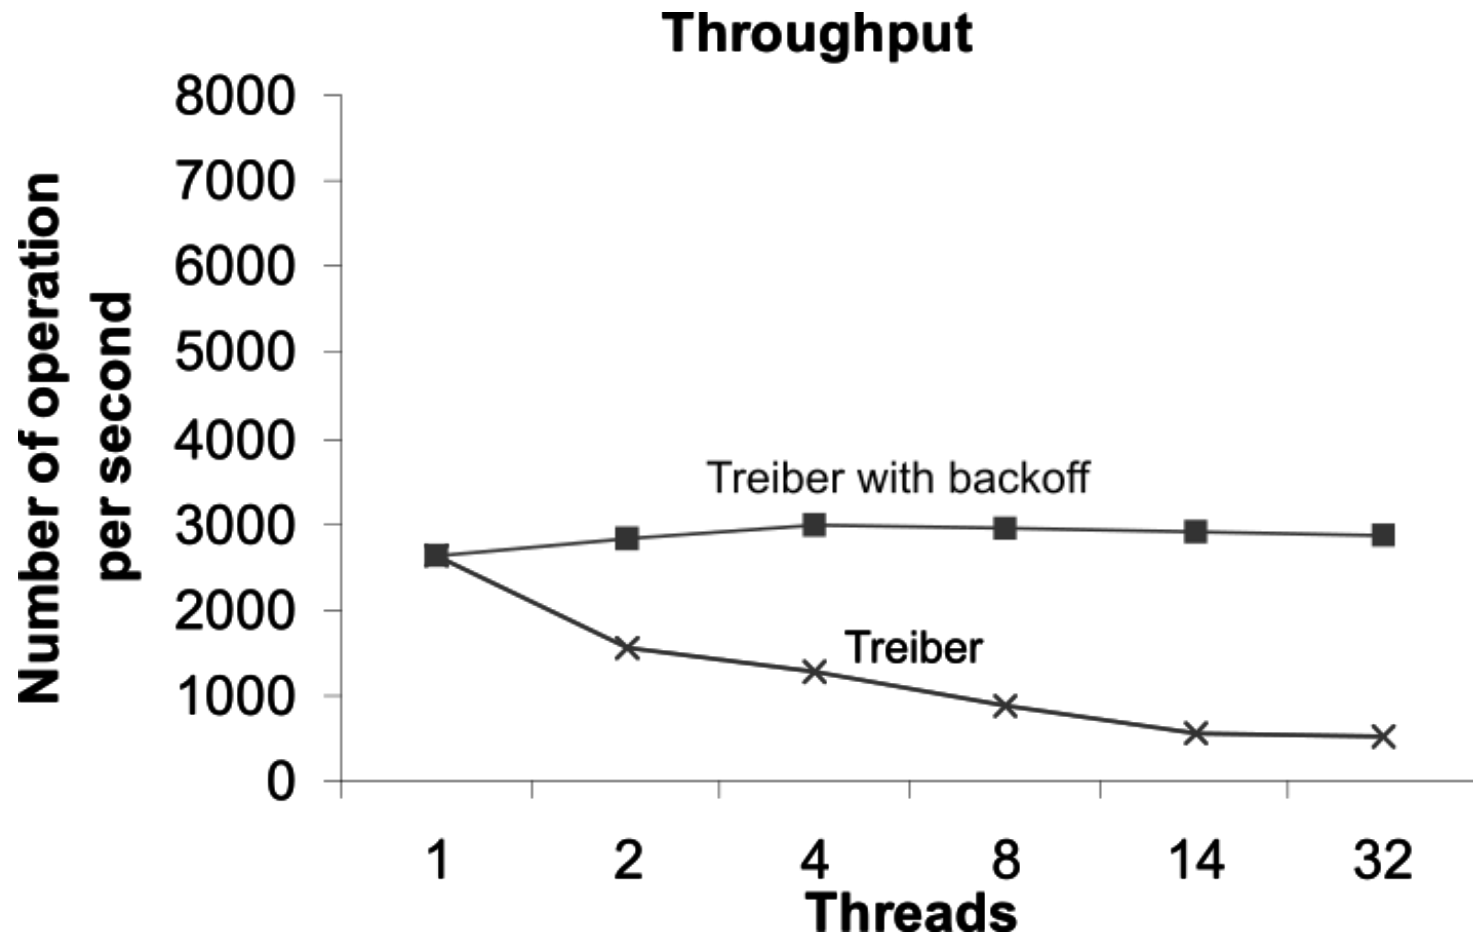
Delete:

1. Get head next
2. Disconnect the item with a CAS
3. If CAS fails, ~~restart~~ backoff and restart



- Is it scalable?

Non-blocking stack – Attempt 2 [Treiber+BO]

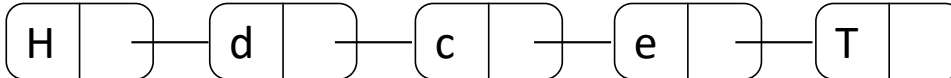


Concurrent stack implementations

- Resort to a global lock
 - Do not scale
- Resort to a naïve non-blocking approach
 - Do not scale
- Resort to a naïve non-blocking approach + Back off
 - Do not scale, but conflict resilient
- How achieve scalability? Make back-off times useful

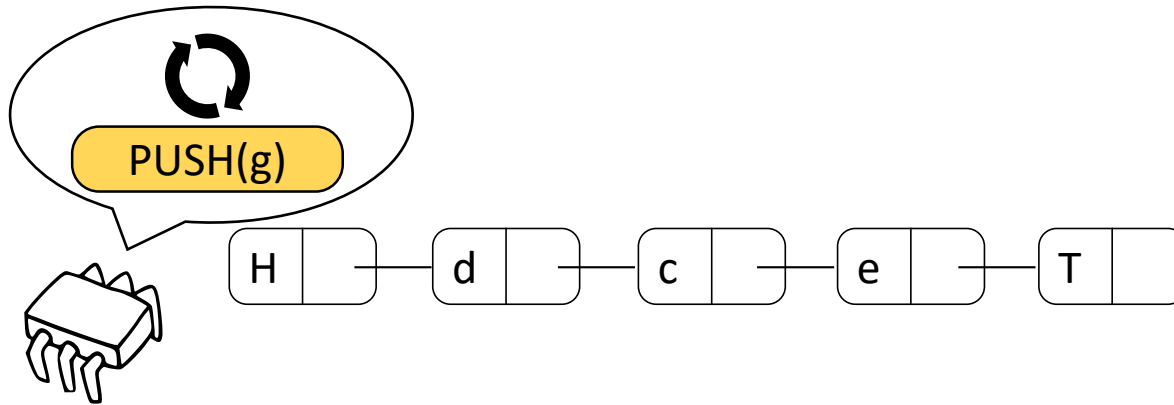
POP()

PUSH(g)



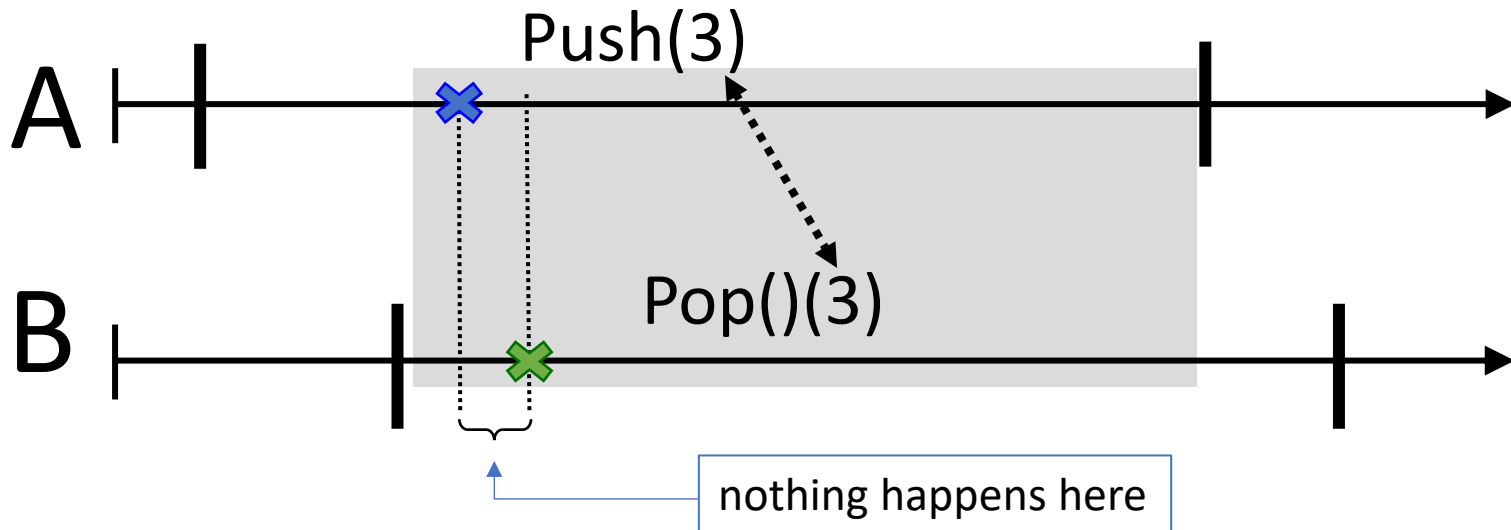
Non-blocking stack – Attempt 3

- How to take advantage of back-off times?



Observation

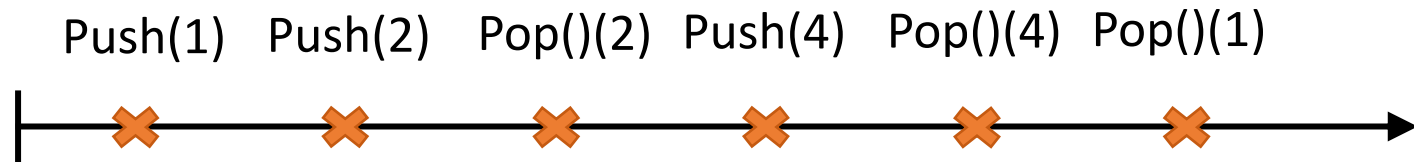
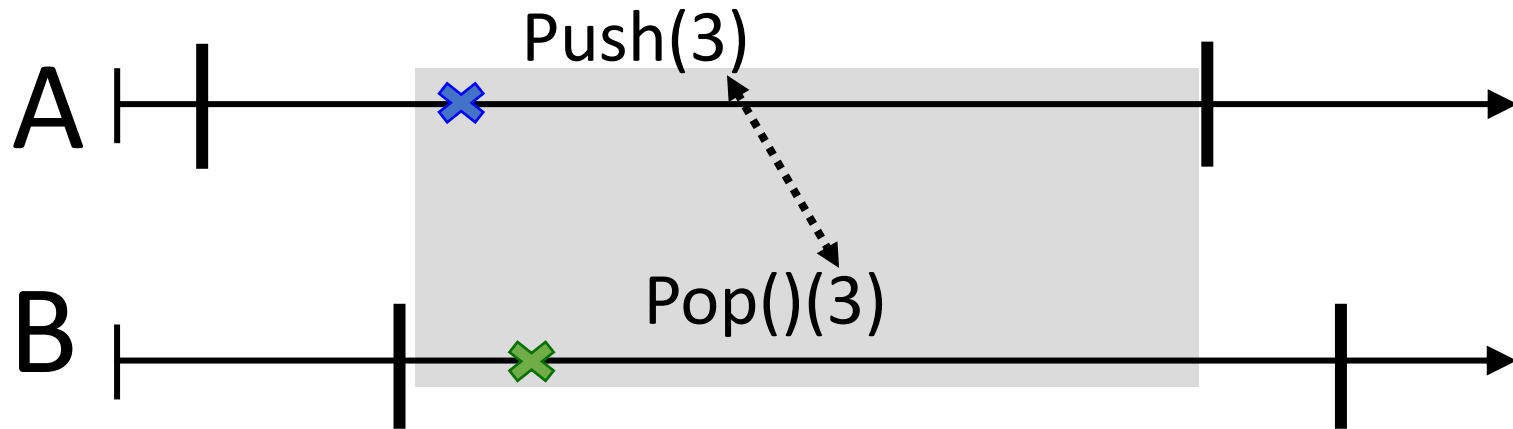
- Concurrent matching push/pop pairs are always linearizable



- A push A and a pop B are:
 - concurrent to each other
 - B returns the item inserted by A
- ⇒ we can always take two points such that:
- A is the last one to insert an item before A linearizes
 - B appears to extract the last item inserted (by A)

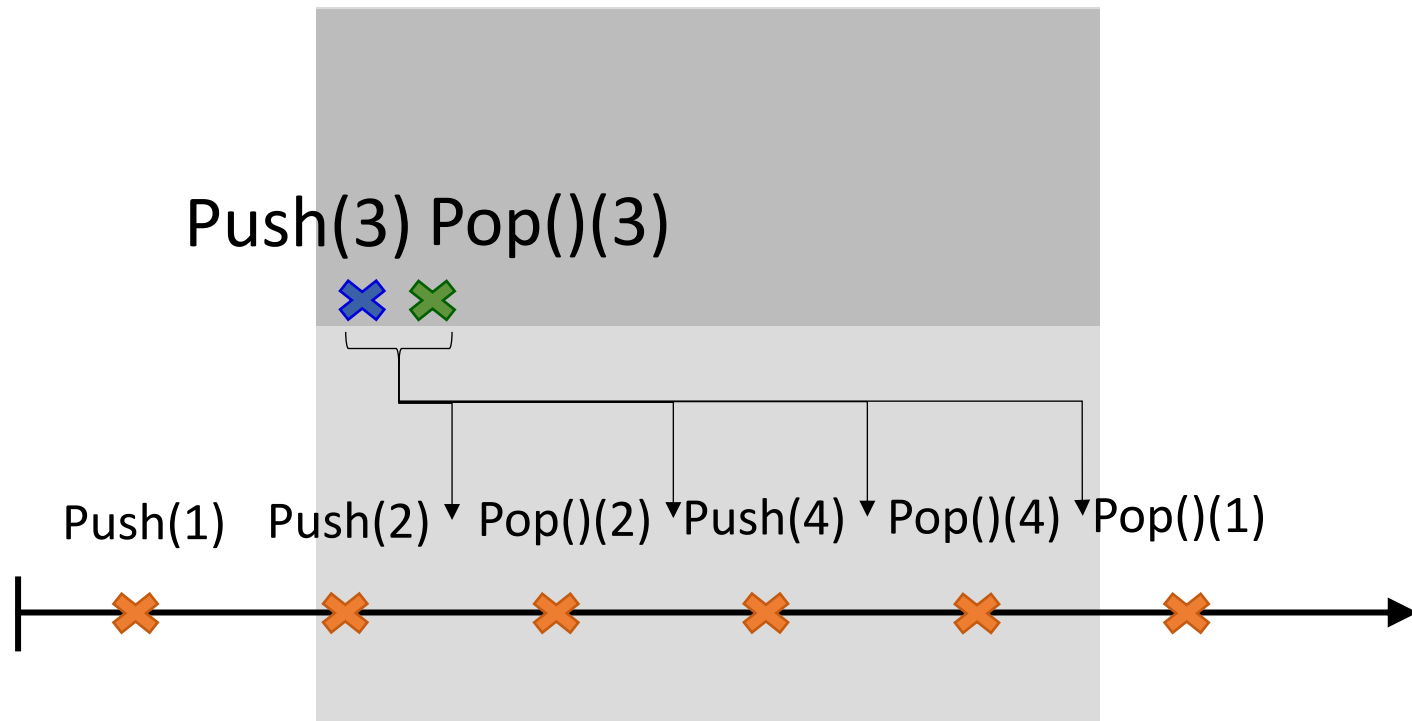
Observation

- Concurrent matching push/pop pairs are always linearizable



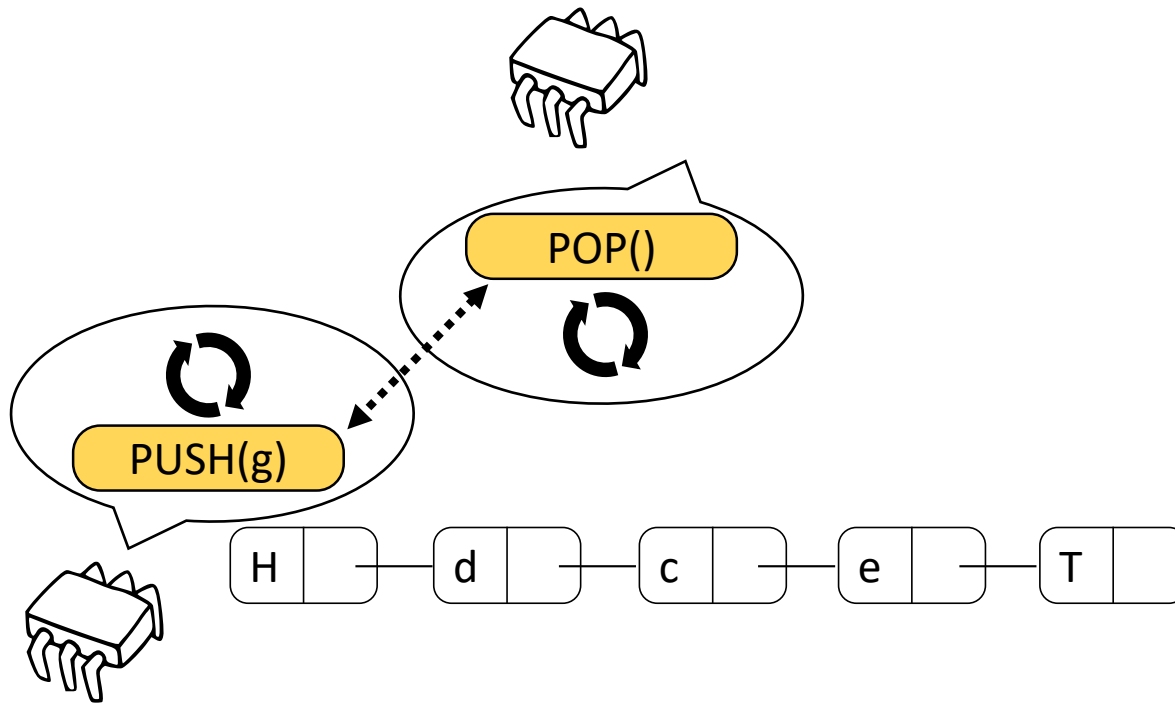
Observation

- Concurrent matching push/pop pairs are always linearizable



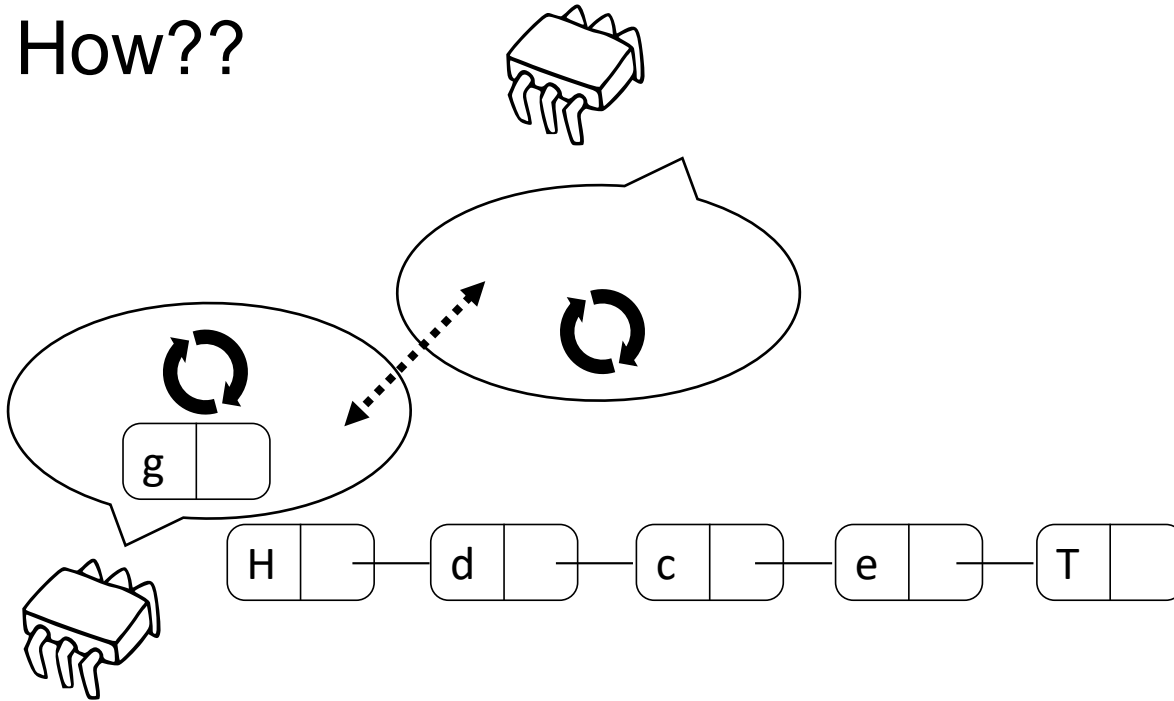
Non-blocking stack – Attempt 3

- How to take advantage of back-off times?
- Hope that an opposite operation arrives while waiting
- Match the two without interacting with the stack



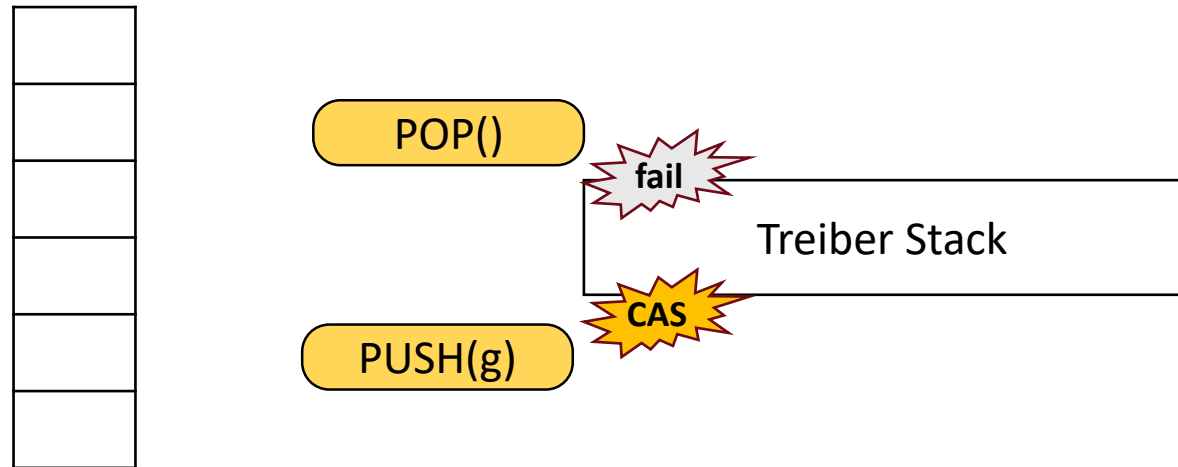
Non-blocking stack – Attempt 3

- How to take advantage of back-off times?
- Hope that an opposite operation arrives while waiting
- Match the two without interacting with the stack
- How??




Non-blocking stack – Elimination stack

- Pair the Treiber stack with an array
- Algorithm:
 1. Update the original stack via CAS
 2. If CAS fails, publish the operation in a random cell of the array

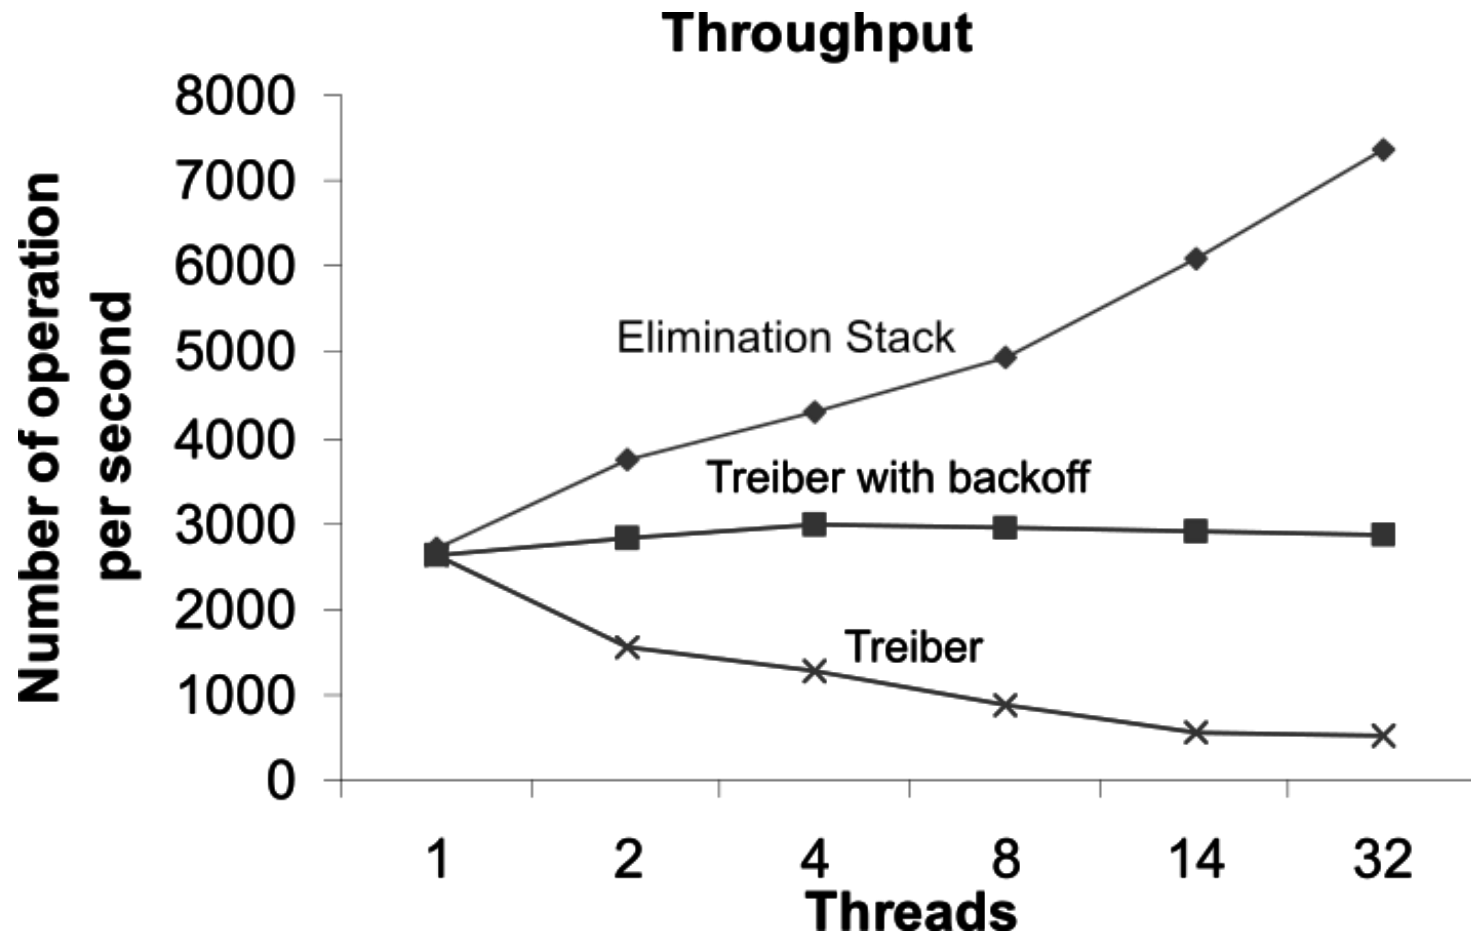


Non-blocking stack – Elimination stack

- Pair the Treiber stack with an array
- Algorithm:
 1. Update the original stack via CAS
 2. If CAS fails, publish the operation in a random cell of the array
 3. Wait for a matching operation
 4. If no matching op, GOTO 1



Non-blocking stack – Attempt 3



Concurrent Data Structures: **Sets**

Set implementations

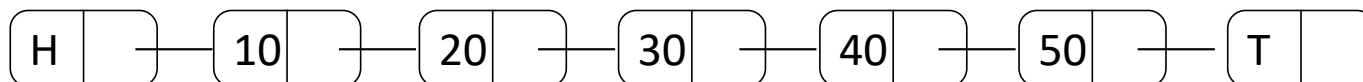
- Set methods:
 - insert(k)
 - delete(k)
 - ~~find(k)~~
- Implemented as an ordered linked list

INSERT(35)

INSERT(25)

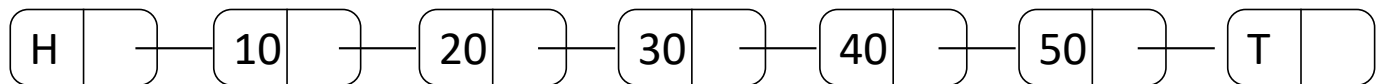
DELETE(40)

INSERT(55)

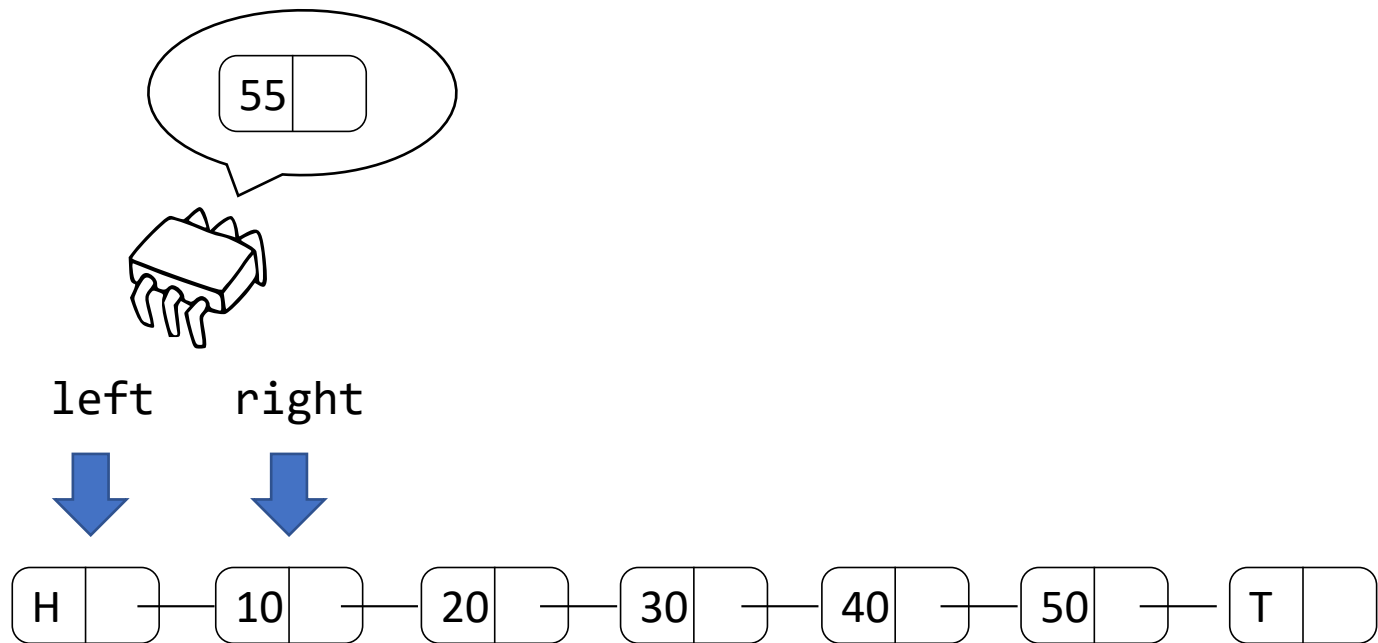


Insert algorithm


INSERT(55)



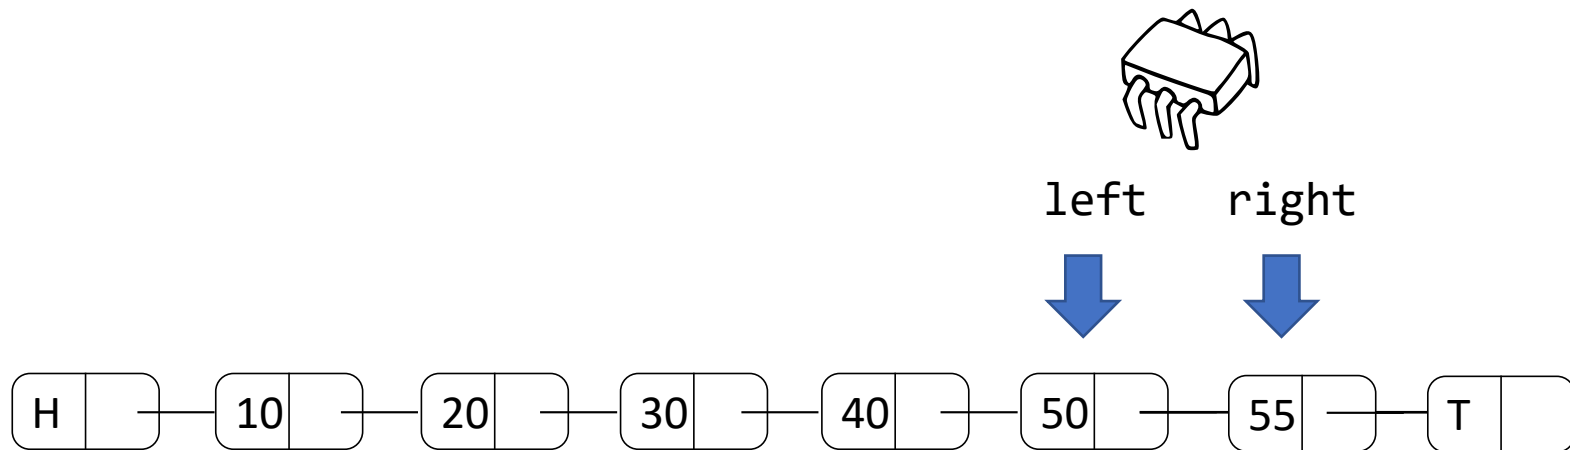
Insert algorithm



Insert algorithm

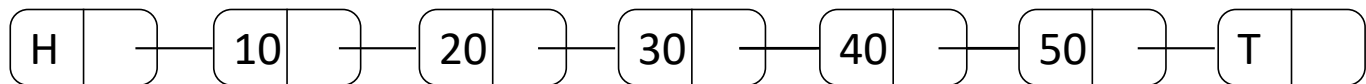


Insert algorithm

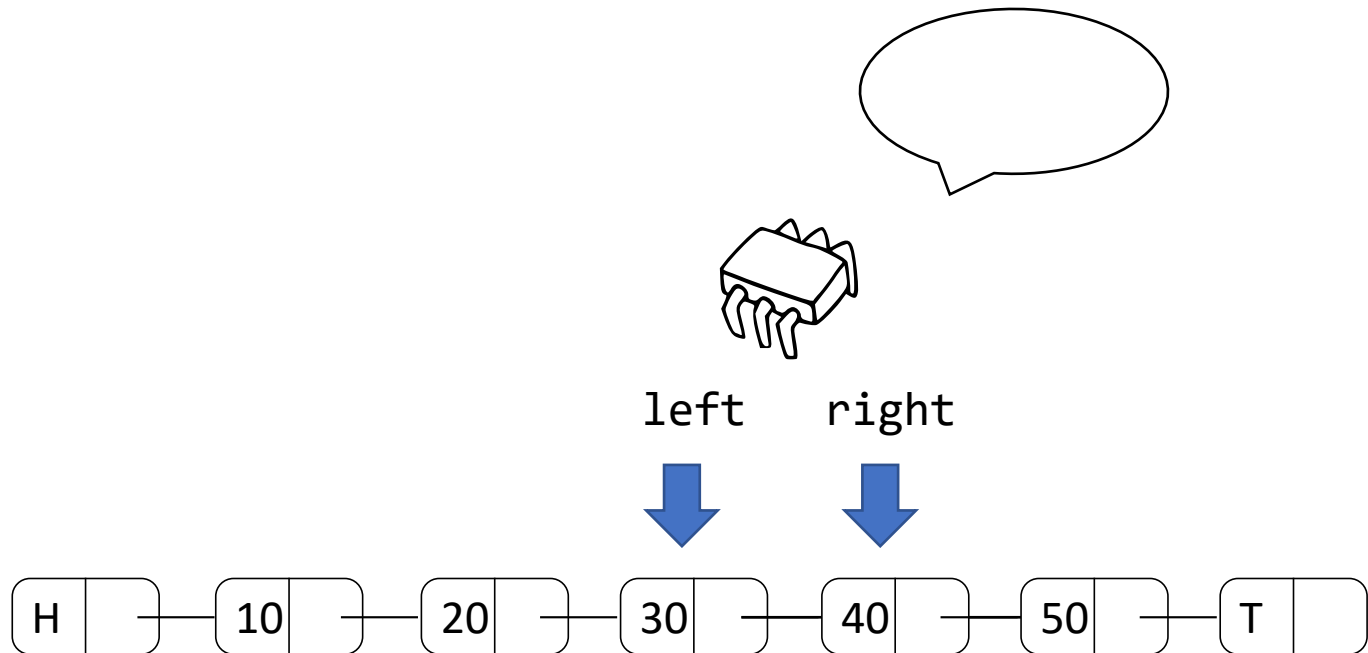


Delete algorithm

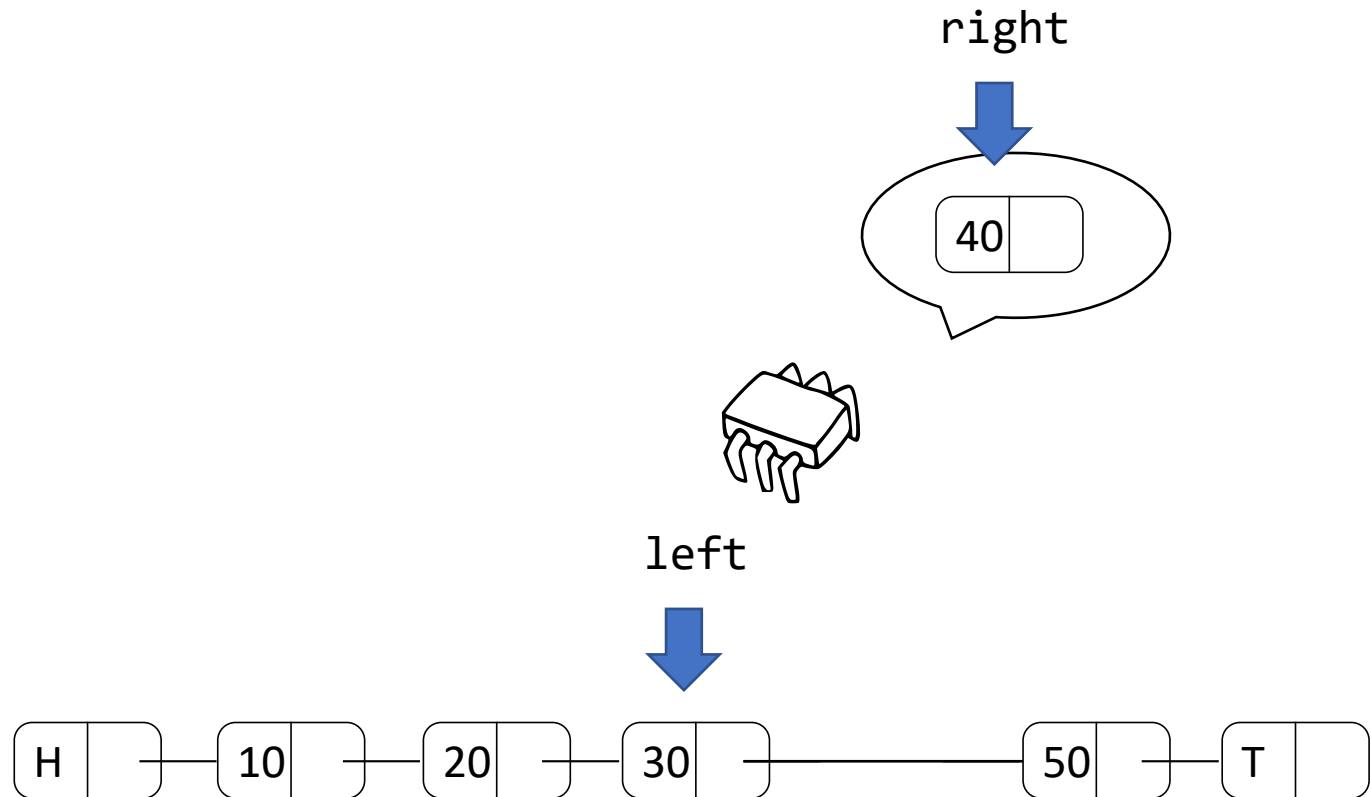
DELETE(40)



Delete algorithm



Delete algorithm



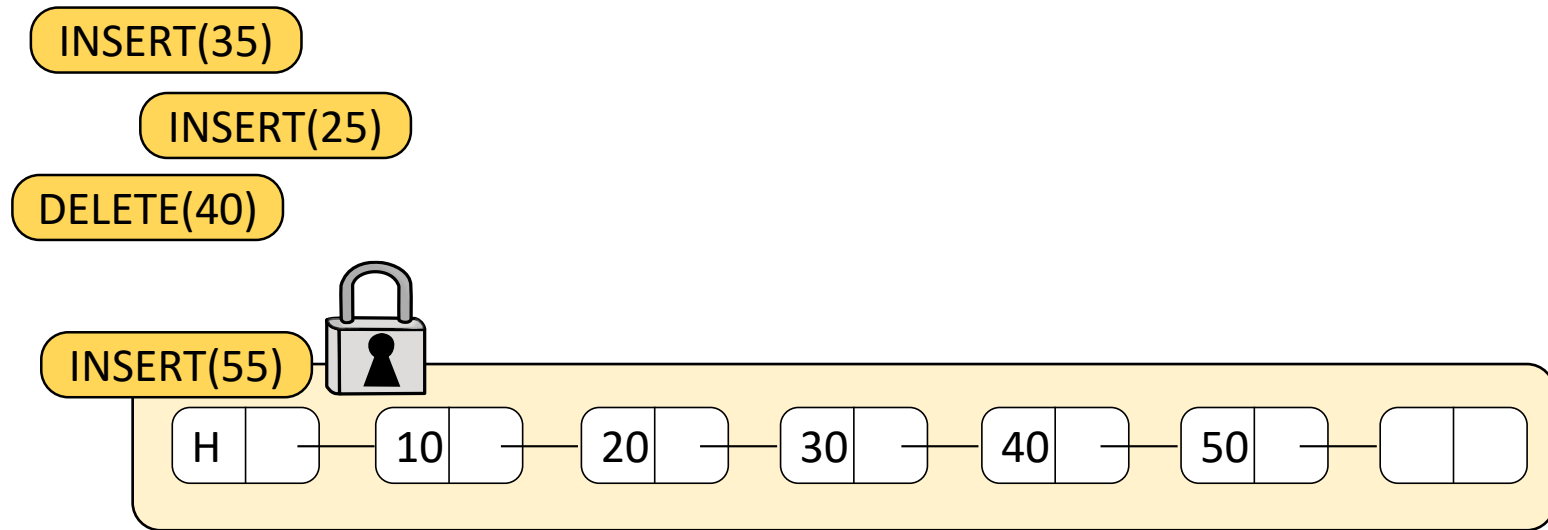
Sequential set implementation

```
1. bool do_operation(int k, int op_type){
2.     bool res = true;
3.     node *l,*r;
4.
5.     l = search(k, &r);
6.     switch(op_type){
7.         case(INSERT):
8.             if(r->key == k)
9.                 res = false;
10.            else
11.                l->next = new node(k,r);
12.            break;
13.        case(DELETE):
14.            if(r->key == k)
15.                l->next = r->next;
16.            else
17.                res = false;
18.            break;
19.    }
20.
21.
22.    return res;
23.}
```

```
1. node* search(int k, node **r){
2.     node *l, *r_next;
3.     l = set->head;
4.
5.     *r = l->next;
6.
7.     r_next = (*r)->next;
8.     while((*r)->key < k){
9.
10.        l = *r;
11.        *r = r_next;
12.
13.        r_next = (*r)->next;
14.    }
15.}
```

Concurrent set – Attempt 1

- PESSIMISTIC approach
- Synchronize via global lock



Concurrent set – Attempt 1 (SRC)

```
1. bool do_operation(int k, int op_type){
2.     bool res = true;
3.     node *l,*r;
4.     LOCK(&glock);
5.     l = search(k, &r);
6.     switch(op_type){
7.         case(INSERT):
8.             if(r->key == k)
9.                 res = false;
10.            else
11.                l->next = new node(k,r);
12.            break;
13.        case(DELETE):
14.            if(r->key == k)
15.                l->next = r->next;
16.            else
17.                res = false;
18.            break;
19.    }
20.    UNLOCK(&glock);
21.
22.    return res;
23.}
```

```
1. node* search(int k, node **r){
2.     node *l, *r_next;
3.     l = set->head;
4.
5.     *r = l->next;
6.
7.     r_next = (*r)->next;
8.     while((*r)->key < k){
9.
10.        l = *r;
11.        *r = r_next;
12.
13.        r_next = (*r)->next;
14.    }
15.}
```

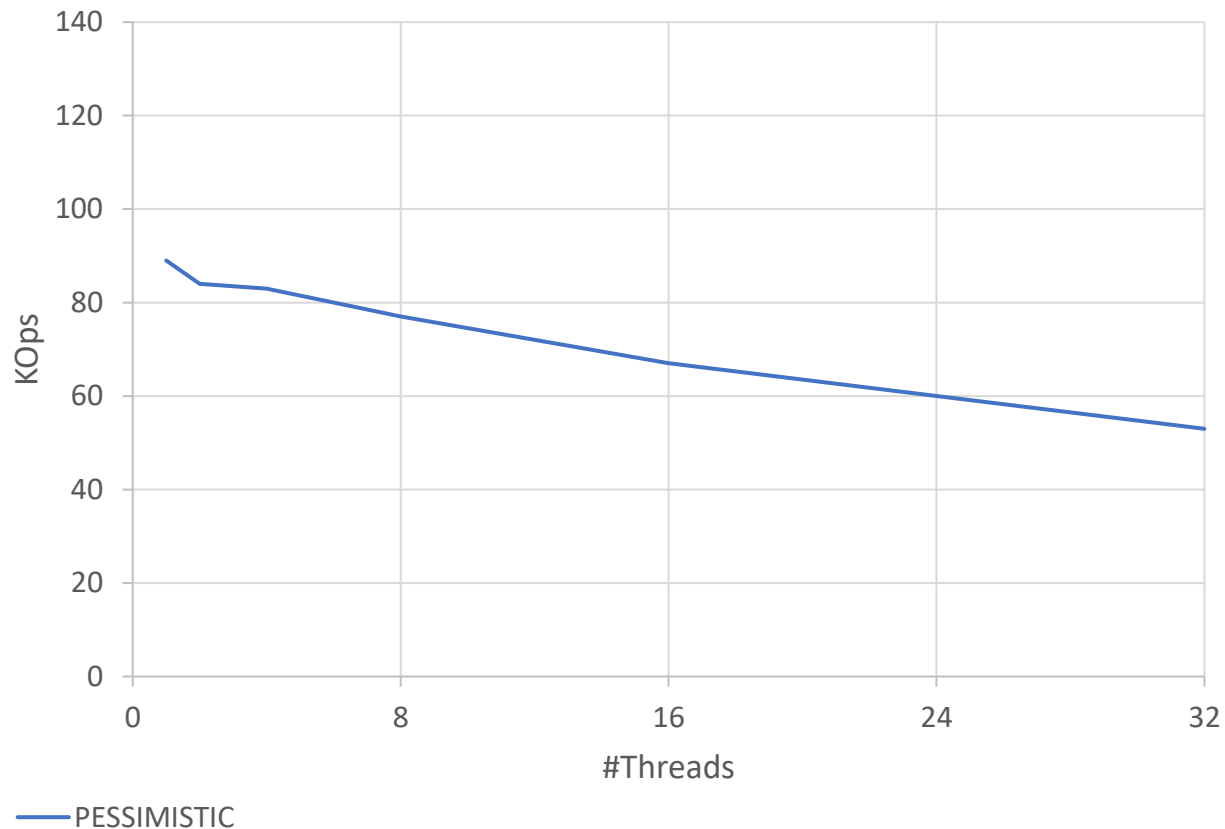
Concurrent set – Attempt 1

AMD Opteron 6128 – 32Cores

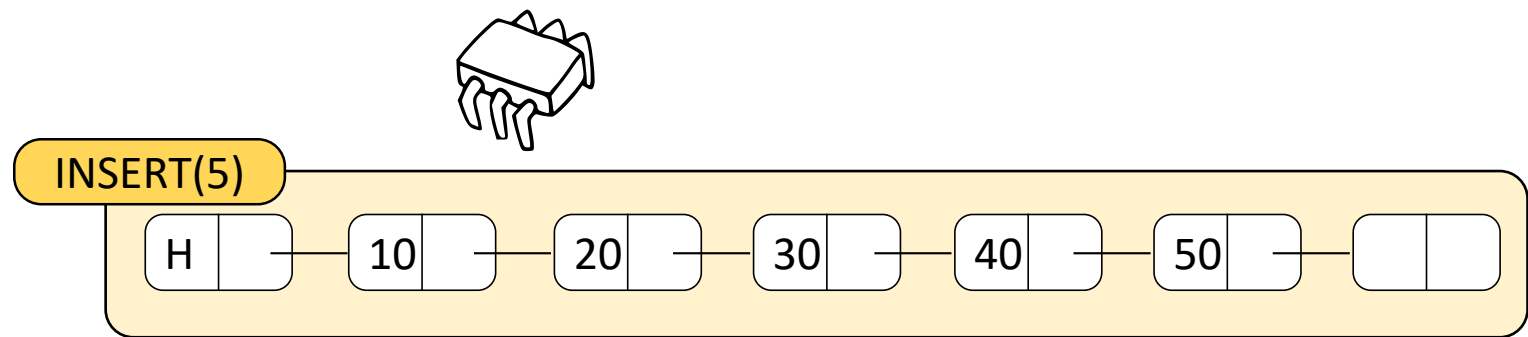
KeyRange = [0,6000]

SetSize = 2400

Update=100%

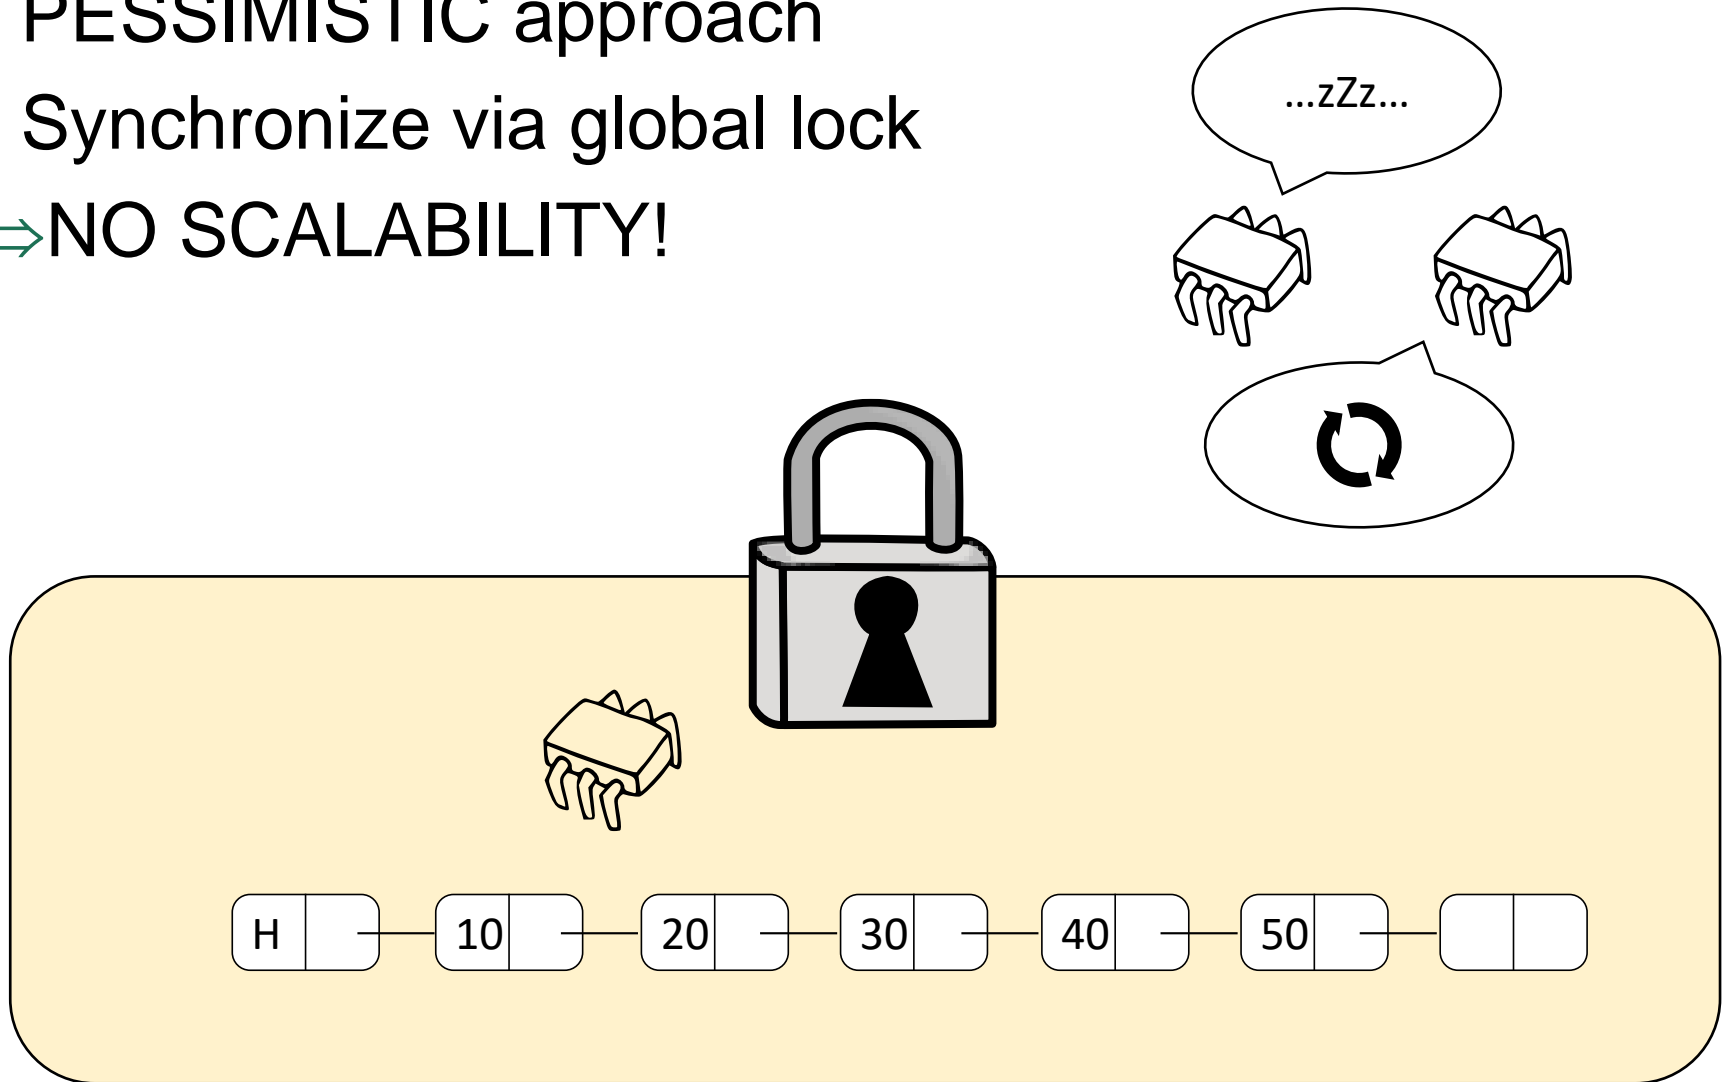


Concurrent set – Attempt 1



Concurrent set – Attempt 1

- PESSIMISTIC approach
 - Synchronize via global lock
- ⇒ NO SCALABILITY!



Concurrent set – Attempt 2

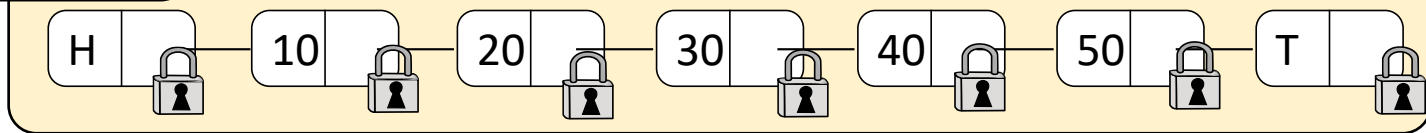
- Fine-grain approach
- Each node has its own lock
- Keep two locks at a time (lock coupling):
 - One on the current node
 - One on its predecessor

INSERT(35)

INSERT(25)

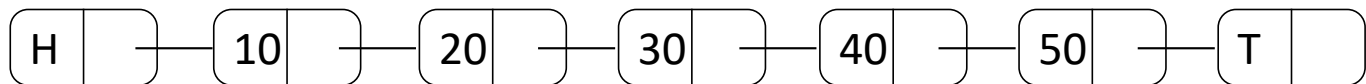
DELETE(40)

INSERT(55)



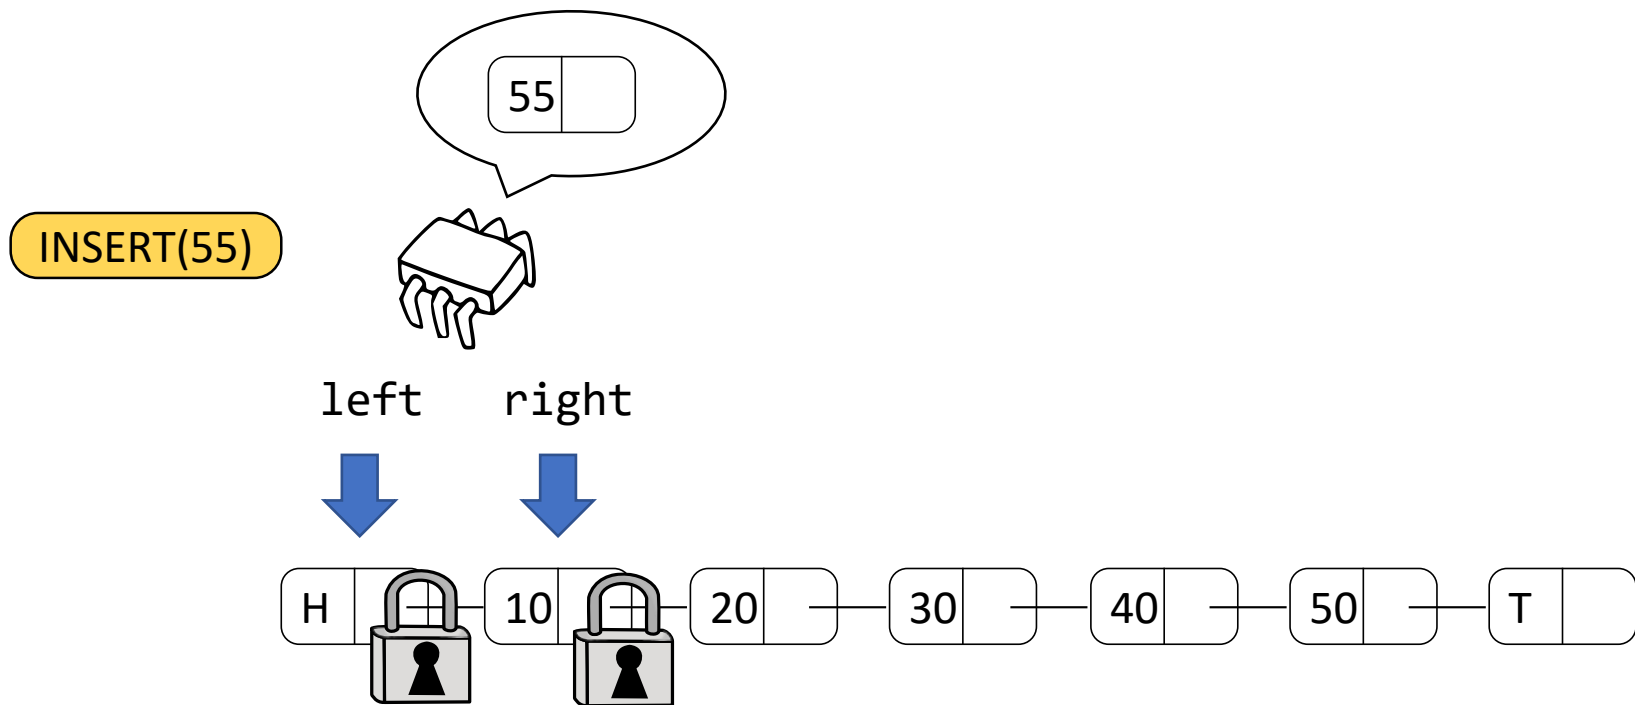
Search algorithm

INSERT(55)



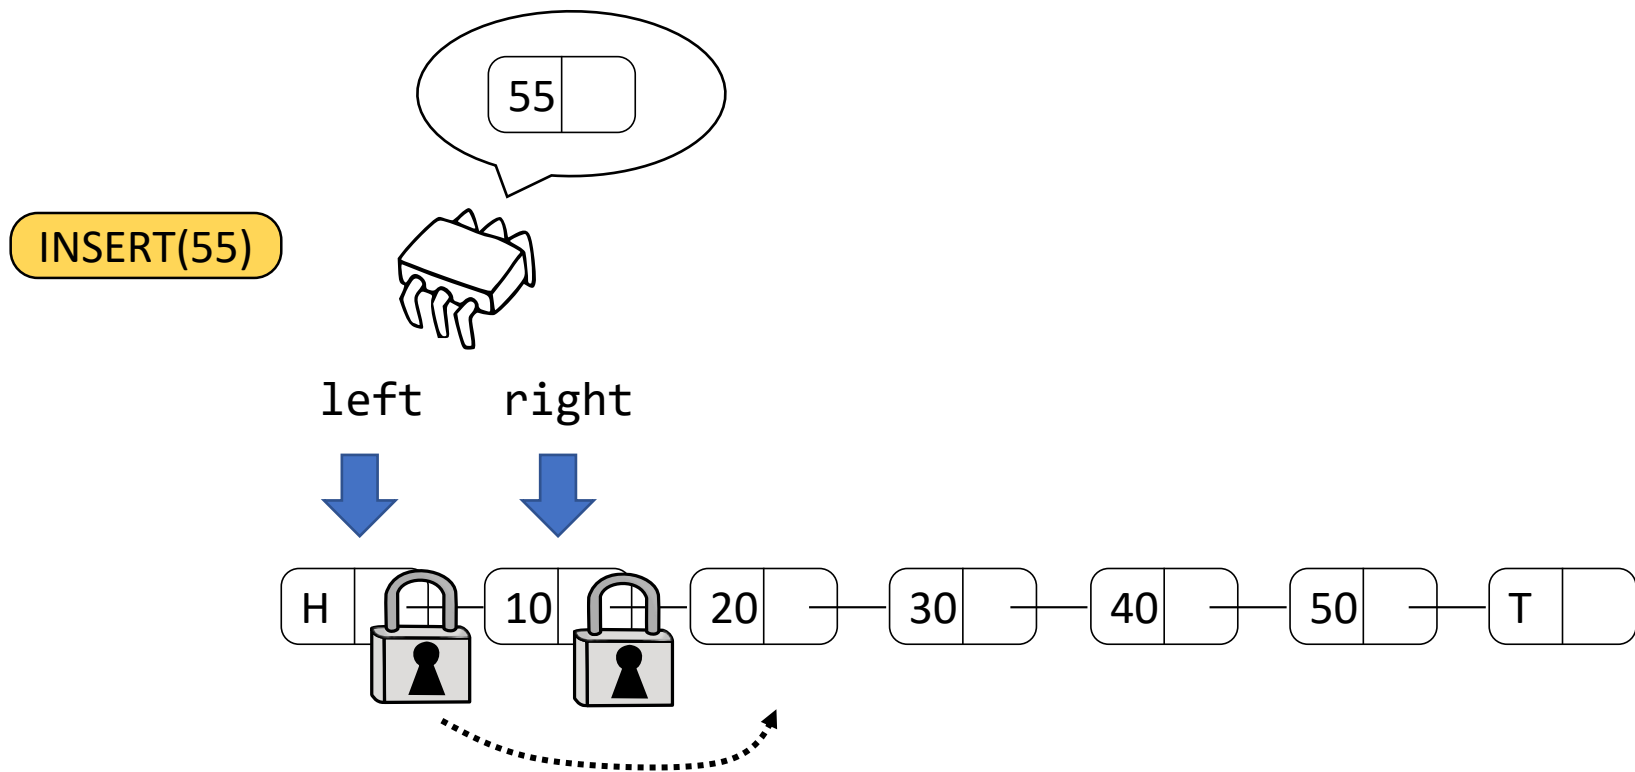
Search algorithm

- Keep two locks at a time (lock coupling):
 - One on the current node
 - One on its predecessor



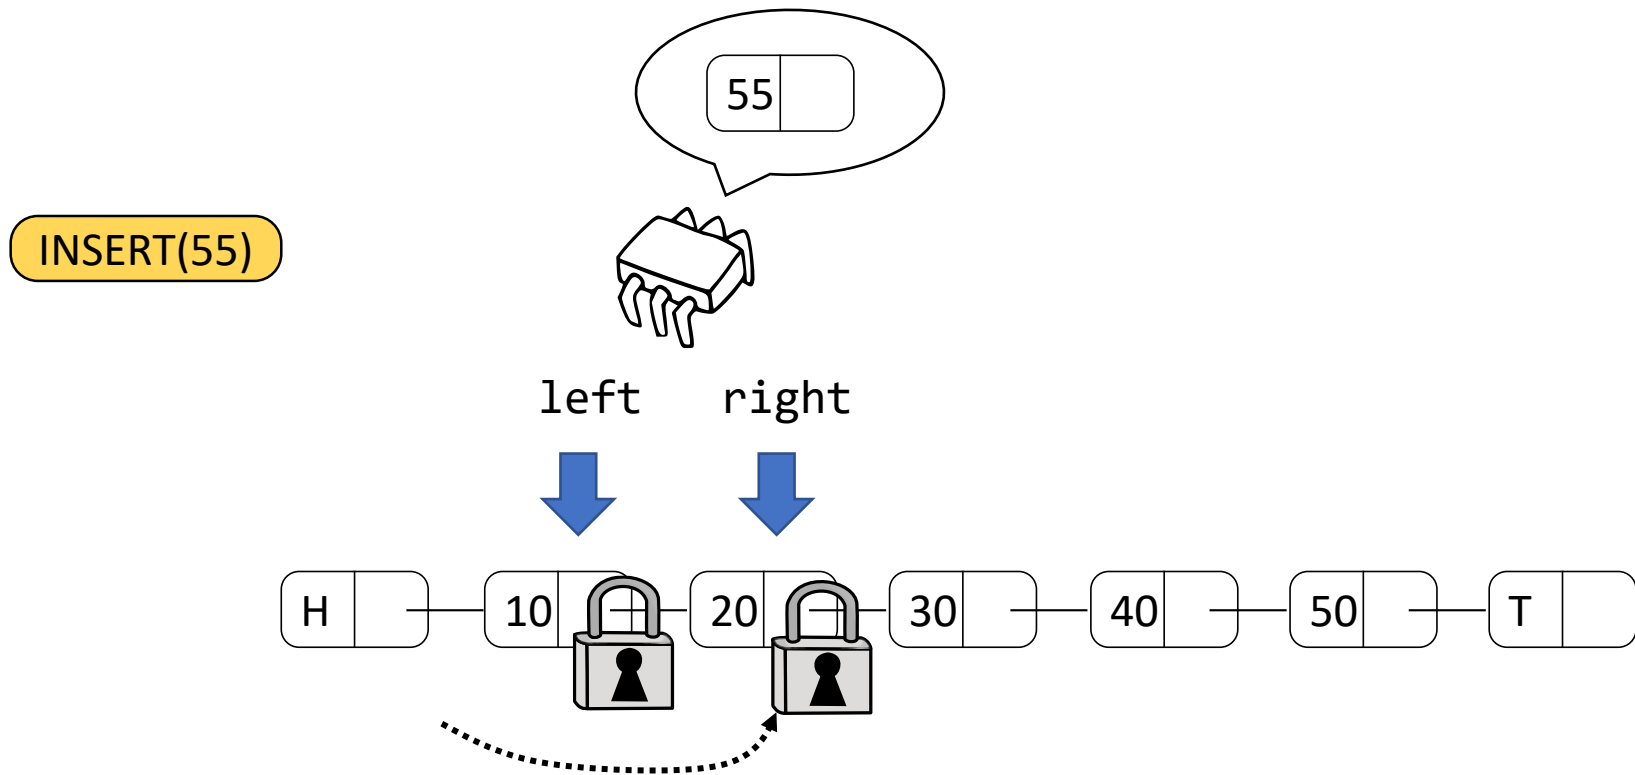
Search algorithm

- Keep two locks at a time (lock coupling):
 - One on the current node
 - One on its predecessor



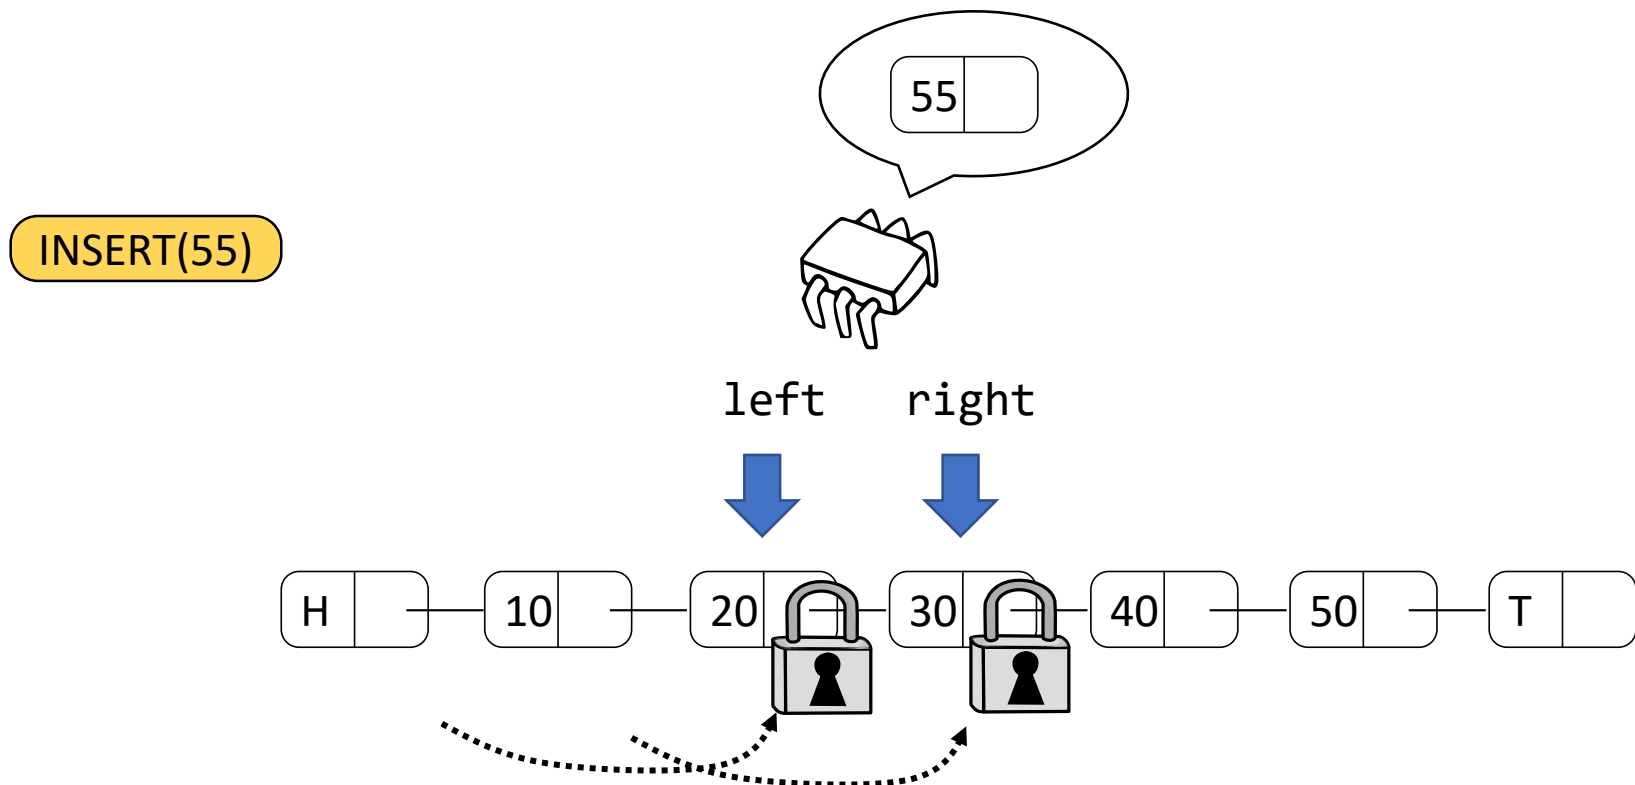
Search algorithm

- Keep two locks at a time (lock coupling):
 - One on the current node
 - One on its predecessor



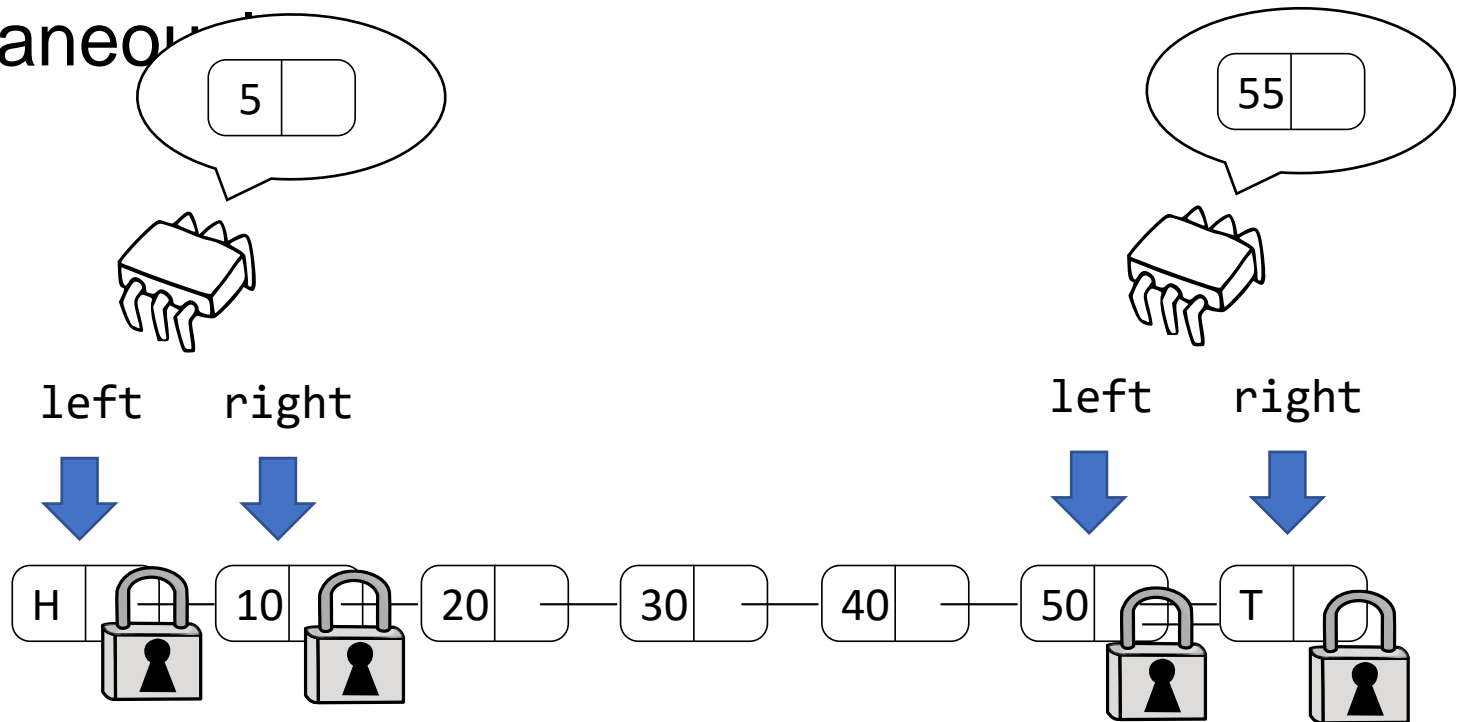
Search algorithm

- Keep two locks at a time (lock coupling):
 - One on the current node
 - One on its predecessor



Search algorithm

- Keep two locks at a time (lock coupling):
 - One on the current node
 - One on its predecessor
- Multiple threads access the data structure simultaneously



Concurrent set – Attempt 2 (SRC)

```
1. bool do_operation(int k, int op_type){
2.     bool res = true;
3.     node *l,*r;
4.     LOCK(&glock);
5.     l = search(k, &r);
6.     switch(op_type){
7.         case(INSERT):
8.             if(r->key == k)
9.                 res = false;
10.            else
11.                l->next = new node(k,r);
12.            break;
13.        case(DELETE):
14.            if(r->key == k)
15.                l->next = r->next;
16.            else
17.                res = false;
18.            break;
19.    }
20.    UNLOCK(&glock);
21.    UNLOCK(&l->lock);
22.    UNLOCK(&r->lock);
23.    return res;
24. }
```

```
1. node* search(int k, node **r){
2.     node *l, *r_next;
3.     l = set->head;
4.     LOCK(&l->lock);
5.     *r = l->next;
6.     LOCK(&(*r)->lock);
7.     r_next = (*r)->next;
8.     while((*r)->key < k){
9.         UNLOCK(&l->lock);
10.        l = *r;
11.        *r = r_next;
12.        LOCK(&(*r)->lock);
13.        r_next = (*r)->next;
14.    }
15. }
```

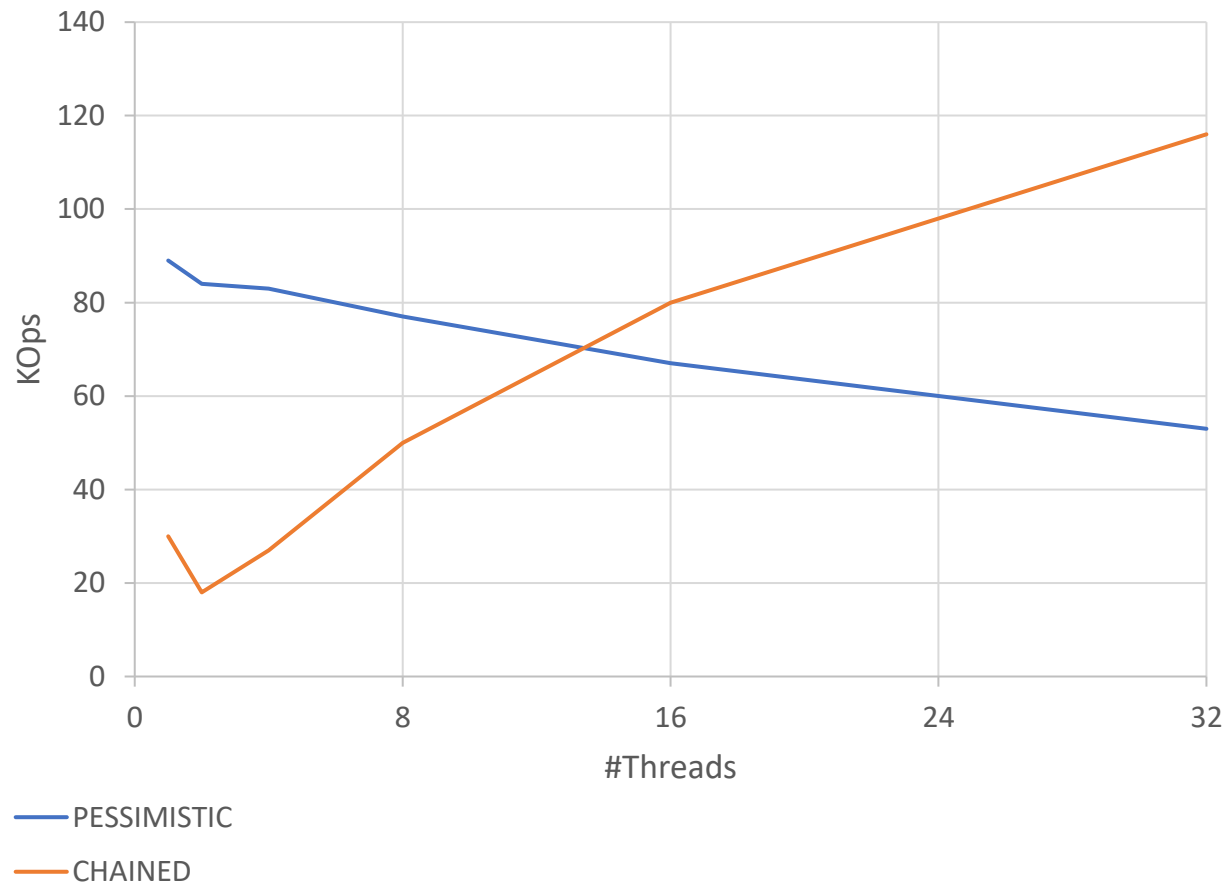
Concurrent set – Attempt 2

AMD Opteron 6128 – 32Cores

KeyRange = [0,6000]

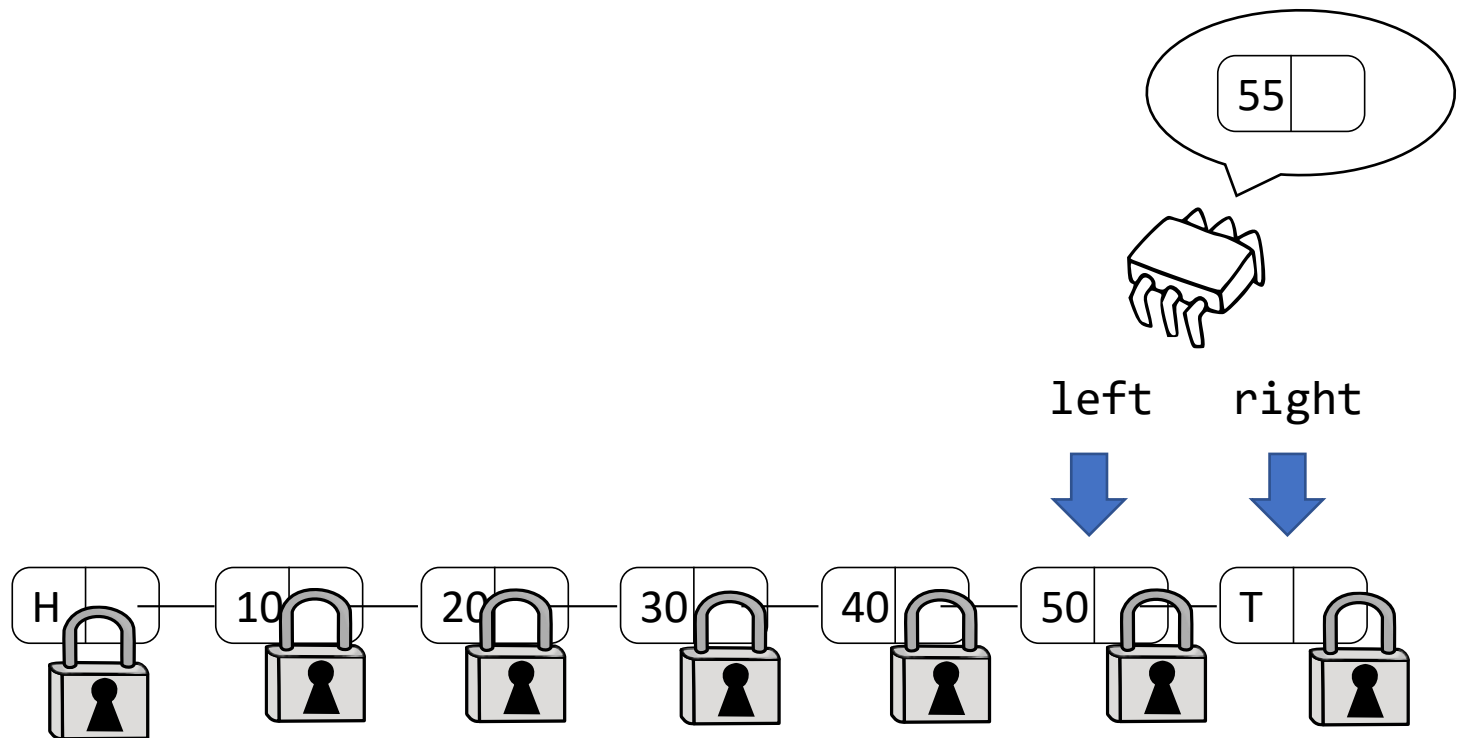
SetSize = 2400

Update=100%



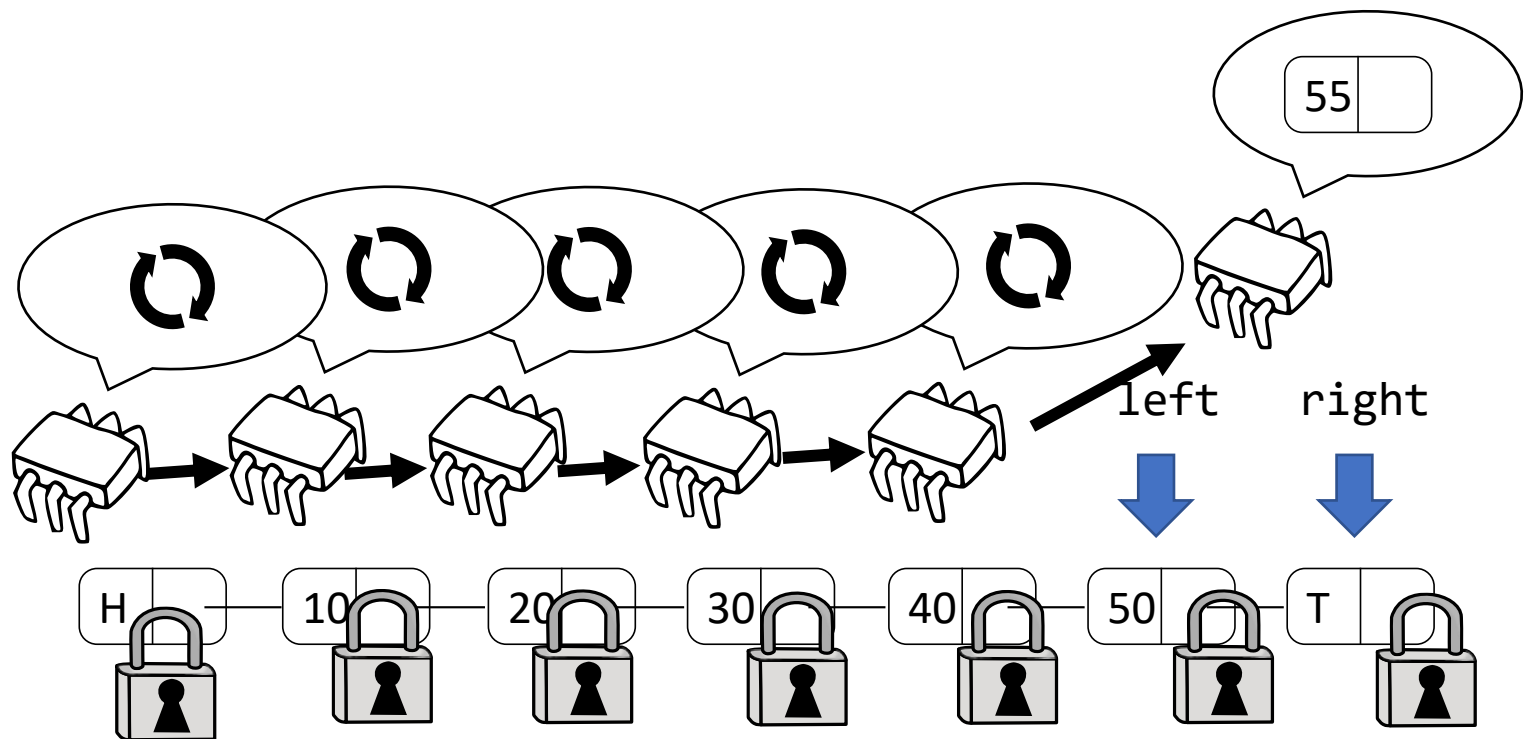
Search algorithm

- Allows an increased parallelism but...



Search algorithm

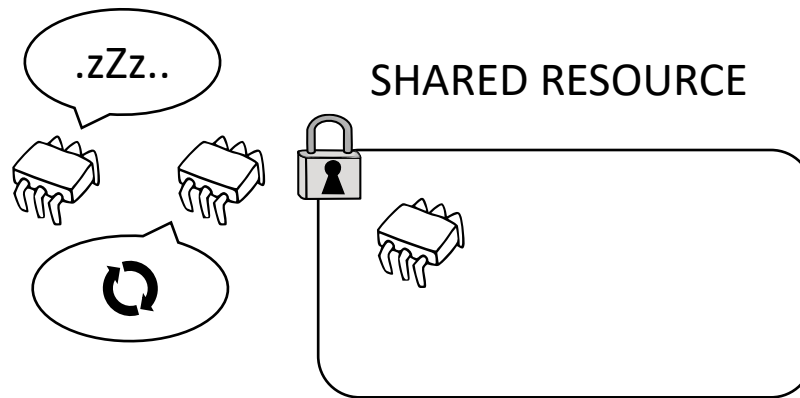
- Allows an increased parallelism but...
- High costs for lock handover



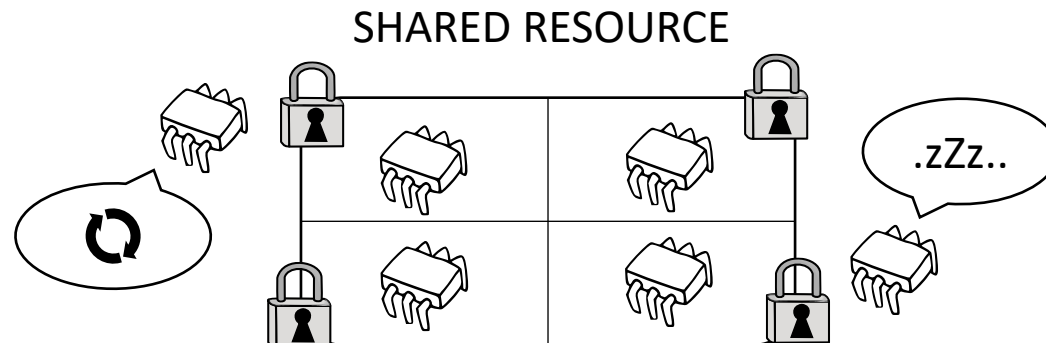
Recap

- Explored two blocking strategies:

1. Global (coarse-grain) lock




2. (Fine-grain) Lock coupling



Concurrent set – Attempt 3

DELETE(40)

INSERT(55)

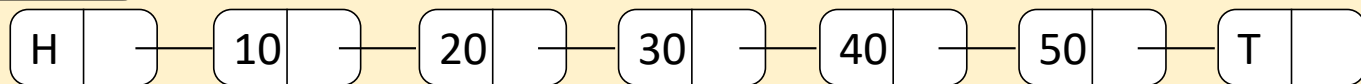


Concurrent set – Attempt 3

- NON-BLOCKING approach [Harris linked list]
- Search without acquiring any lock
- Apply updates with individual atomic instructions

DELETE(40)

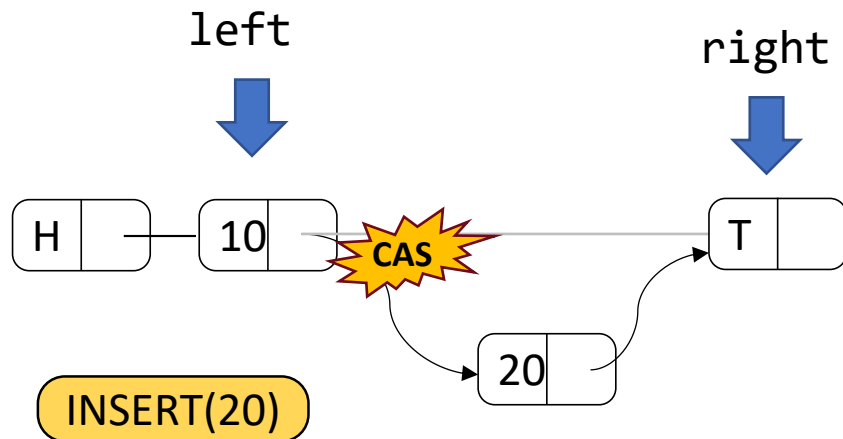
INSERT(55)



Non-blocking insert & delete algorithms


Insert:

1. Search left and right nodes
2. Insert the new item with a CAS
3. If CAS fails restart from 1



Delete:

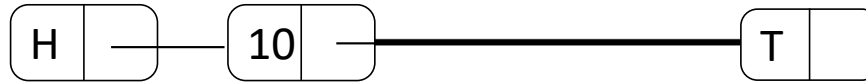
1. Search left and right nodes
2. Disconnect the item with a CAS
3. If CAS fails restart from 1



- Is it correct?


Incorrect delete algorithm

- Edge cases might lead to losing items!



Incorrect delete algorithm

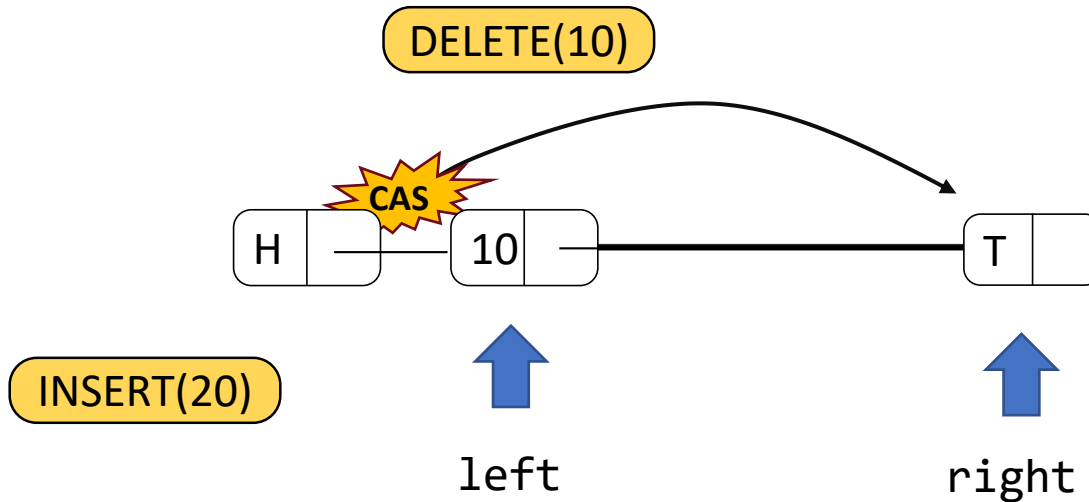
- Edge cases might lead to losing items!



1. Thread A gets left and right node and go to sleep
2. Thread B disconnects the node containing 10
3. Thread A wakes up and add 20 after 10
4. The new item is lost

Incorrect delete algorithm

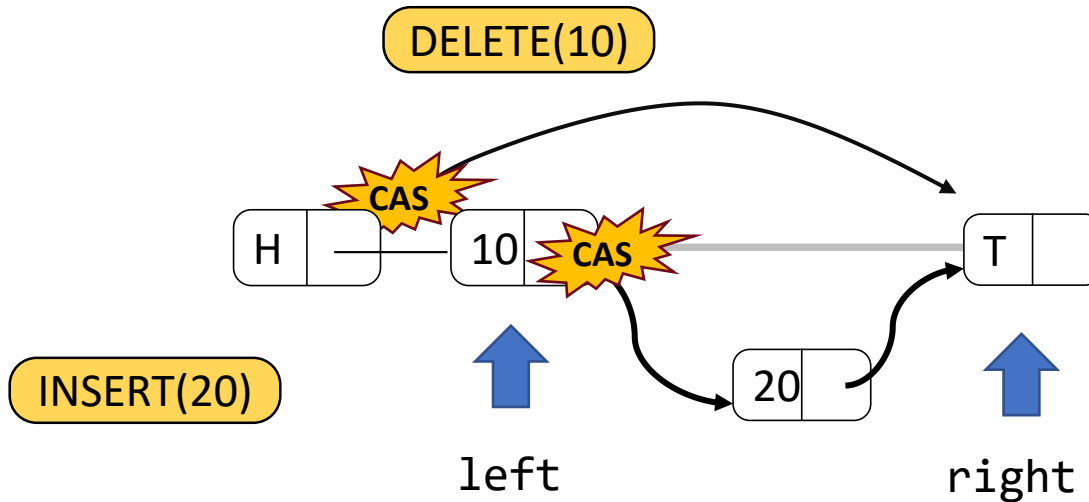
- Edge cases might lead to losing items!



1. Thread A gets left and right node and go to sleep
2. Thread B disconnects the node containing 10
3. Thread A wakes up and add 20 after 10
4. The new item is lost

Incorrect delete algorithm

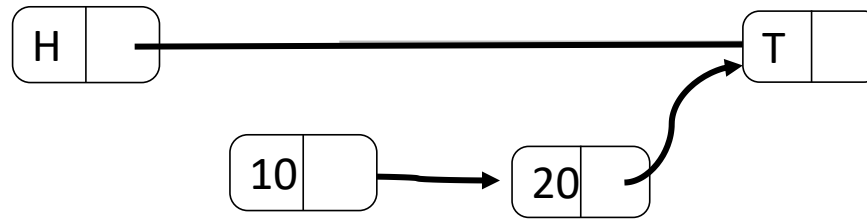
- Edge cases might lead to losing items!



1. Thread A gets left and right node and go to sleep
2. Thread B disconnects the node containing 10
3. Thread A wakes up and add 20 after 10
4. The new item is lost

Incorrect delete algorithm

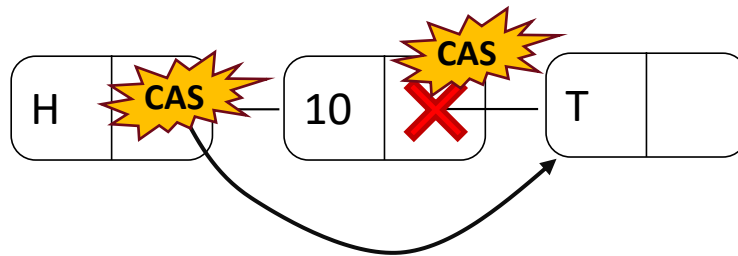
- Edge cases might lead to losing items!



1. Thread A gets left and right node and go to sleep
2. Thread B disconnects the node containing 10
3. Thread A wakes up and add 20 after 10
4. The new item is lost

The correct delete algorithm

- Adopt logical deletion:
 1. Get left and right node
 2. Mark the item as deleted via CAS (*logical* deletion)
 3. If CAS fails GOTO 1
 4. Disconnect the item via CAS (*physical* deletion)
 5. If CAS fails GOTO 4



The correct delete algorithm

- Adopt the correct delete algorithm

1. Get

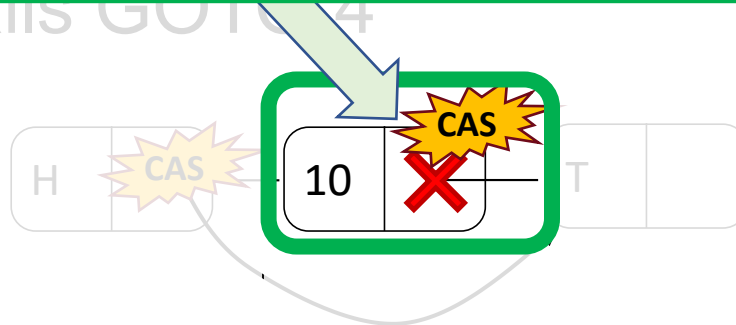
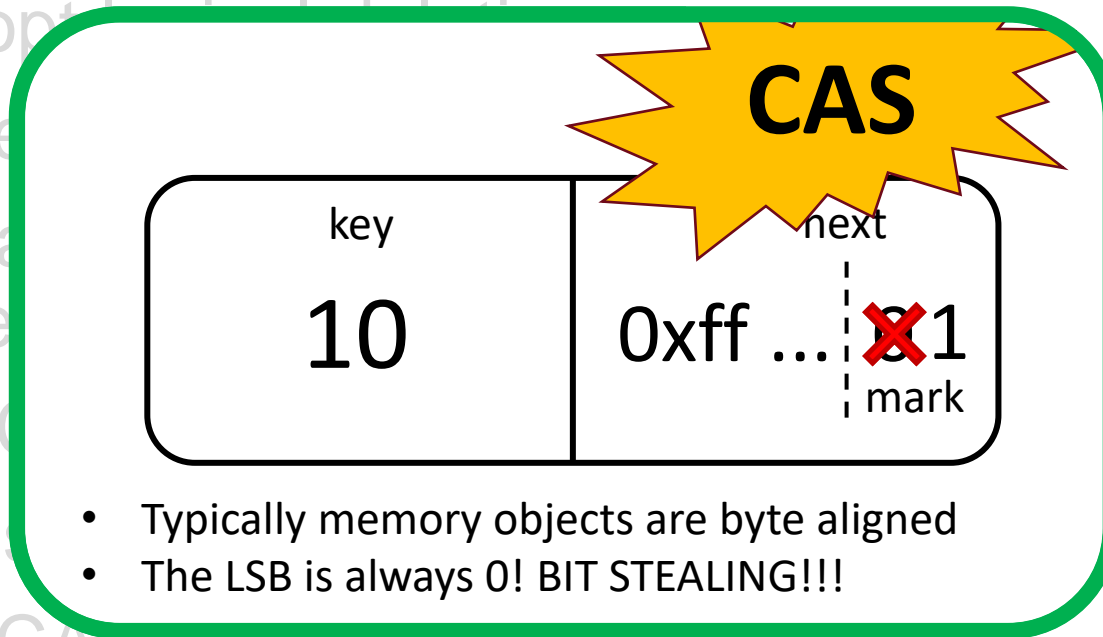
2. Make

de

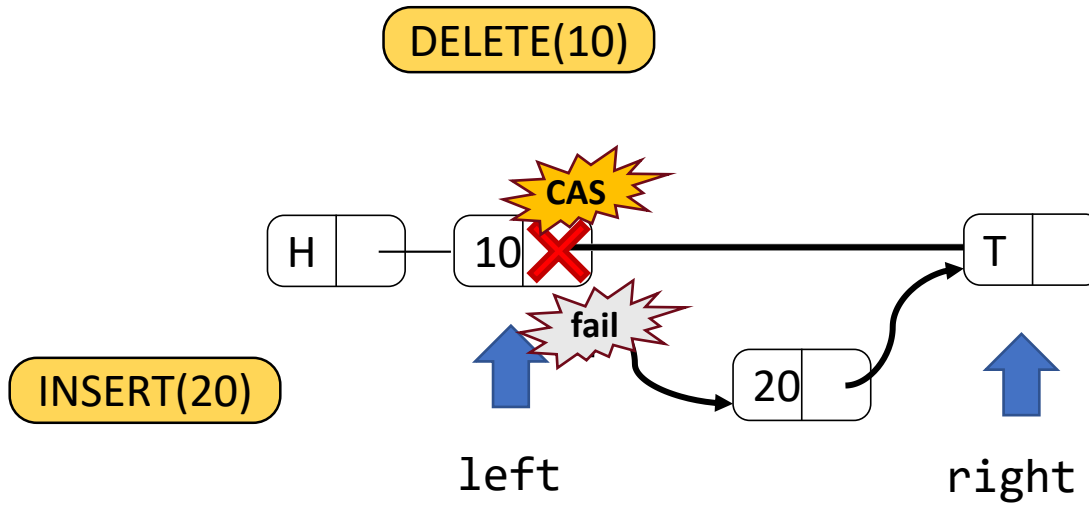
3. If

4. Dis

5. If CAS fails GOTO 4



The correct delete algorithm



- Updates of the "next" field by two opposite concurrent operations cannot both succeed
- What to do upon conflict (failed CAS)? **RESTART FROM SCRATCH!!**


Non-blocking search

- The search returns two adjacent non-marked (left and right) nodes
- Housekeeping: disconnect logically delete items during searches



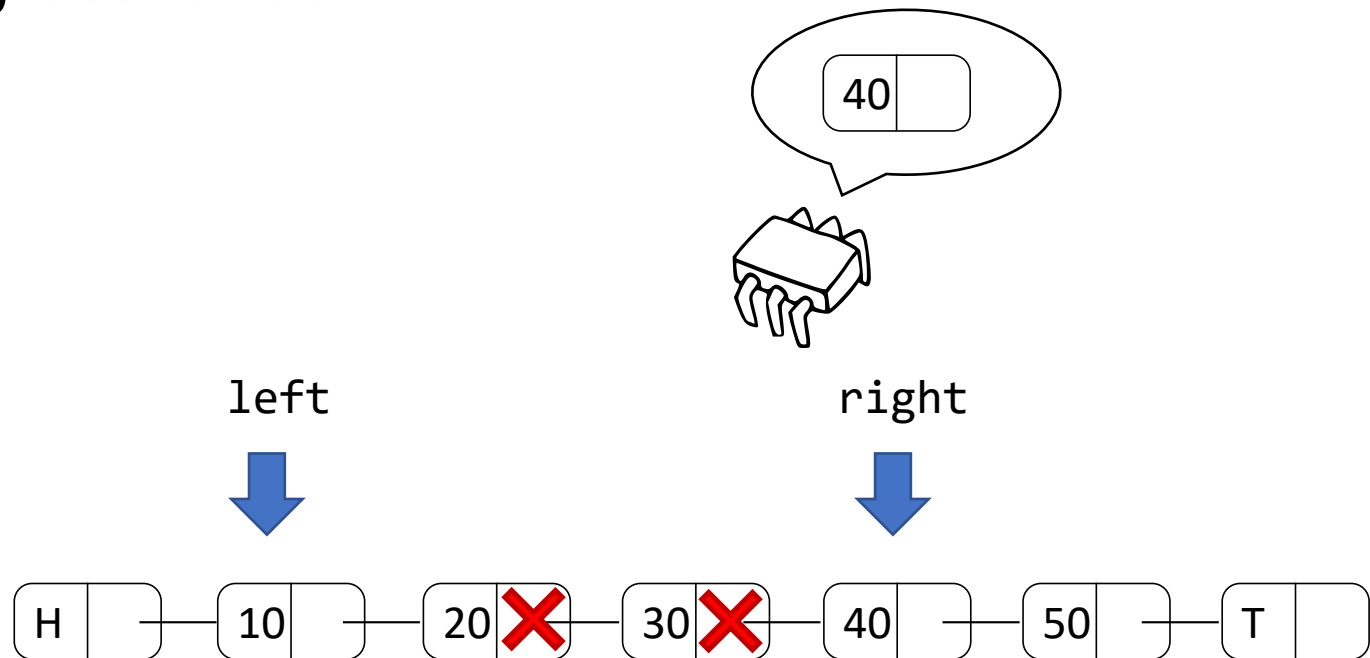
Non-blocking search

- The search returns two adjacent non-marked (left and right) nodes
- Housekeeping: disconnect logically delete items during searches



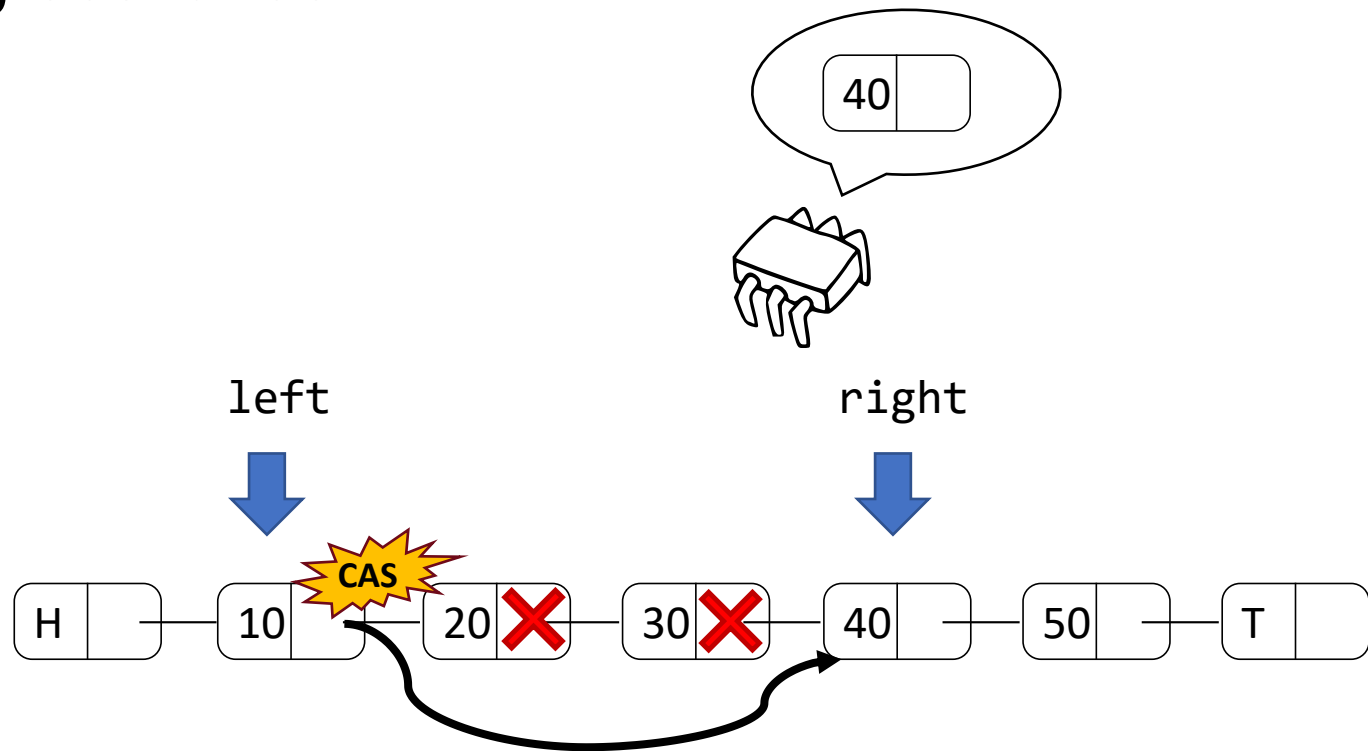
Non-blocking search

- The search returns two adjacent non-marked (left and right) nodes
- Housekeeping: disconnect logically delete items during searches



Non-blocking search

- The search returns two adjacent non-marked (left and right) nodes
- Housekeeping: disconnect logically delete items during searches



Concurrent set – Attempt 3 (SRC)

```
1. bool do_operation(int k, int op_type){
2.     node *l,*r, *n = new node(k);
3.     l = search(k, &r);                /* get left and right node */
4.     switch(op_type){
5.         case(INSERT):
6.             if(r->key == k) return false; /* key present in the set */
7.             n->next = r;
8.             l->next = n;                /* insert the item */
9.
10.
11.         break;
12.         case(DELETE):
13.             if(r->key != k) return false; /* key not present */
14.             l->next = r->next;           /* remove the key */
15.
16.
17.
18.         break;
19.     }
20.     return true;
21. }
```


Concurrent set – Attempt 3 (SRC)

```
1. bool do_operation(int k, int op_type){
2.     node *l,*r, *n = new node(k);
3.     l = search(k, &r);                               /* get left and right node */
4.     switch(op_type){
5.         case(INSERT):
6.             if(r->key == k) return false; /* key present in the set */
7.             n->next = r;
8.             l->next = n; /* insert the item */
9.             if(!CAS(&l->next, r, n))
10.                goto 3; /* insertion failed the item -> restart */
11.             break;
12.         case(DELETE):
13.             if(r->key != k) return false; /* key not present */
14.             l->next = n->next; /* remove the key */
15.             if(is_marked_ref((l=r->next)) || !CAS(&r->next, l, mark(l)))
16.                goto 3; /* insertion failed the item -> restart */
17.             search(k,&r); /* repeat search */
18.             break;
19.     }
20.     return true;
21. }
```

Concurrent set – Attempt 3 (SRC)

```
1. node* search(int k, node **r){
2.     node *l, *t, *t_next, *l_next;
3.     *t = set->head;
4.     t_next = t->head->next;
5.     while(1){                                /* FIND LEFT AND RIGHT NODE */
6.         if(!is_marked(t_next)){
7.             l = t;
8.             l_next = t_next;
9.         }
10.        t = get_unmarked_ref((t_next));
11.        t_next = t->next;
12.        if(!is_marked_ref(t_next) && t->key >= k) break;
13.    }
14.    *r = t;
15.    /* DEL MARKED NODES */
16.    if(l_next != *r && !CAS(&l->next, l_next, *r) goto 3;
17.    return l;
18.}
```

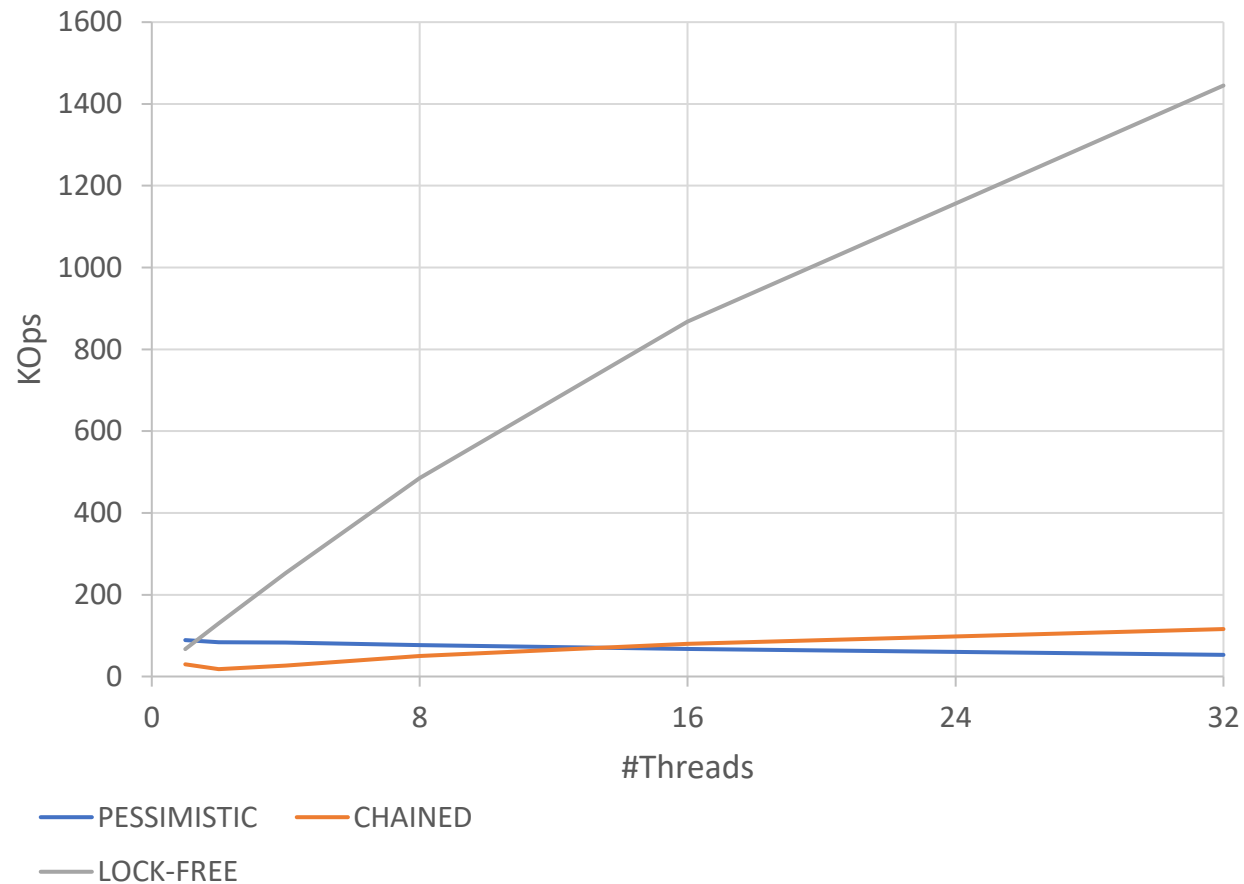
Concurrent set – Attempt 3

AMD Opteron 6128 – 32Cores

KeyRange = [0,6000]

SetSize = 2400

Update=100%



Safety and liveness guarantees

- The algorithm is **NON-BLOCKING (LOCK-FREE)**:
 - If a thread A is stuck in a retry loop => a CAS fails each time
 - If a CAS fail, it is because of another CAS that was successfully executed by a thread B
 - Thread B is making progress
- The algorithm is **LINEARIZABLE**:
 - Each method execution take effect in an atomic point (a successful CAS) contained between its invocation and reply
 - The order obtained by using these points has been proved to be compliant with the Set semantic

Progress (Lock freedom)

- Each method update method has two main steps
 - A search, which might end with a CAS
 - A CAS to insert delete a node
- 1. Suppose an update method is stuck in a search:
 - The key range is finite, so the number of node is finite
 - It continuously fails to disconnect marked nodes
 - It means that new nodes have been both inserted and marked!
 - Other threads have completed update methods
- 2. Suppose an updated method always fails its last step (insertion or marking)
 - Other threads have modified the target next pointer
 - If it is due to the disconnection of marked nodes, see point 1
 - If it is due to the updated step other methods have completed

Safety (Linearizability)

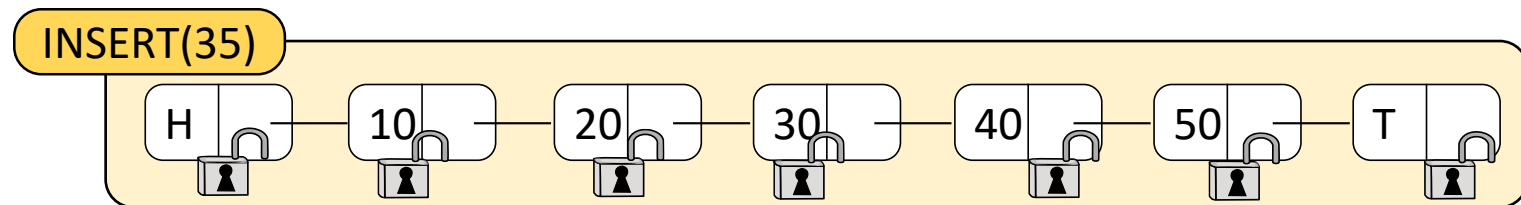
1. The search returns 2 adjacent nodes in an atomic point
 1. The read of next field of the left node
 2. The CAS that make left and right adjacent
- It is like that the search made a snapshot of interested key interval
2. Find, unsuccessful delete and unsuccessful insert linearize with the search (1.1 or 1.2)
3. Insert linearizes with the successful CAS to connect a new node
4. Delete linearizes with the successful CAS to mark a node

Problems

- It is not possible to flip a bit of a reference on memory-managed languages (e.g. JAVA)
- How to solve?

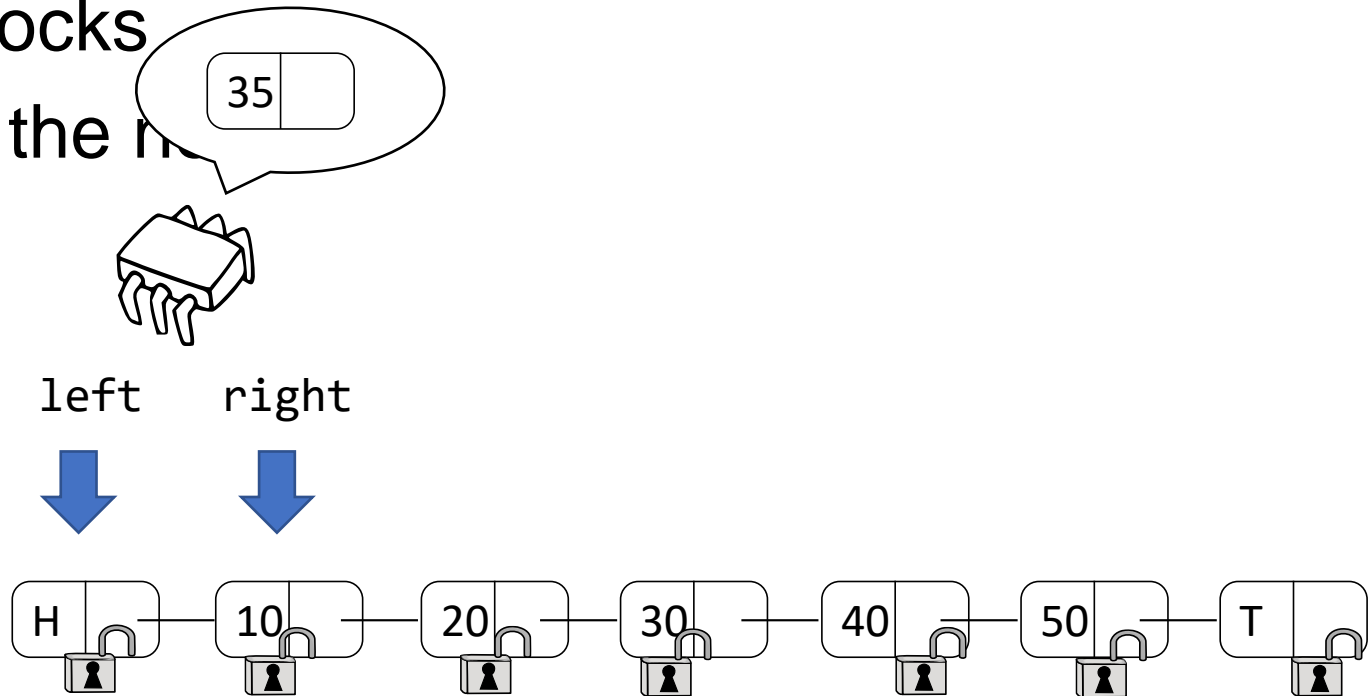
Locks + Optimism

- Use one lock per node
- Move “marked” to a dedicated field

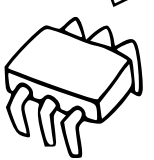


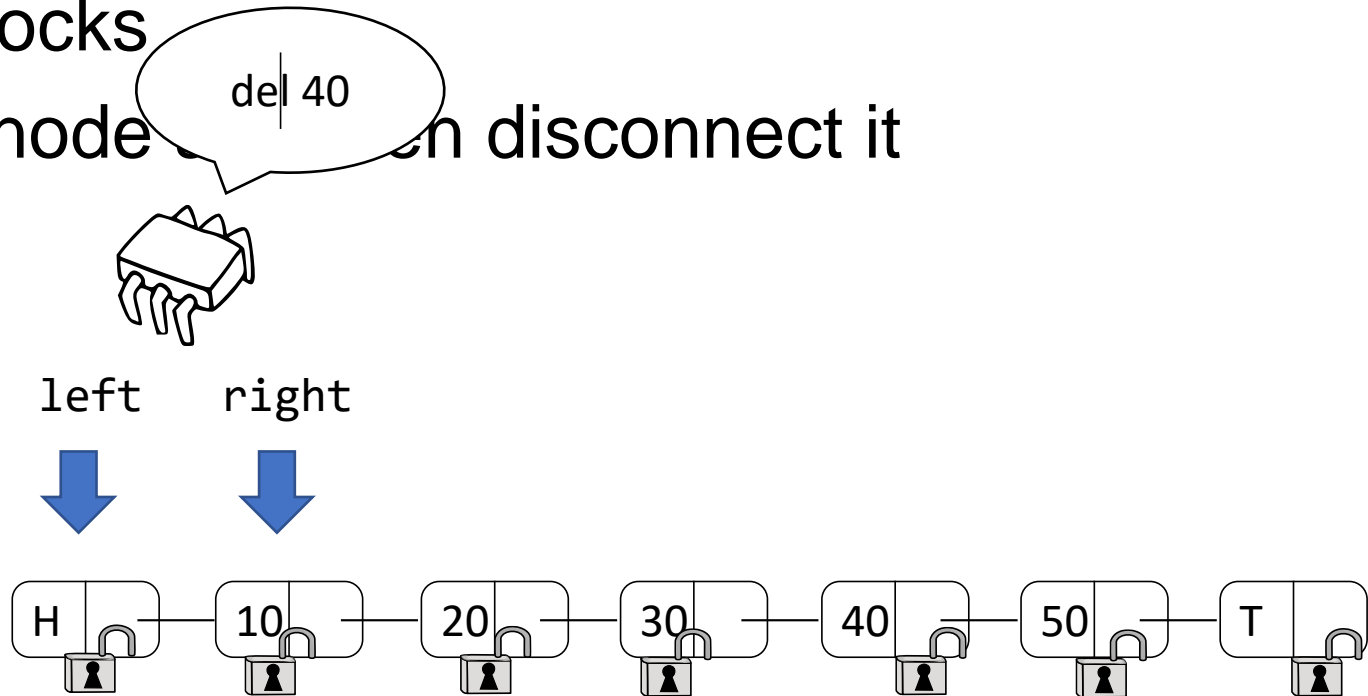
Locks + Optimism (insert)

- Use one lock per node
- Move “marked” to a dedicated field
- Find left and right without taking locks!
- Take locks
- Insert the node



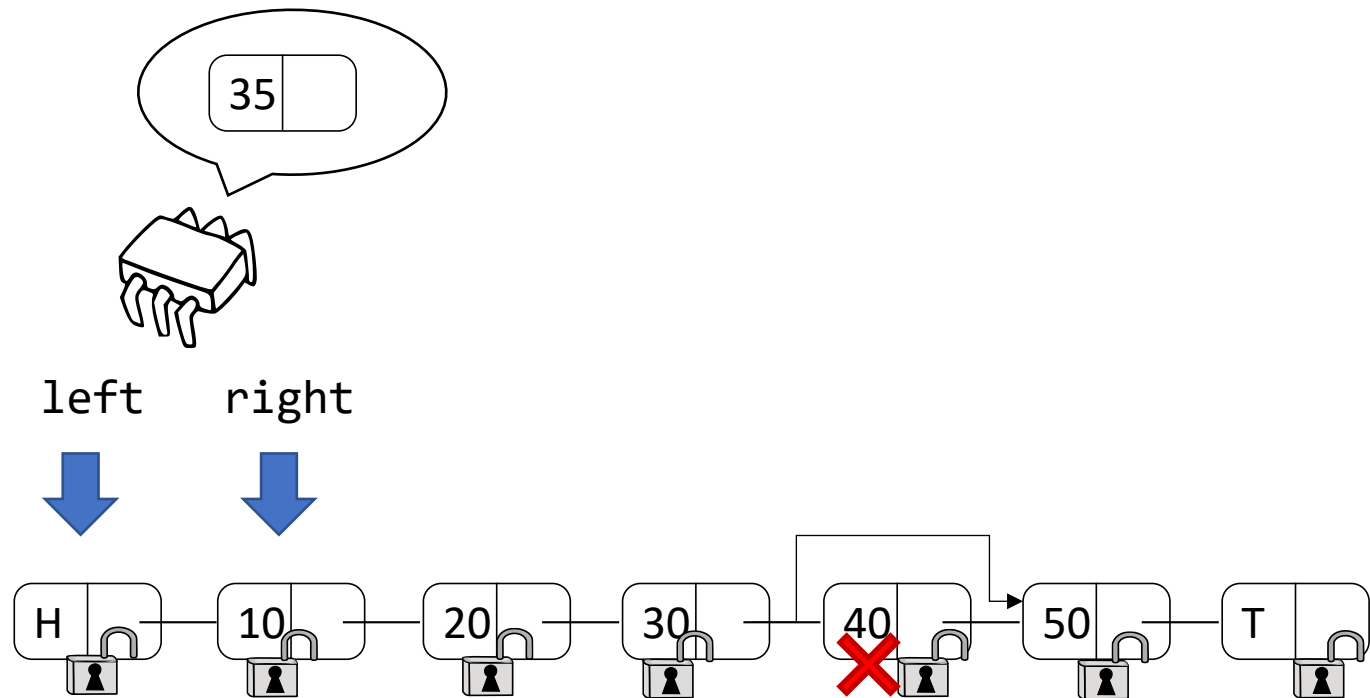
Locks + Optimism (delete)

- Use one lock per node
- Move “marked” to a dedicated field
- Find left and right without taking locks!
- Take locks
- Mark node  when disconnect it



Locks + Optimism (delete)

- Why “optimistic”? Do work (search) and hope nothing wrong happens!
- What could go wrong?



Locks + Optimism (delete)

- Why “optimistic”? Do work (search) and hope nothing wrong happens!
- What could go wrong?
 - Left and/or right being marked
 - Left and right not adjacent
- How to solve?
- Validation of search results:
 - Left unmarked
 - Right unmarked
 - Left.next = right

Locks + Optimism (delete)

- Why “optimistic”? Do work (search) and hope nothing wrong happens!
- What could go wrong?
 - Left and/or right being marked
 - Left and right not adjacent
- How to solve?
- Validation of search results:
 - Left unmarked
 - Right unmarked
 - Left.next = right

Locks + Optimism = Lazy List

- What about correctness?
- What about progress?

Can we do better?


- Costs: $O(n)$
- Starting from scalable “simple” set implementation we can build faster set implementations
 - Hash table: $O(1)$
 - Array of buckets
 - Buckets are concurrent ordered-list based sets
- We know that a search in an ordered set could be more efficient $O(\log(n))$
- How?

Skip list [Pugh 1990]

- Generalization of sorted linked lists
- Randomized data structure
- Costs: $O(\log(n))$
- Idea:
 1. Maintain a core sorted linked list L_0
 2. Use additional sorted linked lists L_i such that:
 1. $L_i \subset L_{i-1}$
 2. $|L_i| \approx |L_{i-1}|/2$
 3. Searches use lists in decreasing order

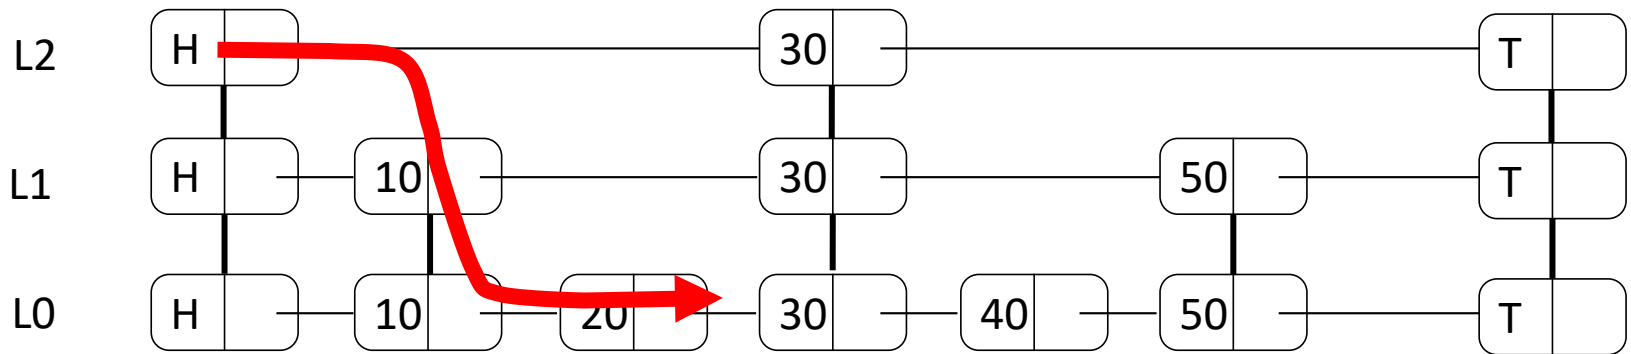
Skip list [Pugh 1990]

Search(50)



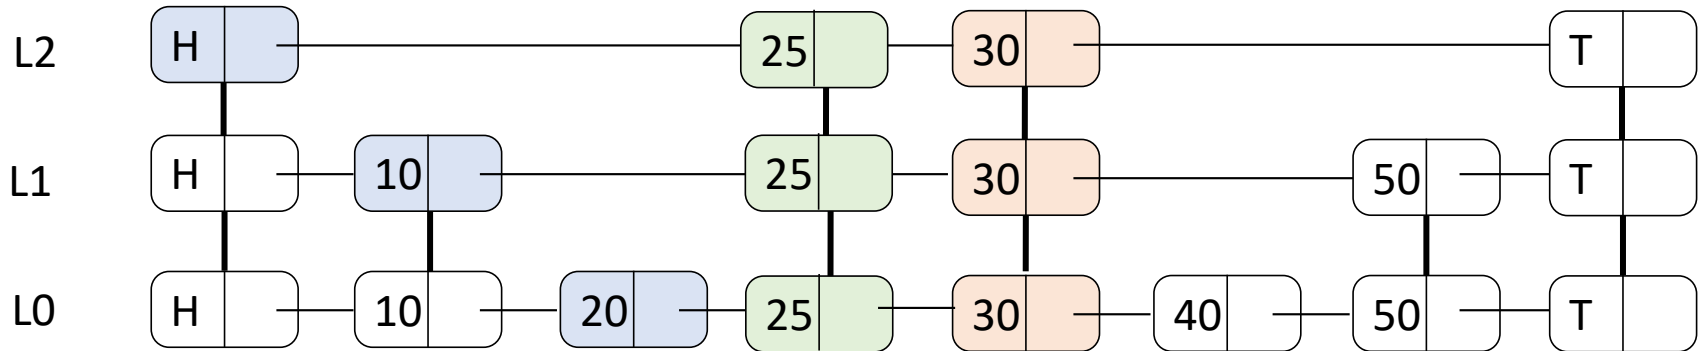
Skip list [Pugh 1990]

Insert(25)



Skip list [Pugh 1990]

Insert(25)

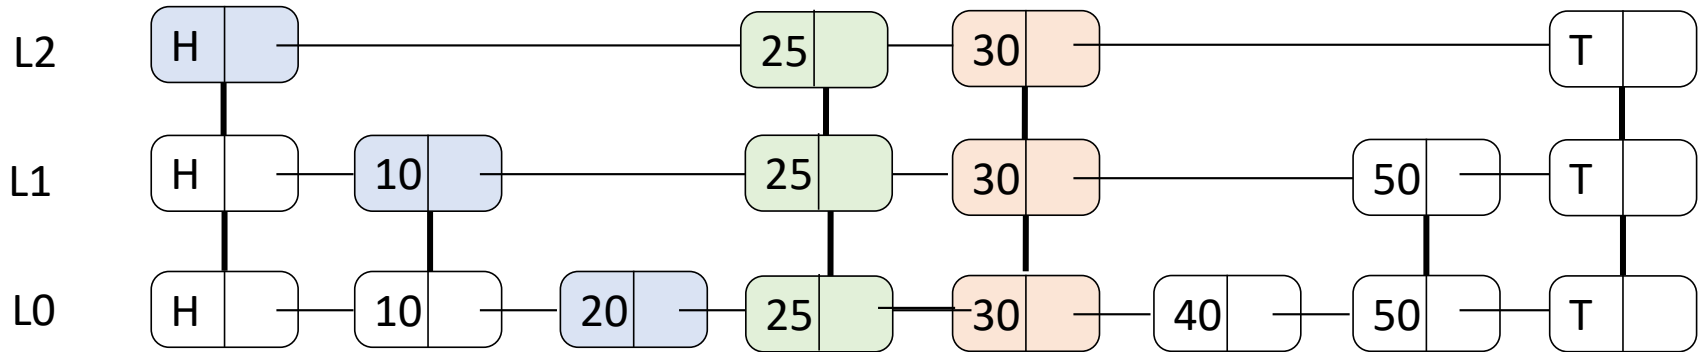


Should I insert 25 at L1? Flip a coin!

Should I insert 25 at L2? Flip a coin!

Skip list [Pugh 1990]

Delete(25)



Skip list [Pugh 1990]

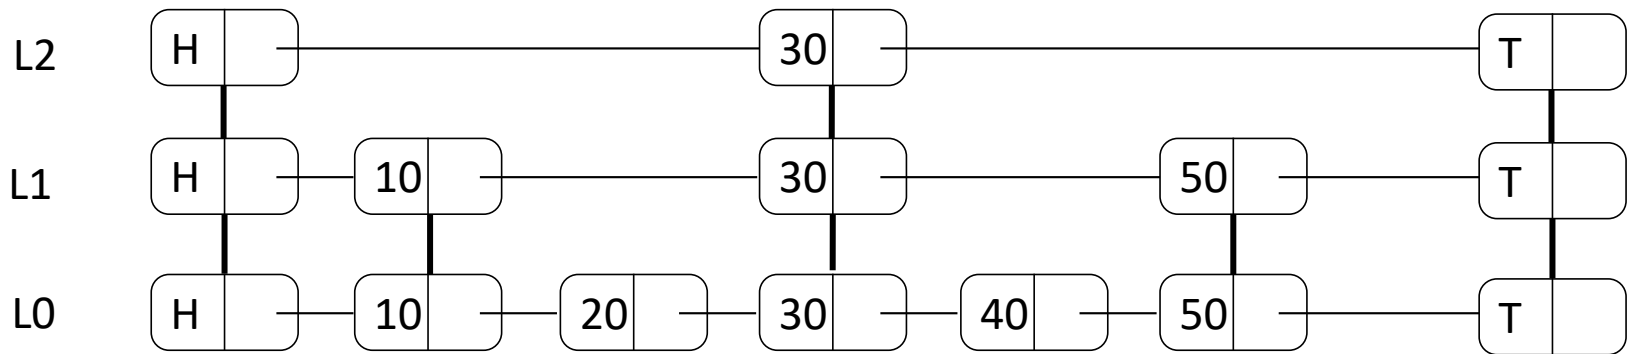
- How many (expected) keys for each level?
- $L_0 = N$
- $L_1 = N/2$
- $L_2 = N/4$
- ...
- $L(\log N) = 1$

Skip list [Pugh 1990]

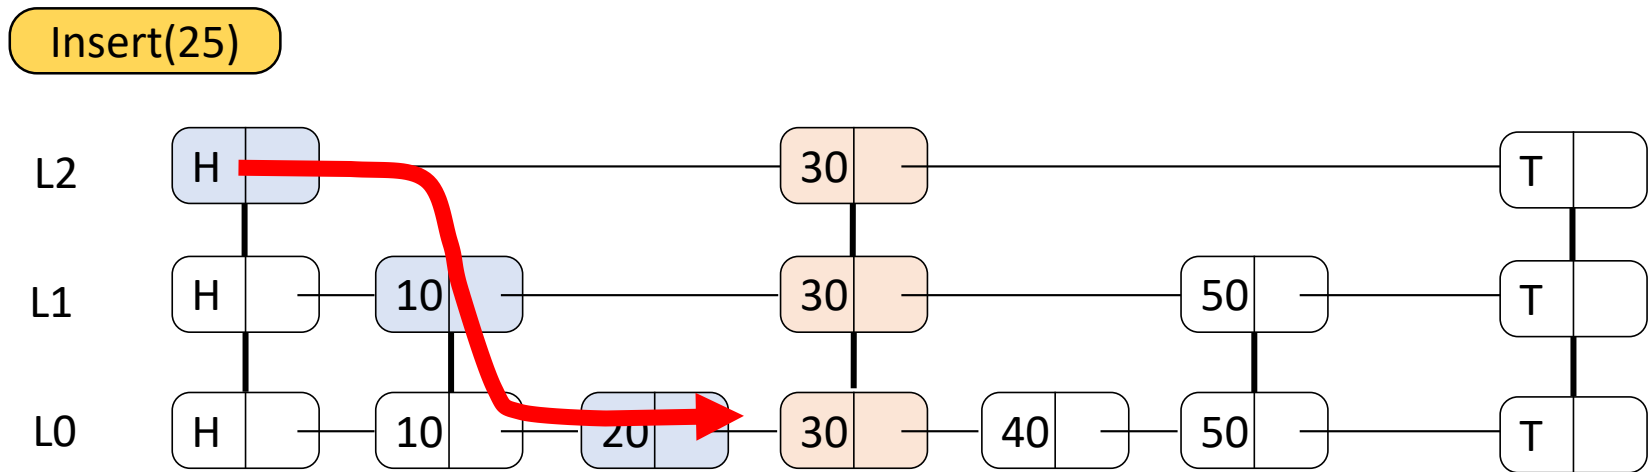
- How many steps per level?

Non-blocking Skip list [Fraser2004]

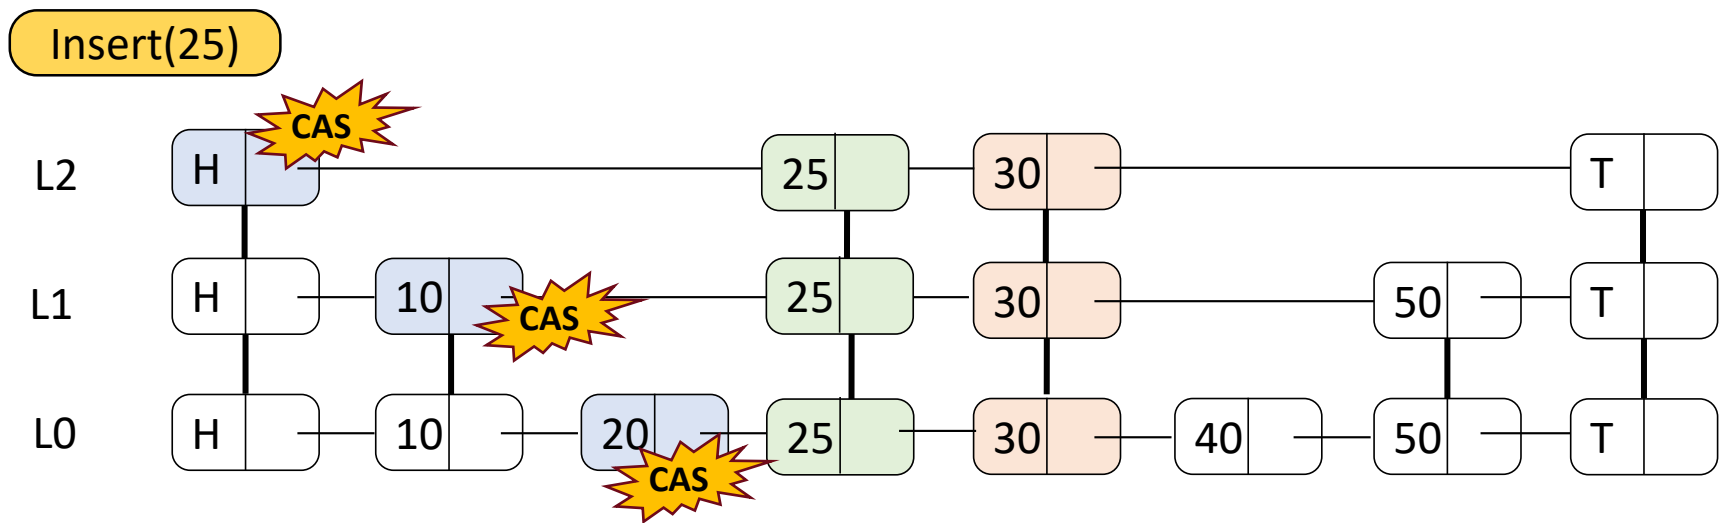
Insert(25)



Non-blocking Skip list [Fraser2004]




Non-blocking Skip list [Fraser2004]

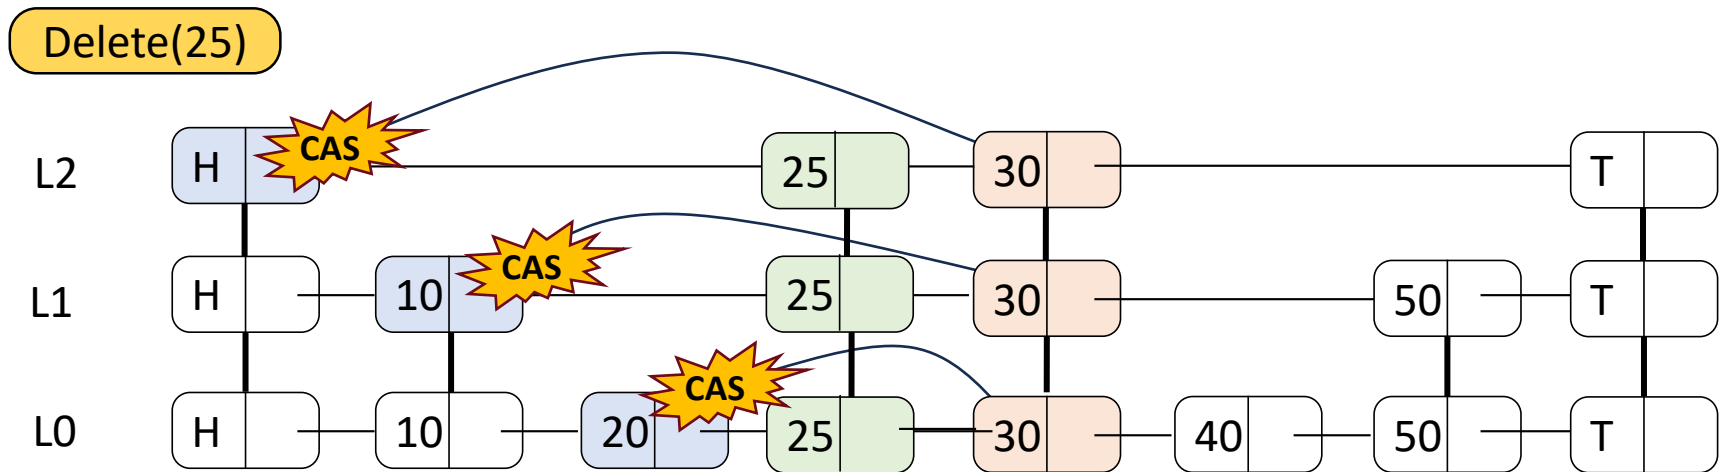


Non-blocking Skip list [Fraser2004]

Delete(25)



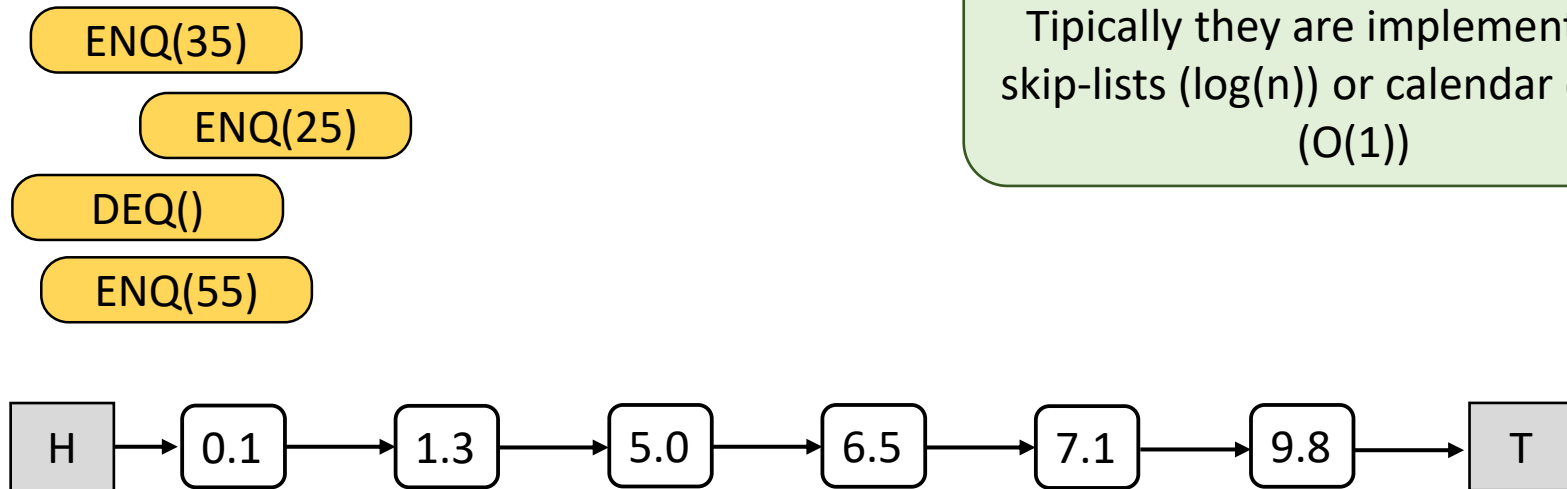
Non-blocking Skip list [Fraser2004]



Concurrent Data Structures: **Priority queues**

Priority queue implementations


- Priority Queue methods:
 - `enqueue(k)`: adds a new item
 - `dequeue()`: returns and remove the highest priority item
- Implemented as an ordered linked list




This is a huge simplification.
Typically they are implemented as
skip-lists ($\log(n)$) or calendar queues
($O(1)$)

Priority queue – Attempt 1

- Enqueue: works as insertions in the non-blocking Set
 - Connect via CAS
- Dequeues: work as deletions in the non-blocking Set
 - Mark as logically deleted, but
 - DISCONNECT IMMEDIATELY
- Is it scalable?

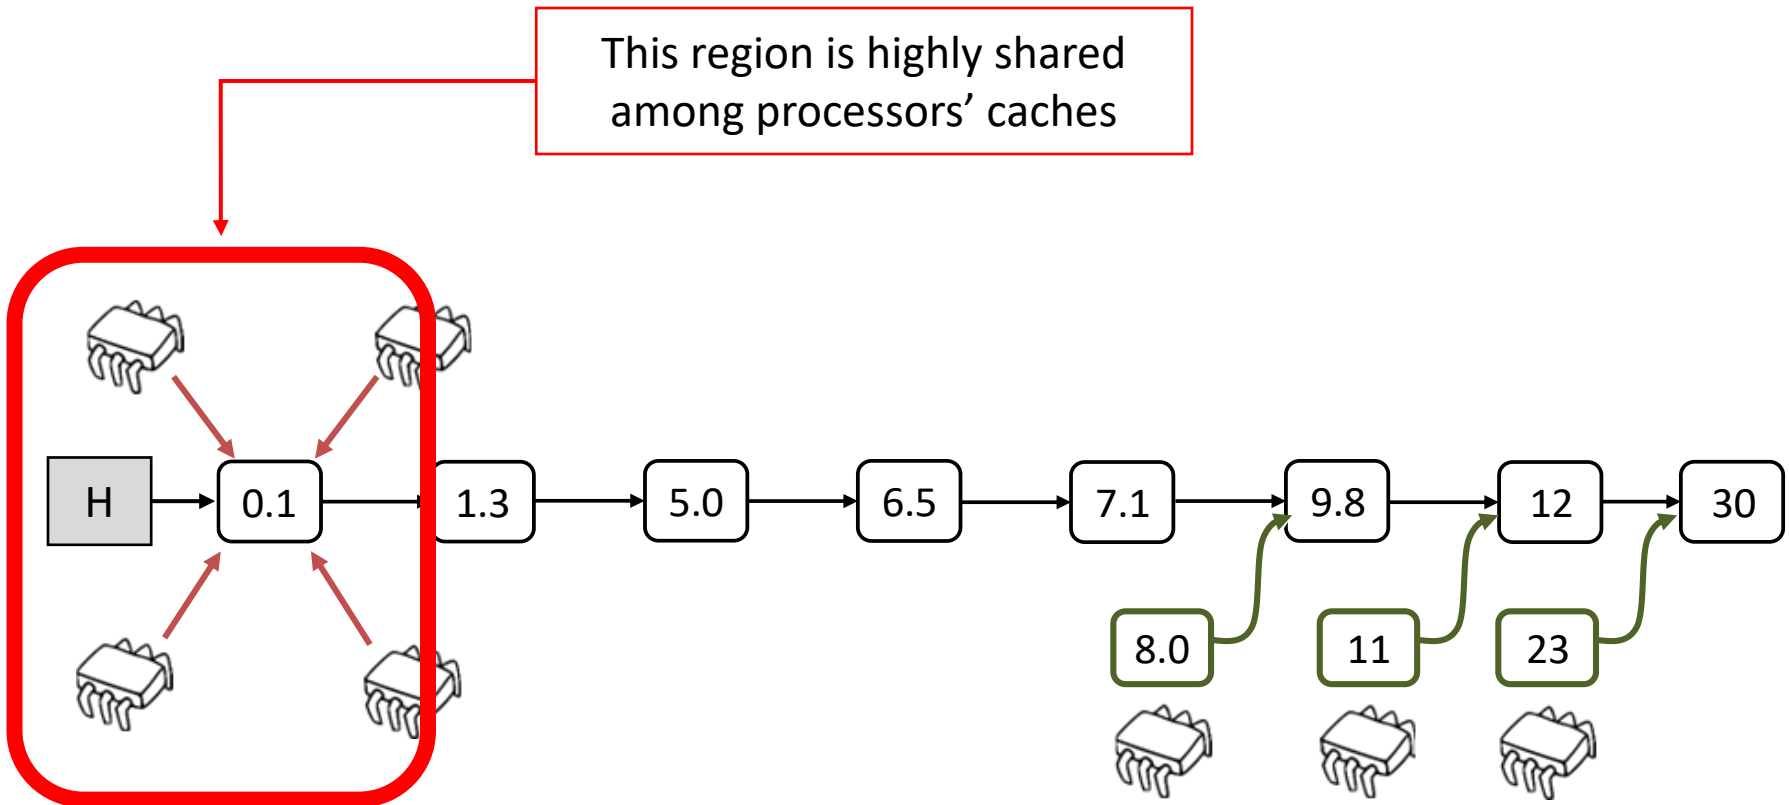


Priority queue – Attempt 1





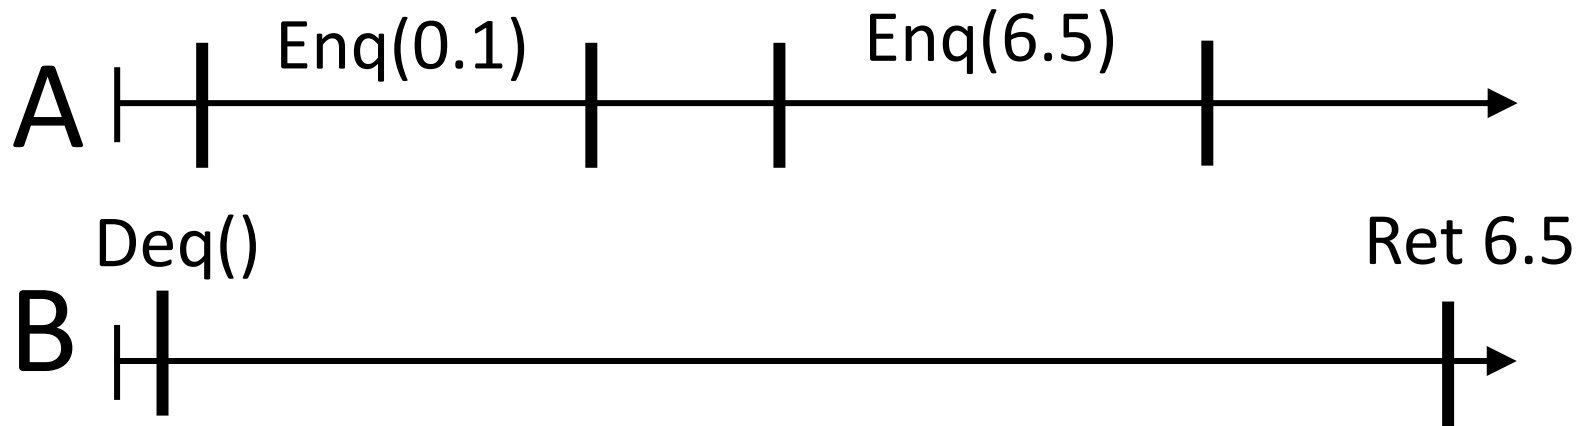
Priority queues: an inherently “sequential” semantic

- Enqueue offers a high level of disjoint access parallelism
- Dequeues are prone to conflicts



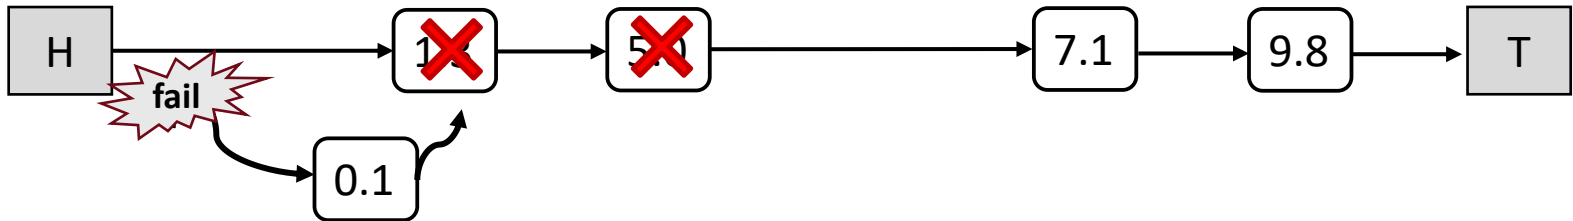
Lazy deletion within priority queues

- If we use lazy deletion “as is”, we might obtain non-linearizable extractions



Correct lazy deletion within priority queues

- To implement correct extractions with lazy deletions there are two main approaches
 1. Move the logical mark of a node in the field “next” of its predecessor

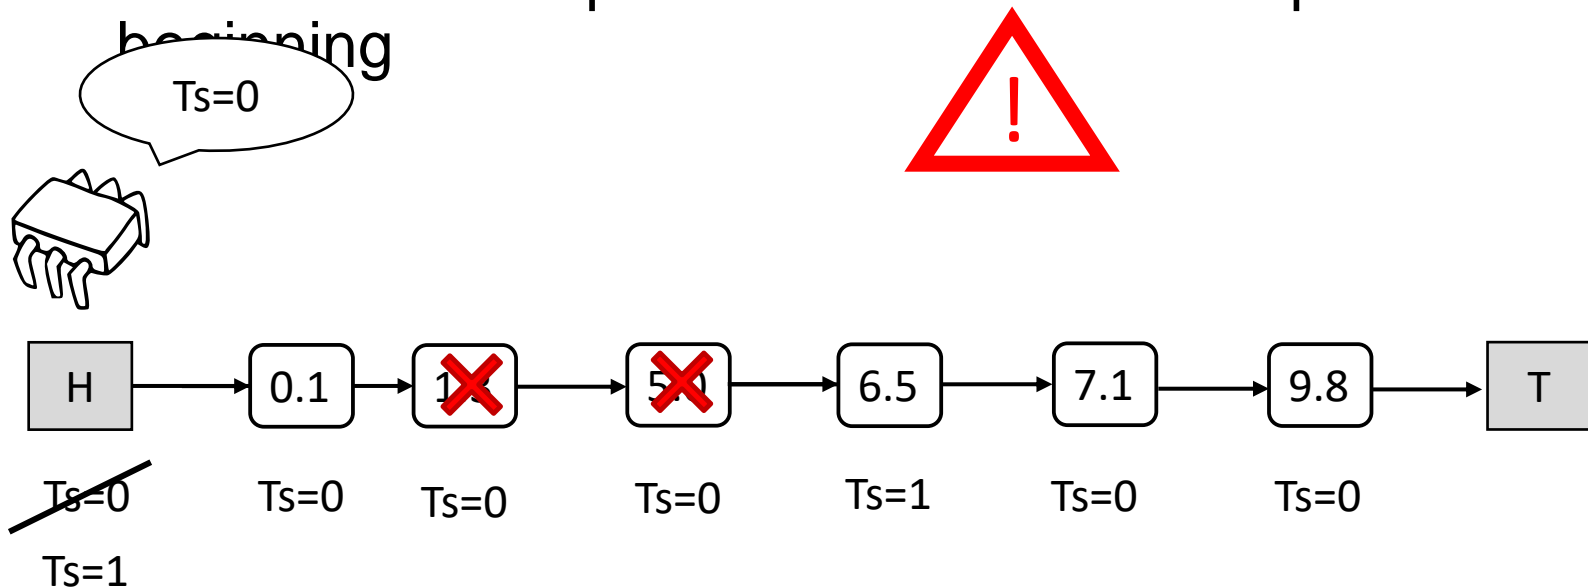


Correct lazy deletion within priority queues

- To implement correct extractions with lazy deletions there are two main approaches

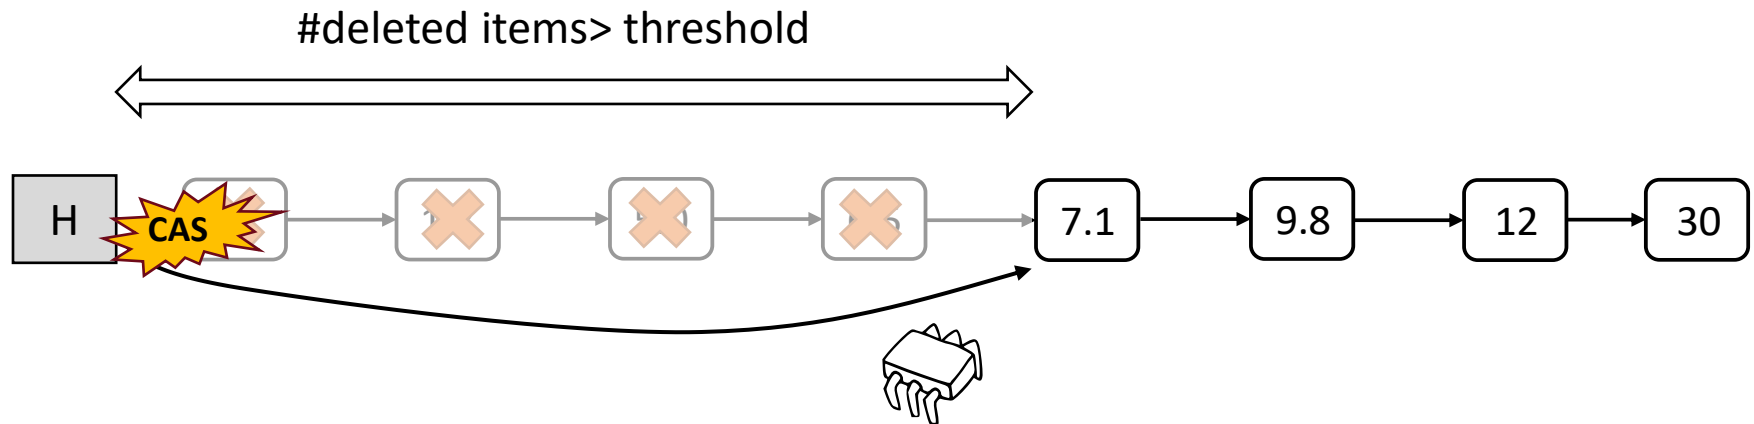
2. Use logical timestamps:

- incremented each time a new minimum has been inserted
- extract item compatible with the timestamp read at the beginning

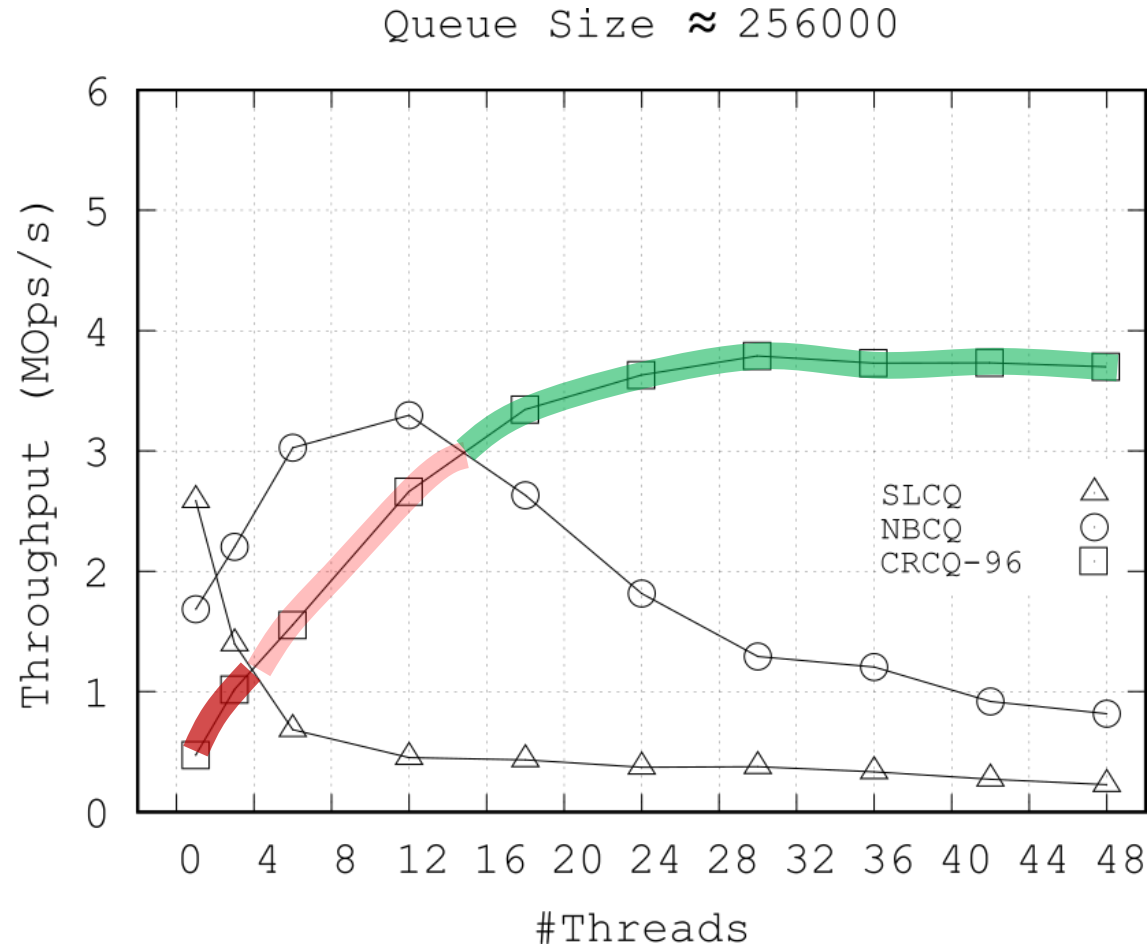


PQ – Attempt 2 - Introducing Conflict Resiliency

- Lazy deletion
- Skip logically deleted items \Rightarrow IT INCREASES THE NUMBER OF STEPS \Rightarrow EXPENSIVE IN TERMS OF IMPACT ON CACHE
- Periodic Housekeeping

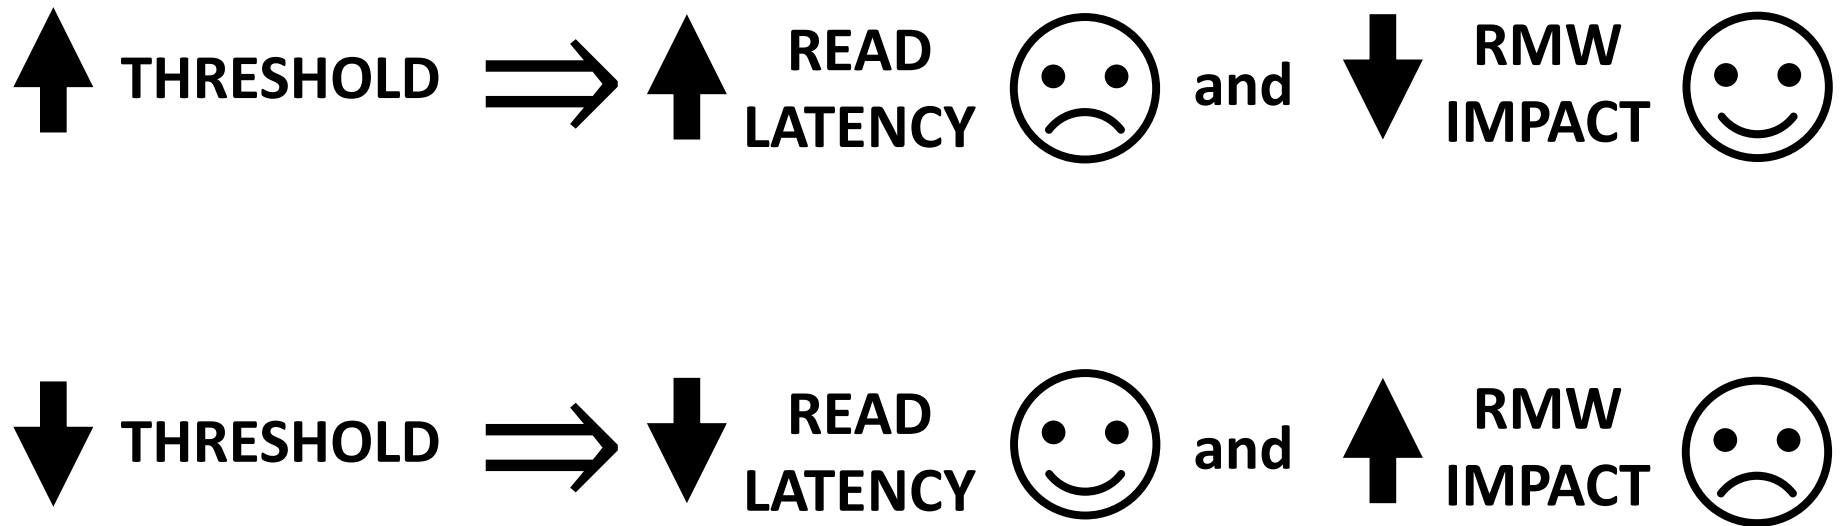


Priority queue – Attempt 2

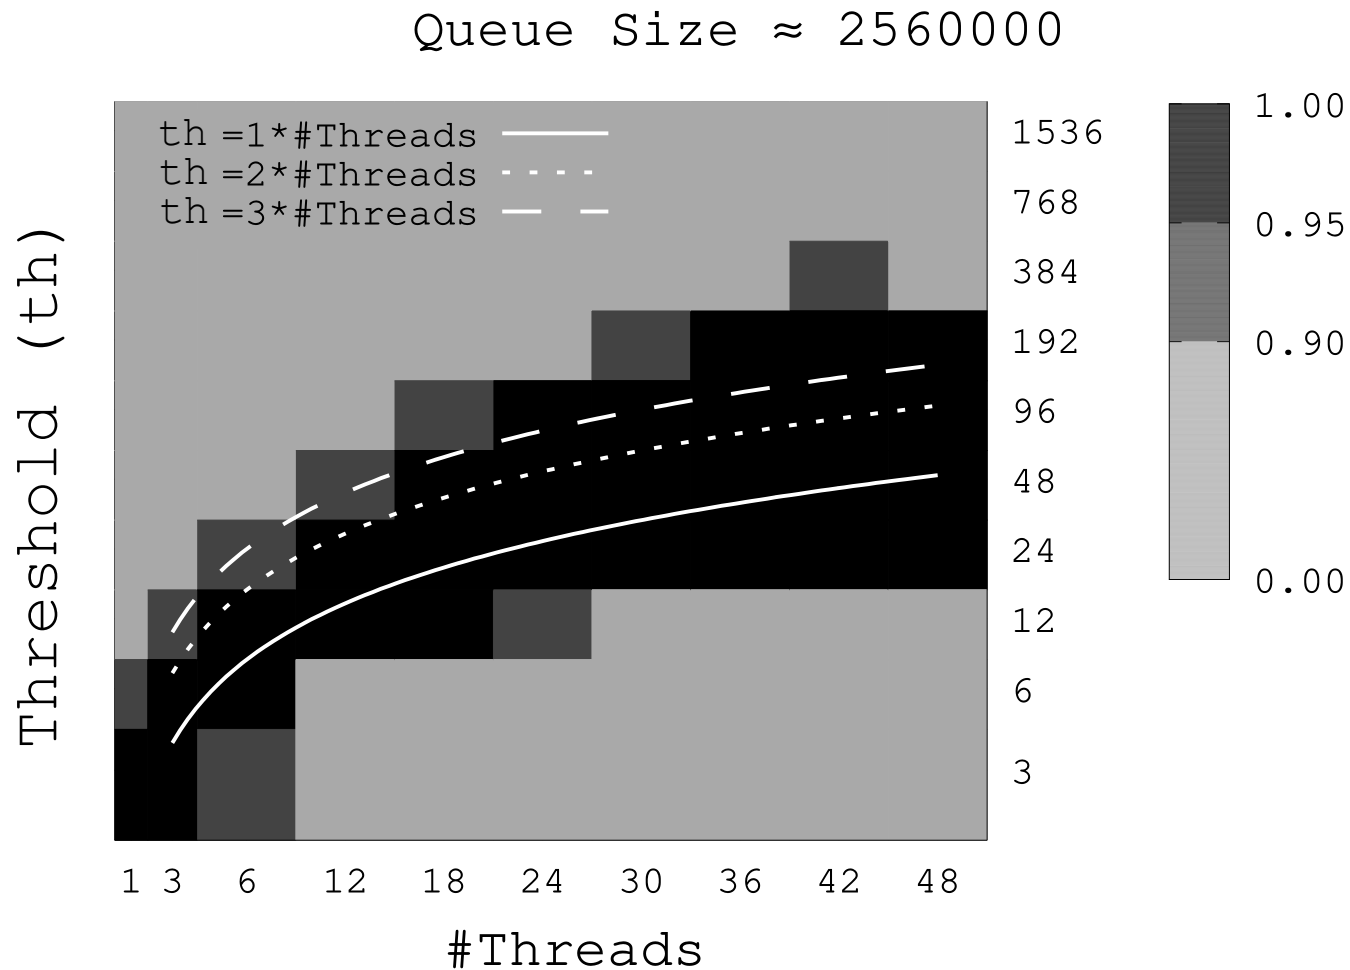


On the conflict resiliency trade off

- The number of steps per dequeue and costs of housekeeping are dependent:



Conflict resiliency trade offs



Priority queues – Attempt 3

