

Programmazione concorrente

Laurea Magistrale in Ingegneria Informatica

Università Tor Vergata

Docente: Romolo Marotta

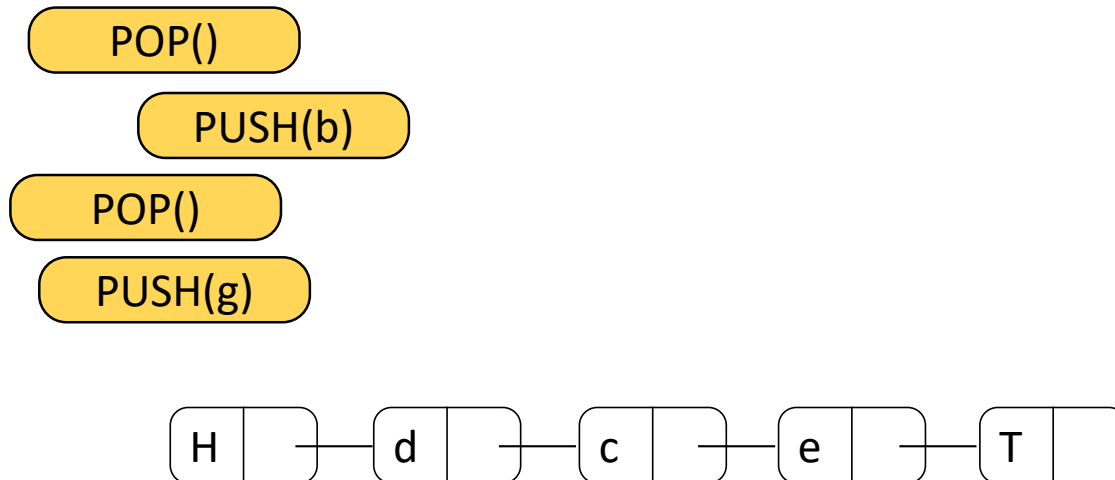
Concurrent data structures

1. Stack

Concurrent Data Structures: **Stacks**

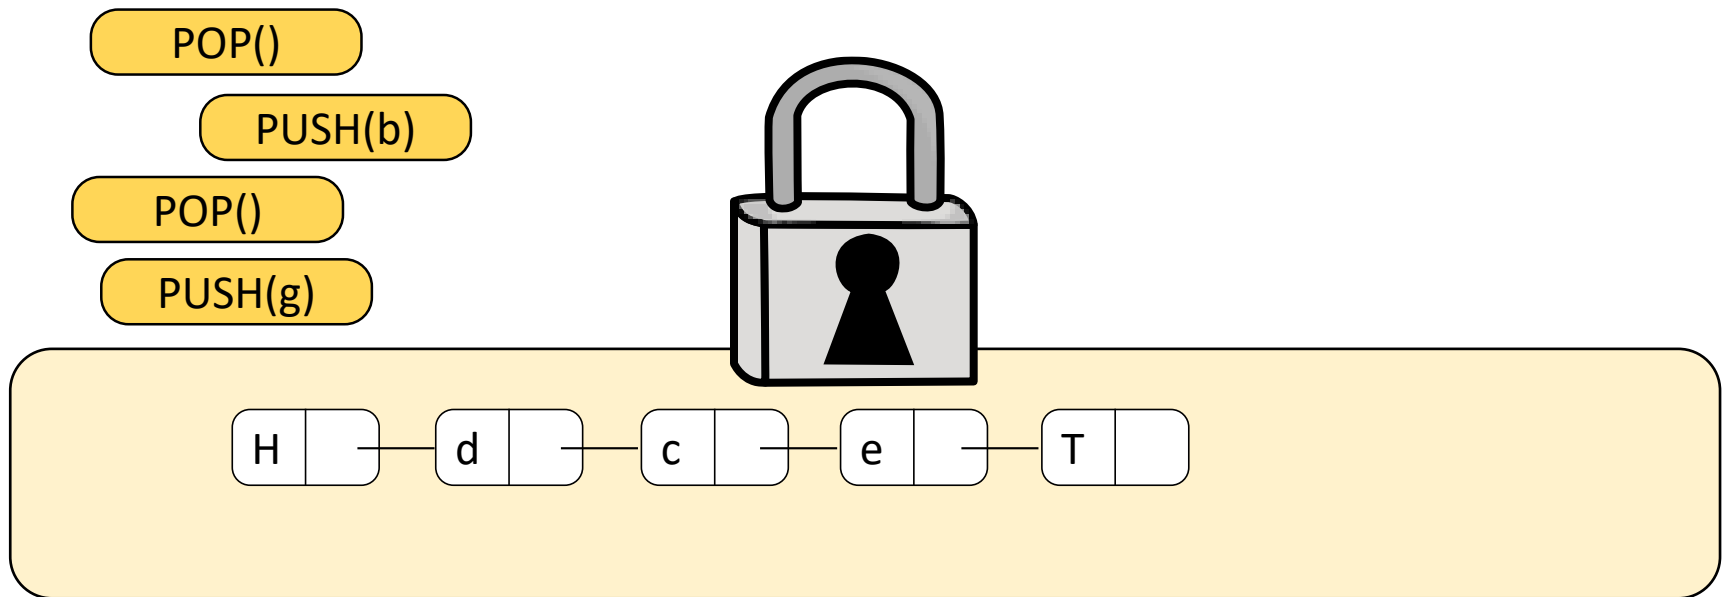
Stack implementation

- Stack methods:
 - `push(v)`
 - `pop()`
- Implemented as a linked list



Concurrent stack implementations

- Resort to a global lock



Read-Modify-Write

- RMW instructions allow to read memory and modify its content in an apparently instantaneous fashion.

```
1. RMW(MRegister *r, Function f){  
2.   atomic{  
3.     old = r;  
4.     *r = f(r);  
5.     return old;  
6.   }  
7. }
```

- Even conventional atomic Load and Store can be seen as RMW operations

Compare-And-Swap

- Compare-and-Swap (CAS) is an atomic instruction used in multithreading to achieve synchronization
 - It compares the contents of a memory area with a supplied value
 - If and only if they are the same
 - The contents of the memory area are updated with the new provided value
- Atomicity guarantees that the new value is computed based on up-to-date information
- If, in the meanwhile, the value has been updated by another thread, the update fails
- This instruction has been introduced in 1970 in the IBM 370 trying to limit as much as possible the use of spinlocks



Attempt 1

Push:

1. Get head next
2. Insert the new item with a CAS
3. If CAS fails, restart


Delete:

1. Get head next
2. Disconnect the item with a CAS
3. If CAS fails, restart



- Is it scalable?

Non-blocking stack – Attempt 2 [Treiber+BO]





Non-blocking stack – Attempt 2 [Treiber+BO]

Push:

1. Get head next
2. Insert the new item with a CAS
3. If CAS fails, ~~restart~~ backoff and restart


Delete:

1. Get head next
2. Disconnect the item with a CAS
3. If CAS fails, ~~restart~~ backoff and restart



- Is it scalable?

Non-blocking stack – Attempt 2 [Treiber+BO]




Concurrent stack implementations

- Resort to a global lock
 - Do not scale
- Resort to a naïve non-blocking approach
 - Do not scale
- Resort to a naïve non-blocking approach + Back off
 - Do not scale, but conflict resilient
- How achieve scalability? Make back-off times useful


POP()

PUSH(g)




Non-blocking stack – Attempt 3

- How to take advantage of back-off times?



Observation


- Concurrent matching push/pop pairs are always linearizable



- A push A and a pop B are:
 - concurrent to each other
 - B returns the item inserted by A
- ⇒ we can always take two points such that:
- A is the last one to insert an item before A linearizes
 - B appears to extract the last item inserted (by A)


Non-blocking stack – Attempt 3

- How to take advantage of back-off times?
- Hope that an opposite operation arrives while waiting
- Match the two without interacting with the stack




Non-blocking stack – Attempt 3

- How to take advantage of back-off times?
- Hope that an opposite operation arrives while waiting
- Match the two without interacting with the stack
- How??




Non-blocking stack – Elimination stack

- Pair the Treiber stack with an array
- Algorithm:
 1. Update the original stack via CAS
 2. If CAS fails, publish the operation in a random cell of the array



Non-blocking stack – Elimination stack

- Pair the Treiber stack with an array
- Algorithm:
 1. Update the original stack via CAS
 2. If CAS fails, publish the operation in a random cell of the array
 3. Wait for a matching operation
 4. If no matching op, GOTO 1



Non-blocking stack – Attempt 3

