

# Programmazione concorrente

Laurea Magistrale in Ingegneria Informatica


Università Tor Vergata

Docente: Romolo Marotta

## Locks

1. Spin locks
2. FIFO Spin locks
3. Hybrid (Spin+Sleep) locks

# Blocking coordination

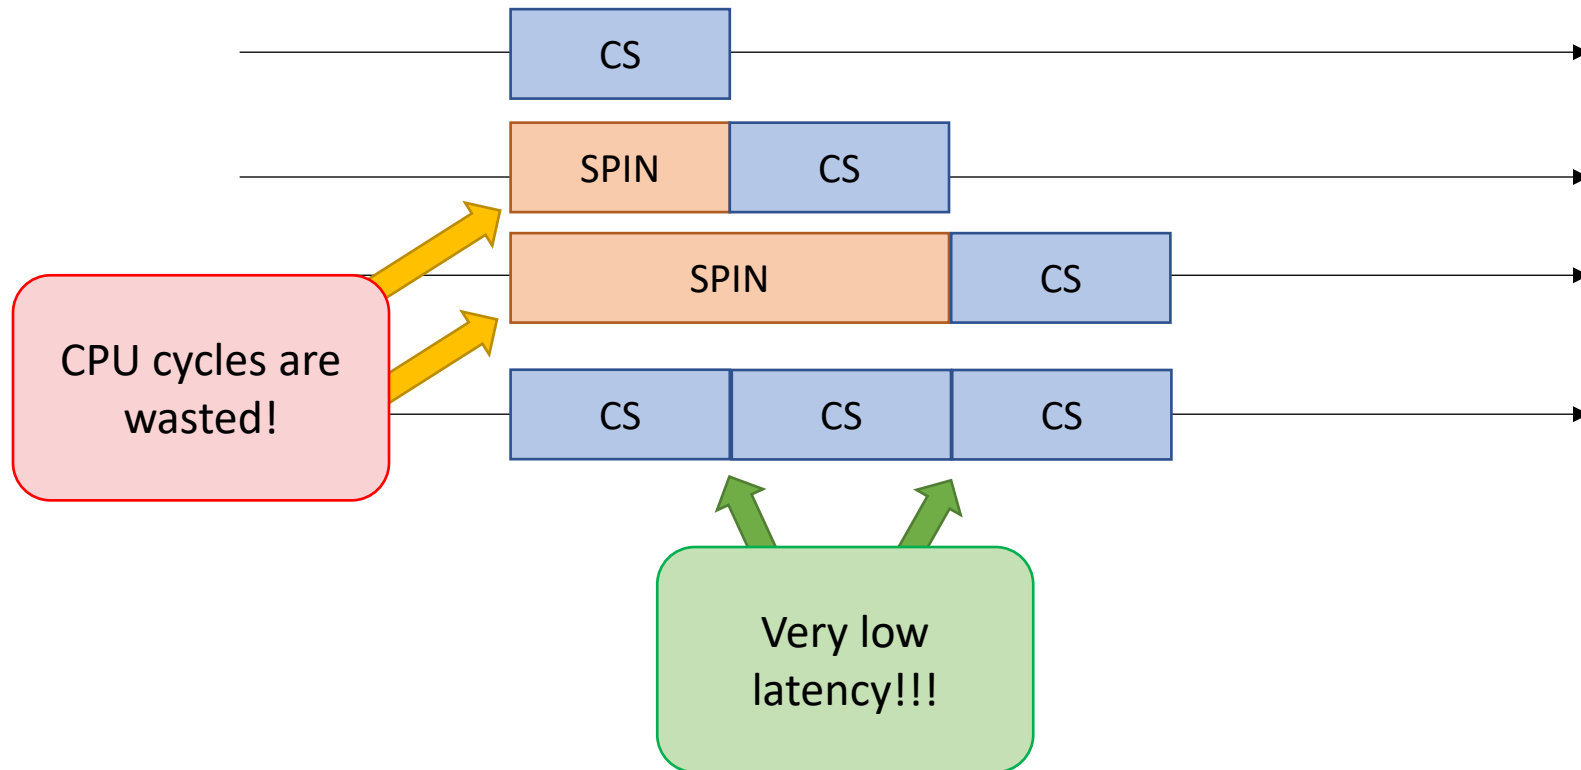


Choosing between them is delegated to developers!

This is a very hard task:  
multiple trade offs!

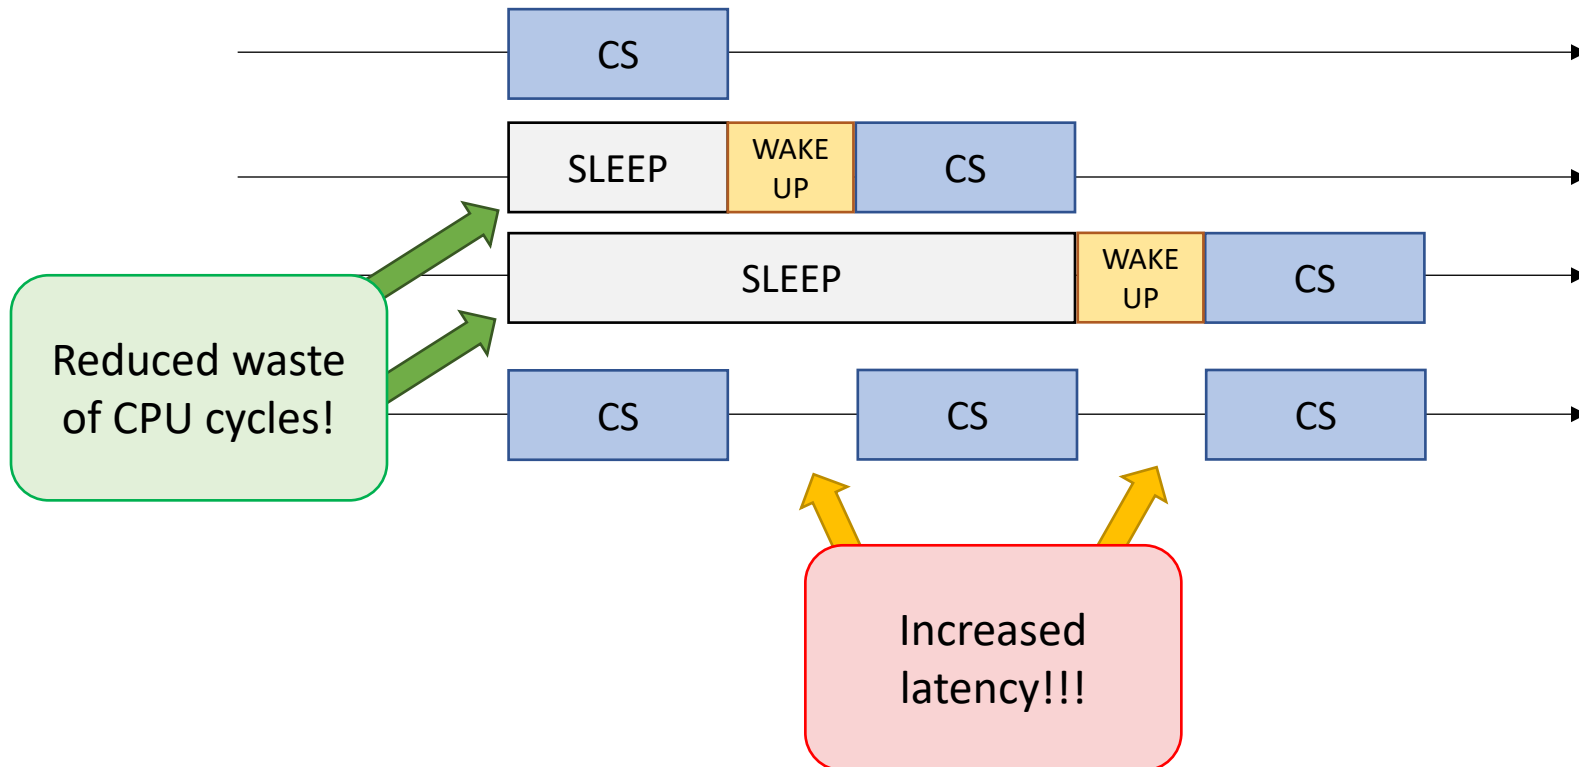
# Spinning vs Sleeping

Benefits	Spinning
Guaranteed low latency	✓
Computing power savings	✗



# Spinning vs Sleeping

Benefits	Waiting Policy	
	Spinning	Sleeping
Guaranteed low latency	✓	✗
Computing power savings	✗	✓
Autonomic Adaptivity	✗	✗



# Spin vs Sleep – is that all?

- Choosing the proper back off scheme is very challenging
- Even implementing a simple spin lock is not trivial
  - Trade off between low and high contention case
  - You should have heard about algorithms for Mutual Exclusion in Distributed Systems lectures
    - E.g. Dijkstra, Bakery algorithm, Peterson...
  - Those algorithm essentially implements spin locks by resorting only on read/write operations
- Here, we will focus on spin locking algorithms that exploit stronger synchronization primitives... RMW!

# Test-and-set spin lock

- Test-and-set lock is the simplest spin lock
- Acquiring threads always try to set a variable via RMW

```
int lock = 0;
```

```
void acquire(int *lock){    void release(int *lock){  
    while(XCHG(lock, 1));    *lock = 0;  
}                             }
```

# A small benchmark

- We have an array of integers
- Each thread reverse the array

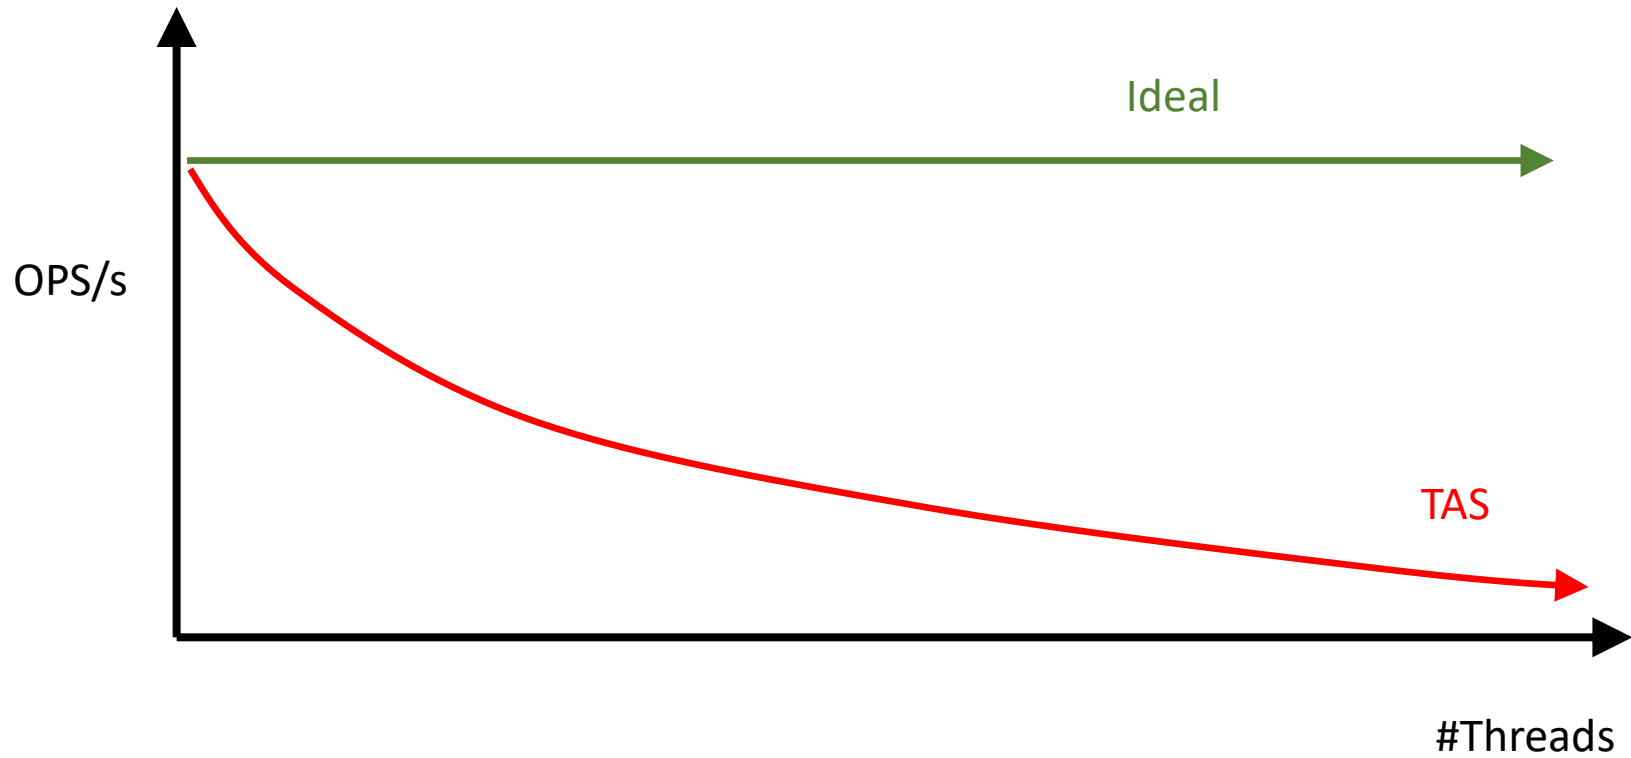


- This is done within a critical section

```
while(!stop){  
    acquire(&lock);  
    flip_array();  
    release(&lock);  
}
```

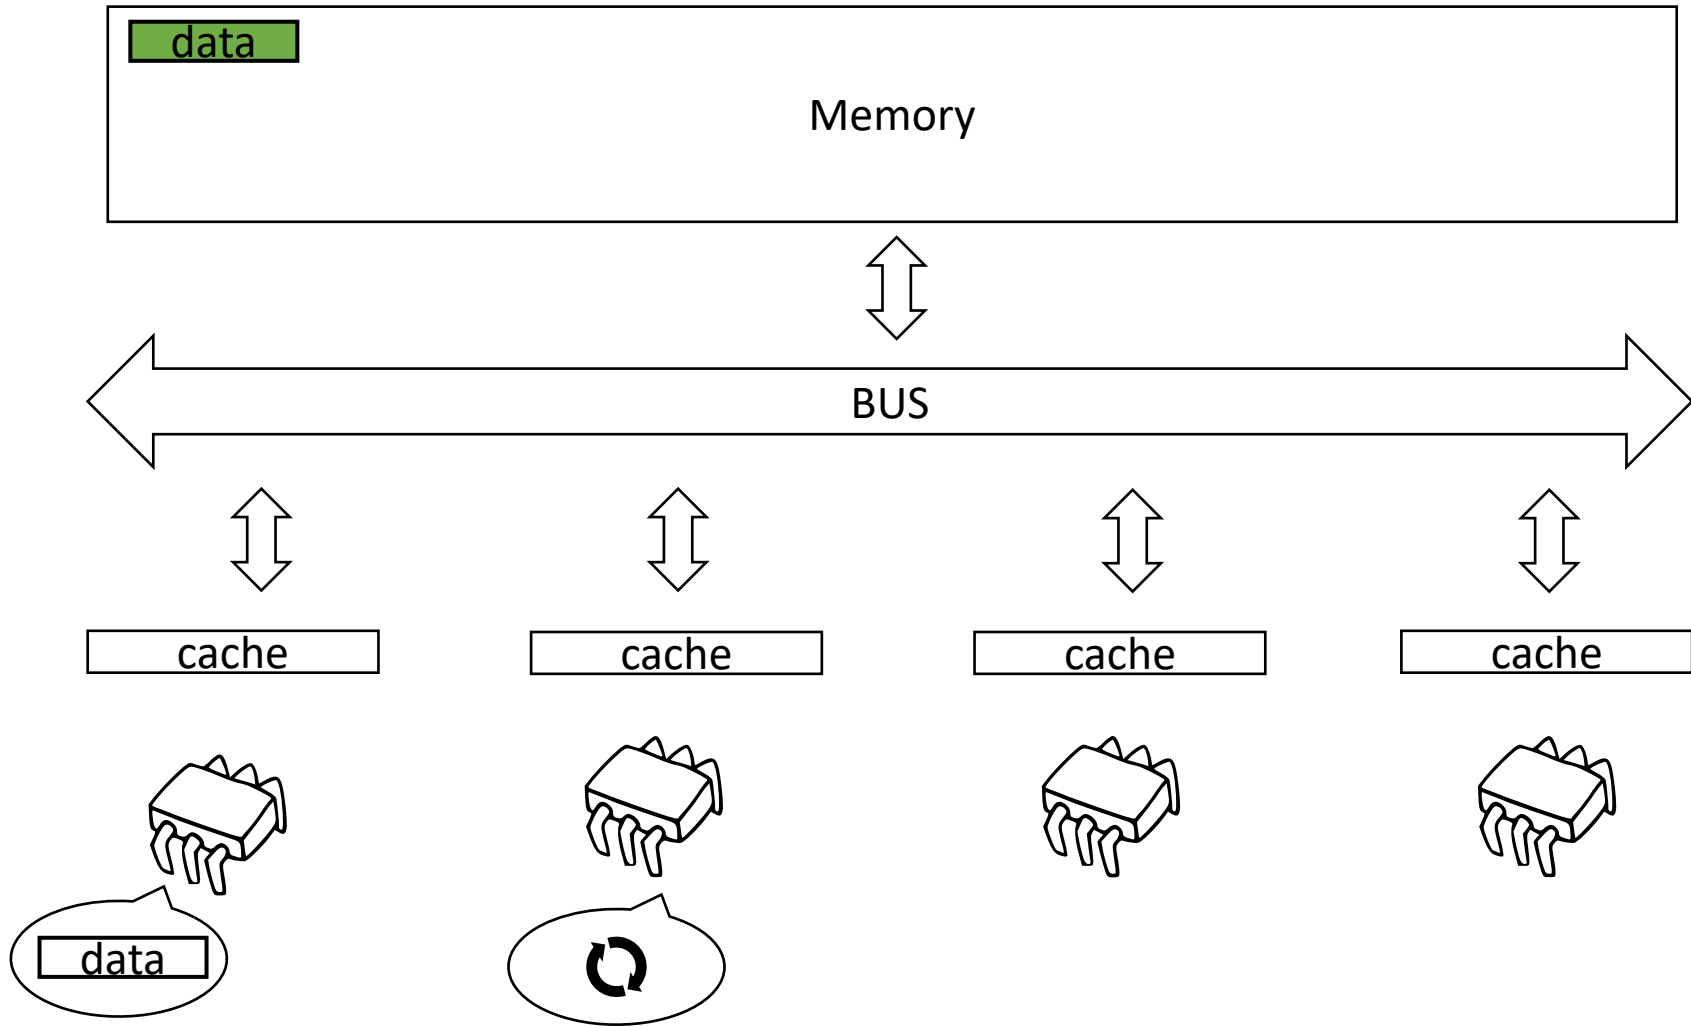
- Performance Metric:
  - Throughput = #Flips per second

# Results






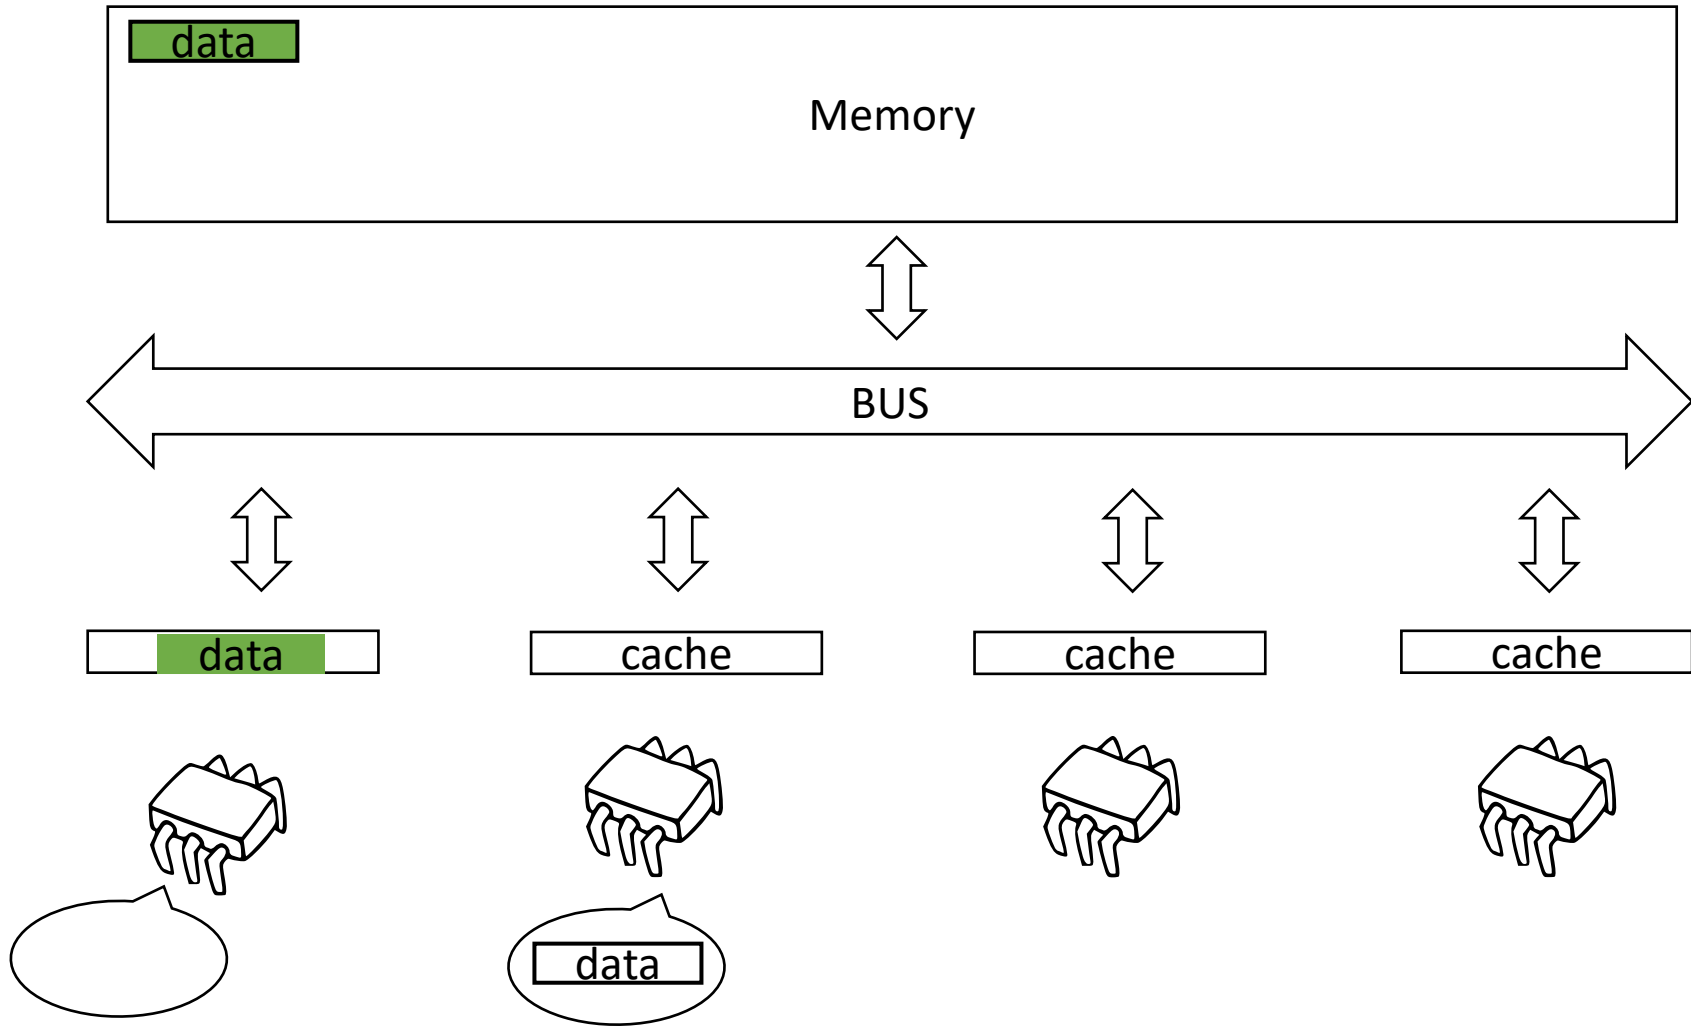
# Memory Model



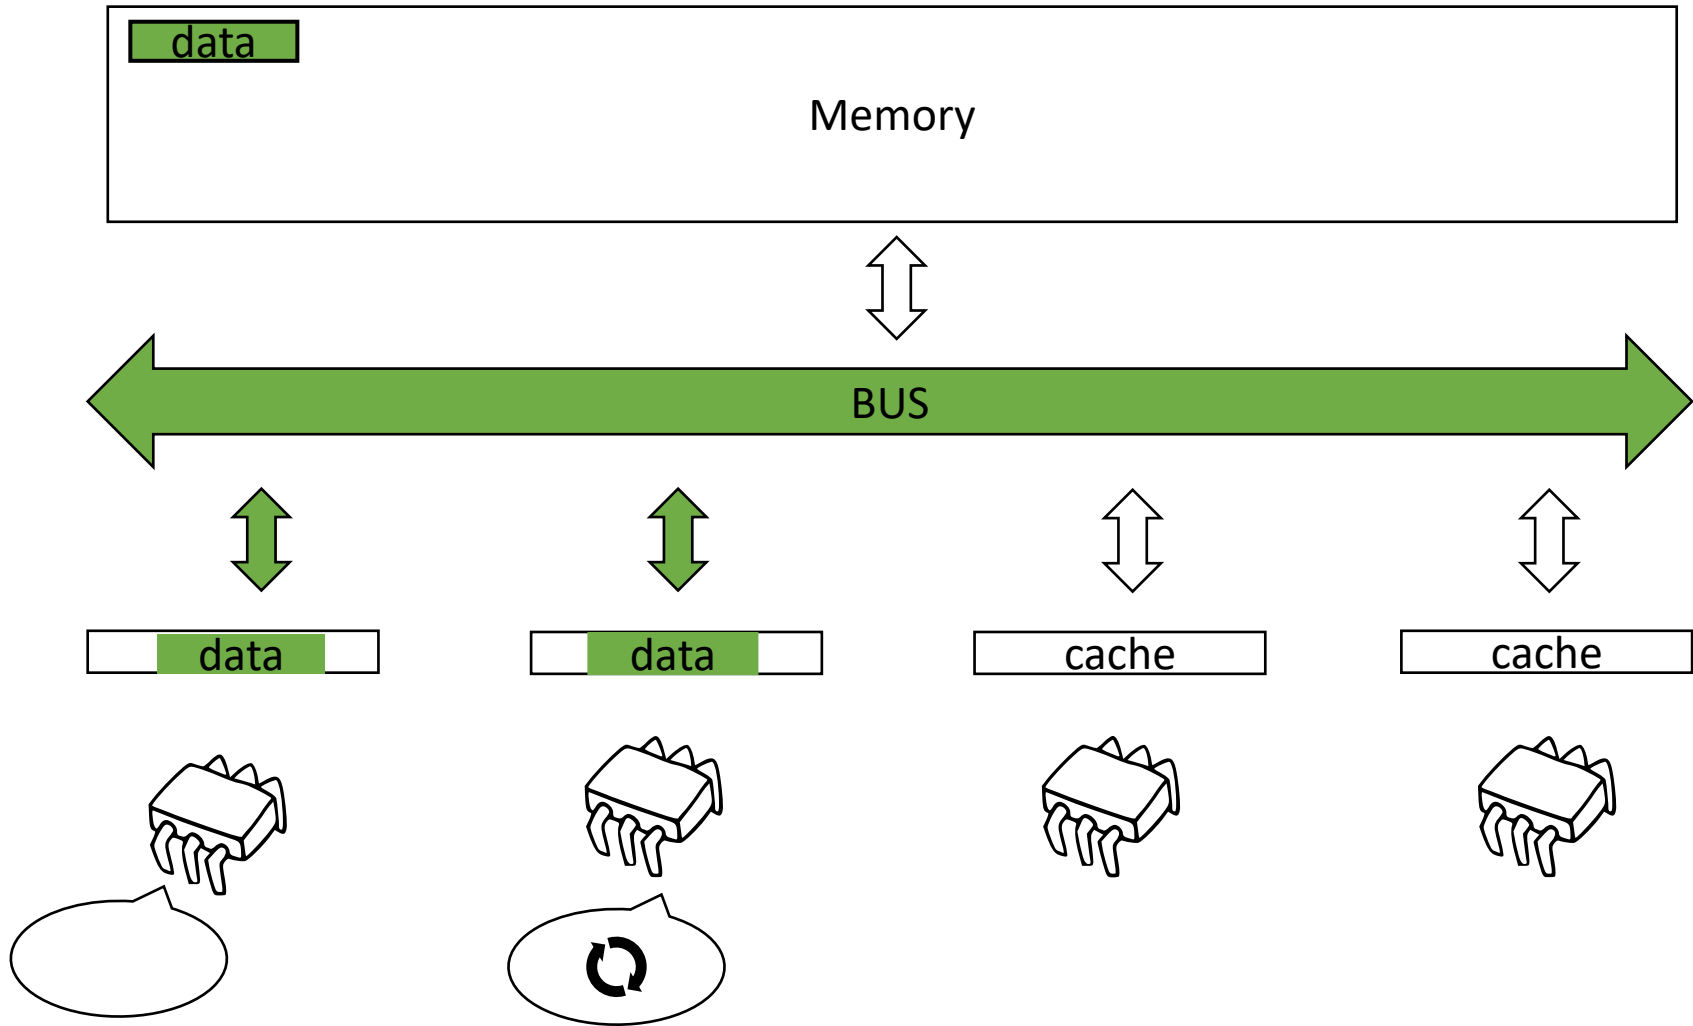
# Memory Model



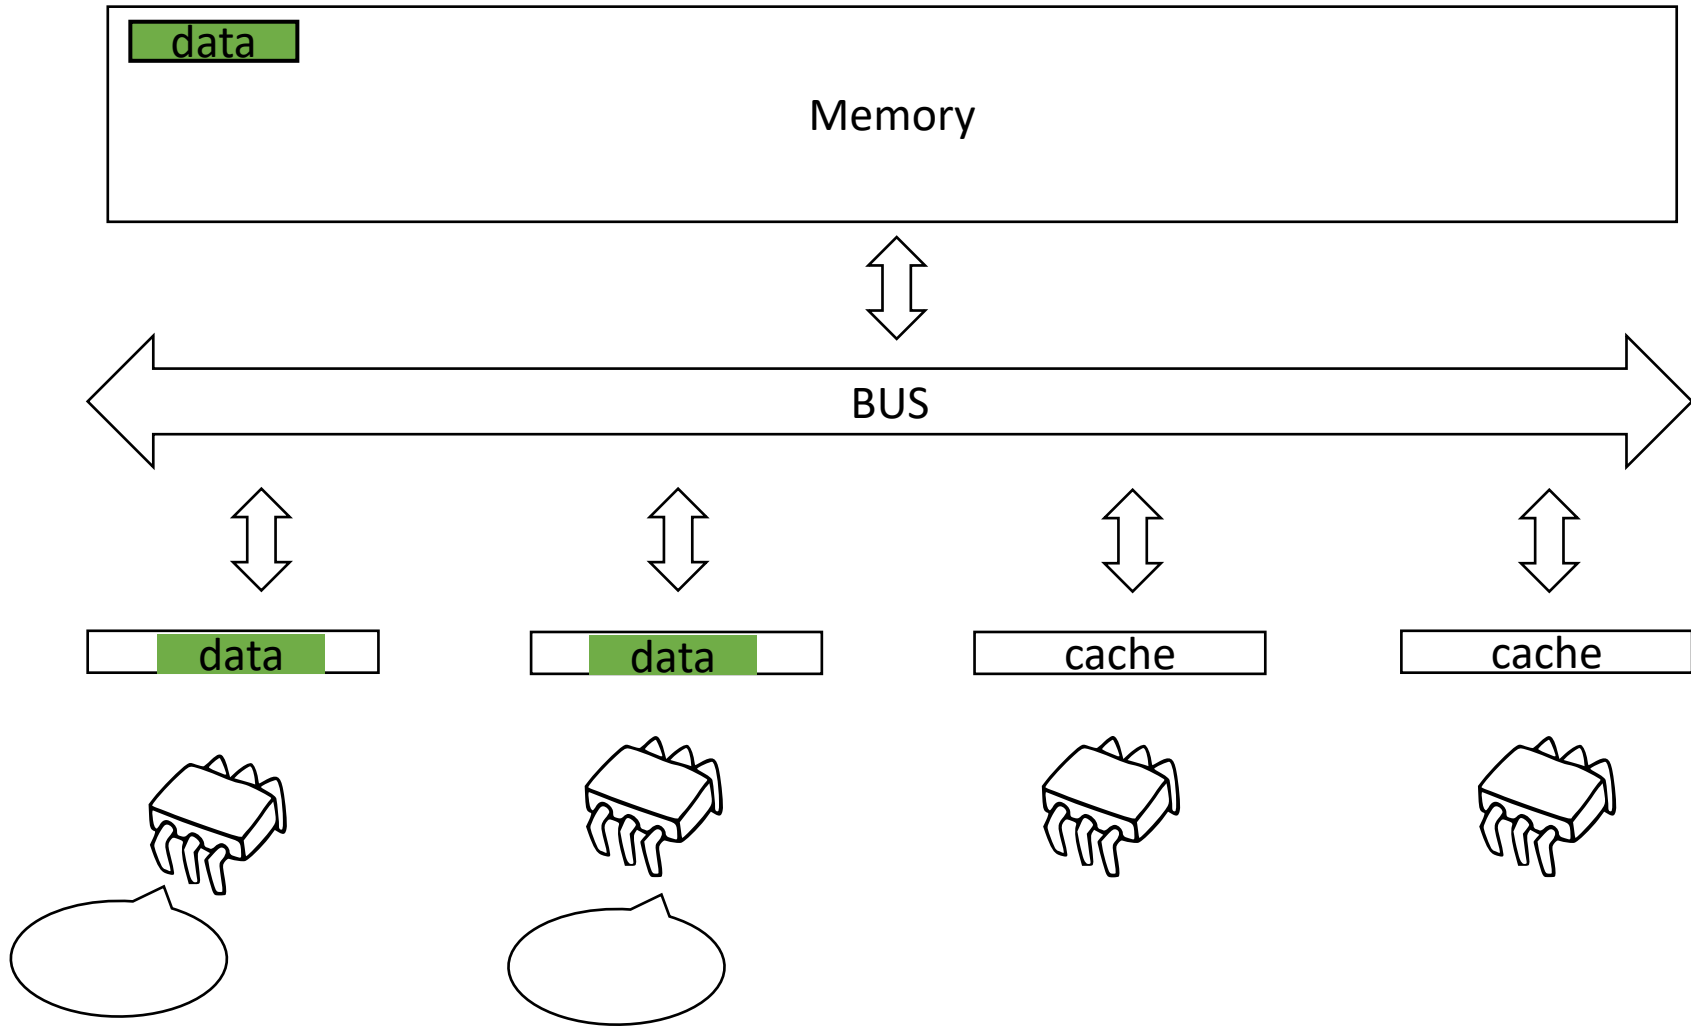
# Memory Model



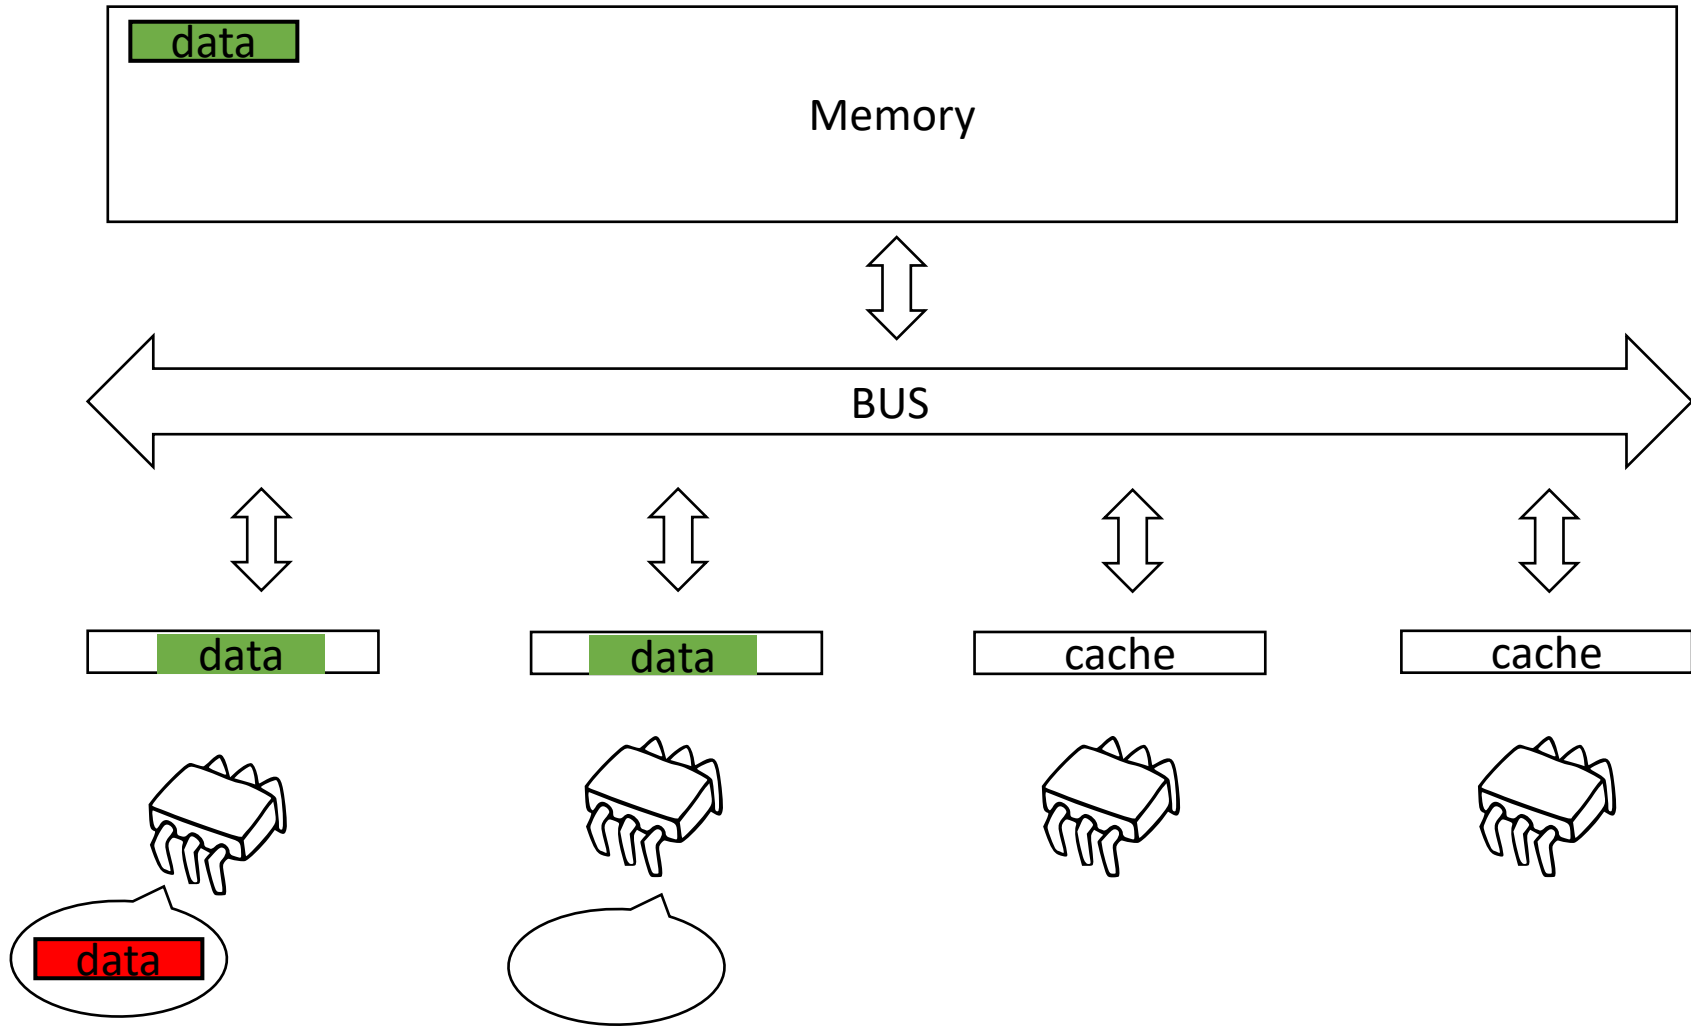
# Memory Model




# Memory Model



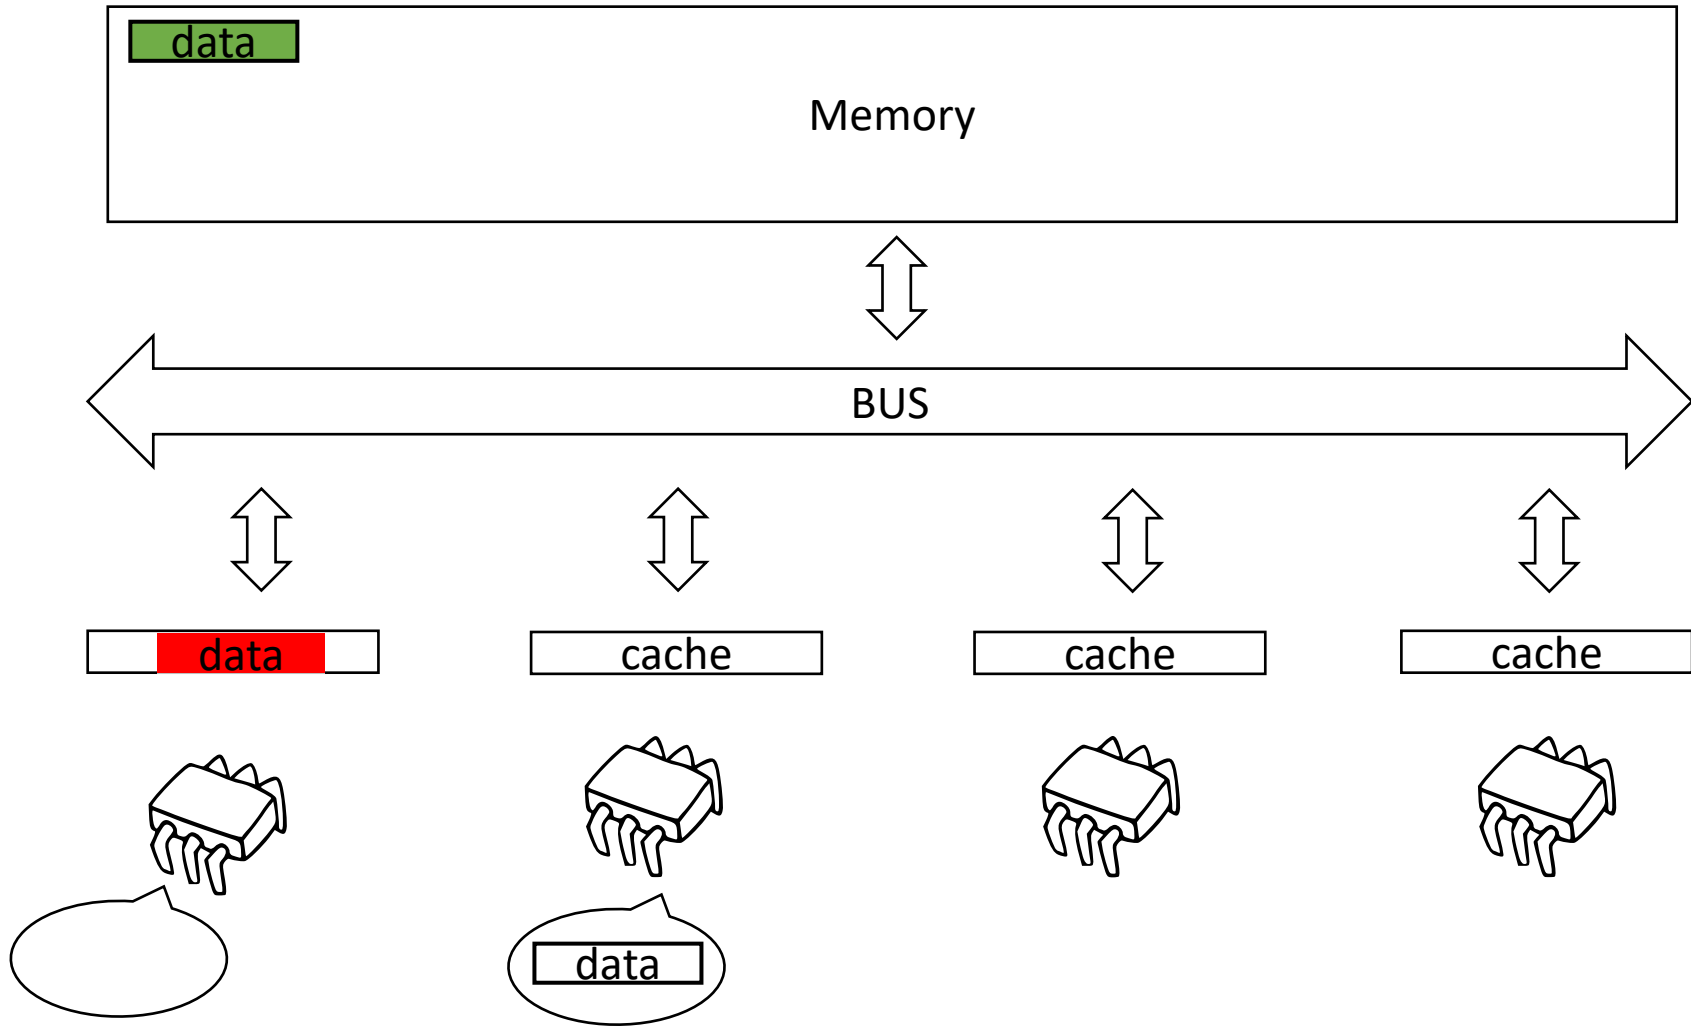
# Memory Model



# Memory Model




# Memory Model






# Memory Model



# Memory Model



# Lock implementations

# Test-and-set spin lock


- Test-and-set lock is the simplest spin lock
- Acquiring threads always try to set a variable via RMW

```
int lock = 0;
```

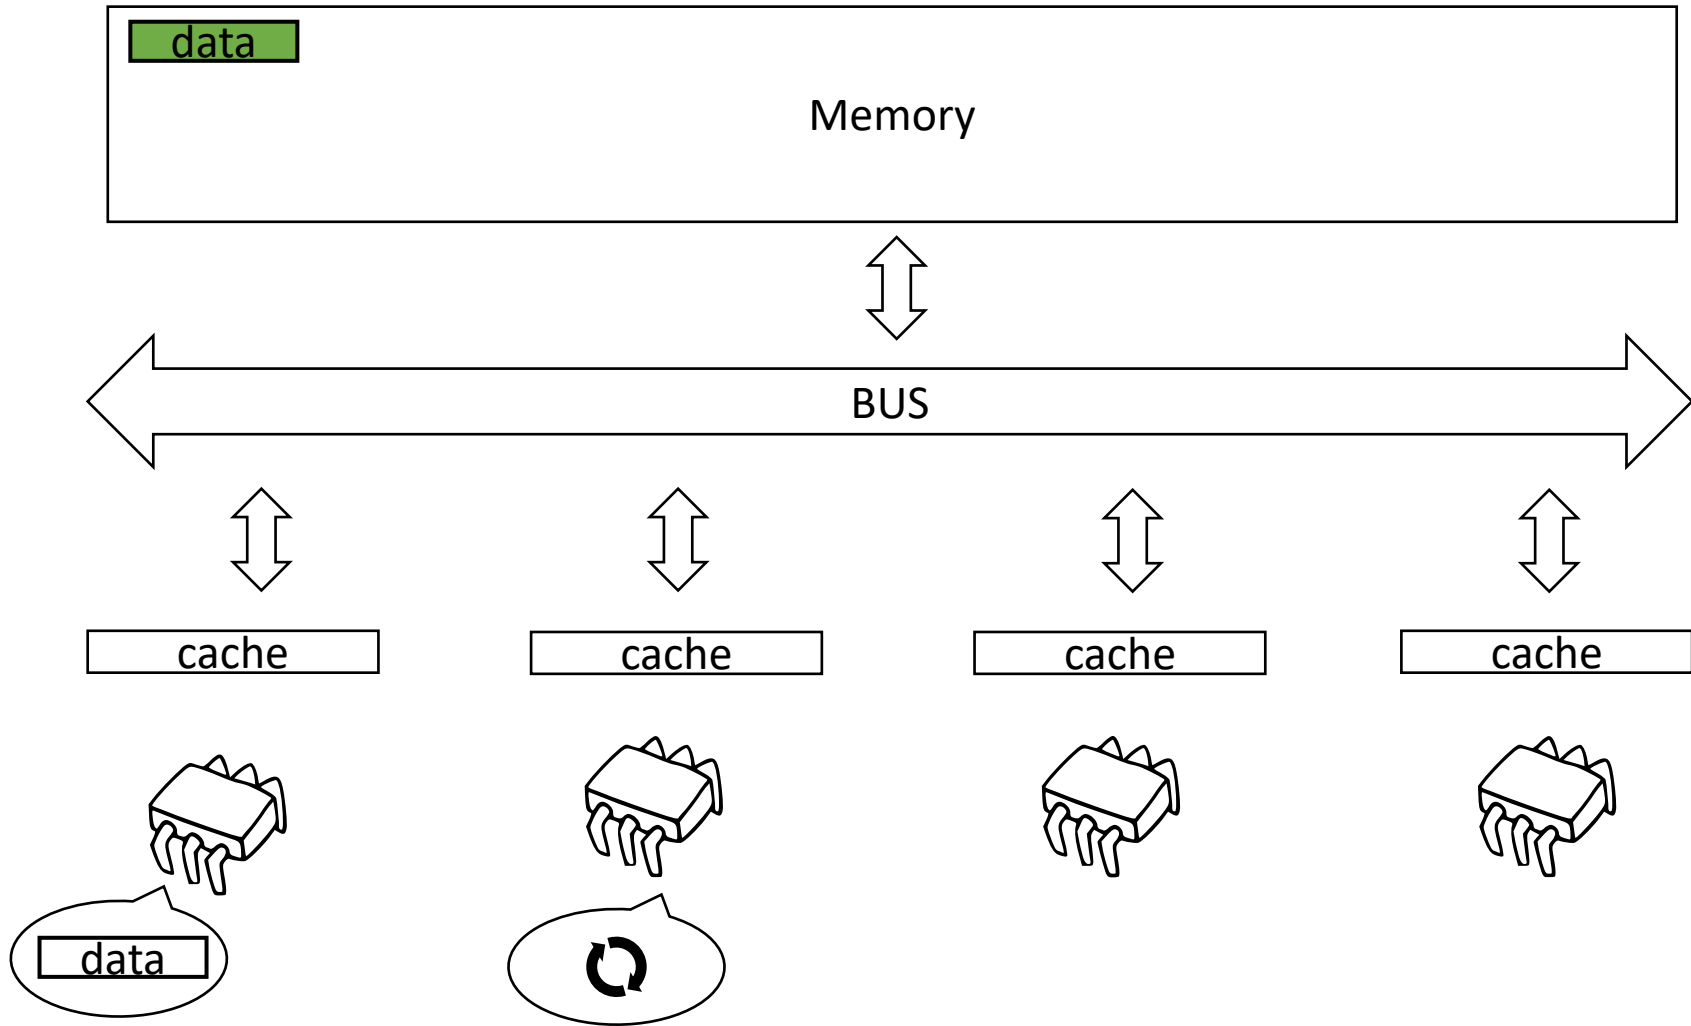
```
void acquire(int *lock){  
    while(XCHG(lock, 1));  
}
```

```
void release(int *lock){  
    *lock = 0;  
}
```

# Results



# Memory Model




# Test-and-set spin lock

- Test-and-set lock is the simplest spin lock
- Acquiring threads always try to set a variable via RMW

```
int lock = 0;
```

```
void acquire(int *lock){  
    while(XCHG(lock, 1));  
}  
  
void release(int *lock){  
    *lock = 0;  
}
```



# Test-and-set spin lock

- Test-and-set lock is the simplest spin lock
- Acquiring threads always try to set a variable via RMW

```
int lock = 0;
```

```
void acquire(int *lock){  
    while(XCHG(lock, 1));  
}  
  
void release(int *lock){  
    *lock = 0;  
}
```

We can reduce the impact of memory traffic by introducing exponential back off!  
But how to set it properly?






# Test-and-test-and-set spin lock


- Like test-and-set, but spins by reading the value of the lock
- Traffic is generated only upon lock handover

```
int lock = 0;
```

```
void acquire(int *lock){  
    while(XCHG(lock, 1))  
        while(*lock);  
}  
  
void release(int *lock){  
    *lock = 0;  
}
```



# Results



# Test-and-test-and-set spin lock

- Like test-and-set, but spins by reading the value of the lock
- Traffic is generated only upon lock handover

```
int lock = 0;
```

```
void acquire(int *lock){  
    while(XCHG(lock, 1))  
        while(*lock);  
}  
  
void release(int *lock){  
    *lock = 0;  
}
```

- Lock handover costs increase with the concurrency level
- Very lightweight for the uncontended case
- Is it feasible reducing handover costs?
- AND IMPROVING FAIRNESS?

# FIFO locks

# Ticket locks

- Similar to the bakery algorithm but it uses RMW instructions
  - Two variables
    - The next available ticket
    - The served ticket
- ```
typedef struct _tck_lock{  
    int ticket  = 0;  
    int current = 0;  
} tck_lock;
```

```
void acquire(tck_lock *lock){  
    int cur_tck;  
    int mytck = fetch&add(lock->ticket, 1);  
    while(mytck != (cur_tck = lock->current) )  
        delay((mytck-cur_tck)*BASE);  
}  
void release(tck_lock *lock){ lock->current += 1; }
```


# Ticket locks

- Ensure fairness
- Similar structure w.r.t. TTAS spinlock
  - One variable updated once at each acquisition (better than TTAS)
  - Write-1-Read-N variable updated at each release (same as TTAS)
- How?

# Anderson queue lock

- Use similar to ticket lock
- Use the ticket to obtain an individual cache line


Ticket = ~~0~~ 1 2 3



# Anderson queue lock

- Use similar to ticket lock
- Use the ticket to obtain an individual cache line

Ticket = ~~0~~ 1 2 3

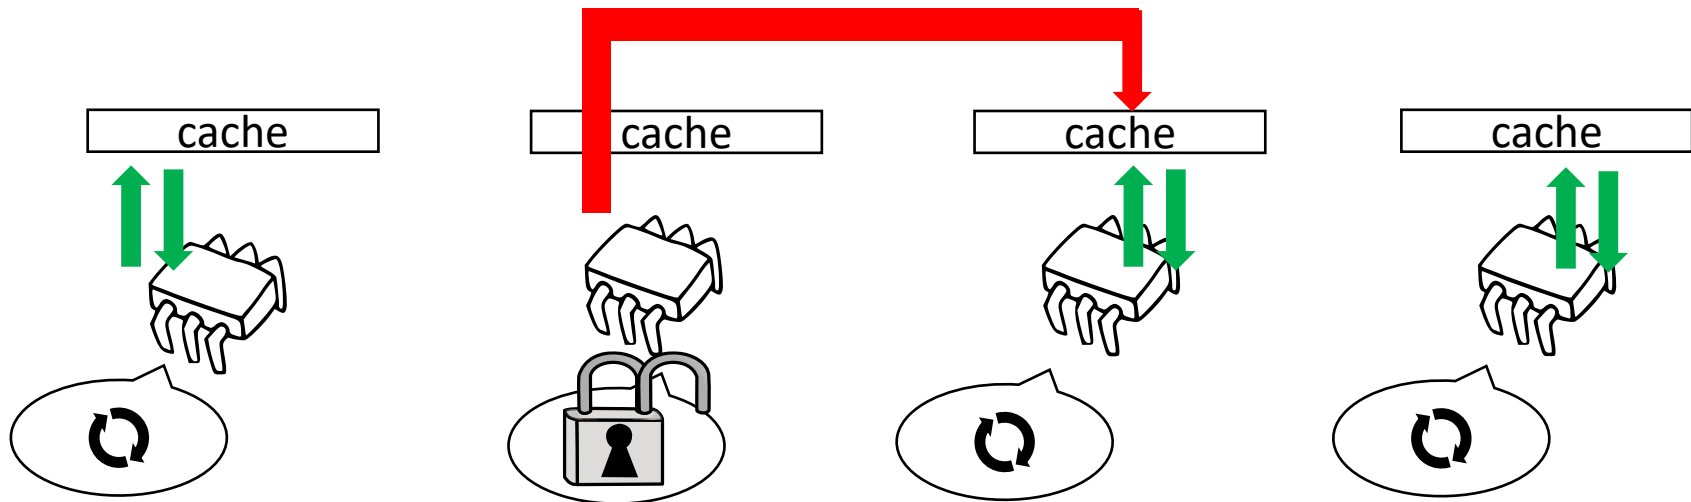




# Anderson queue lock

```
void acquire(android_lock *lock){  
    mytck = fetch&add(lock->ticket, 1);  
    while(lock->array[mytck]);  
    lock->array[mytck] = 1;  
}
```

```
void release(int *lock){  
    lock->array[mytck+1] = 0;  
}
```

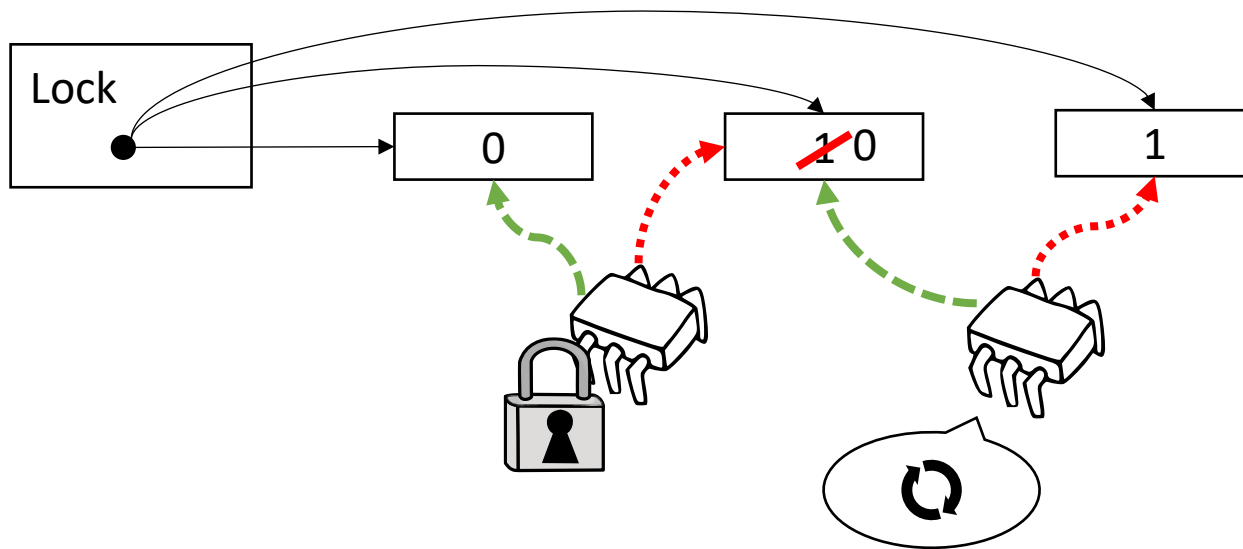


# Anderson queue lock

- Pros:
  - One variable updated once at each acquisition (like Ticket lock)
  - Write-1-Read-1 variable updated once per release (better than (T)TAS and Ticket)
- Cons:
  - Increased memory footprint
  - Each lock needs to know the maximum number of threads
- Let:
  - $T$  be the number of threads
  - $L$  be the number of locks
- Space Usage
  - Anderson =  $O(LT)$
  - TAS, TTAS, Ticket =  $O(L)$

# CLH lock

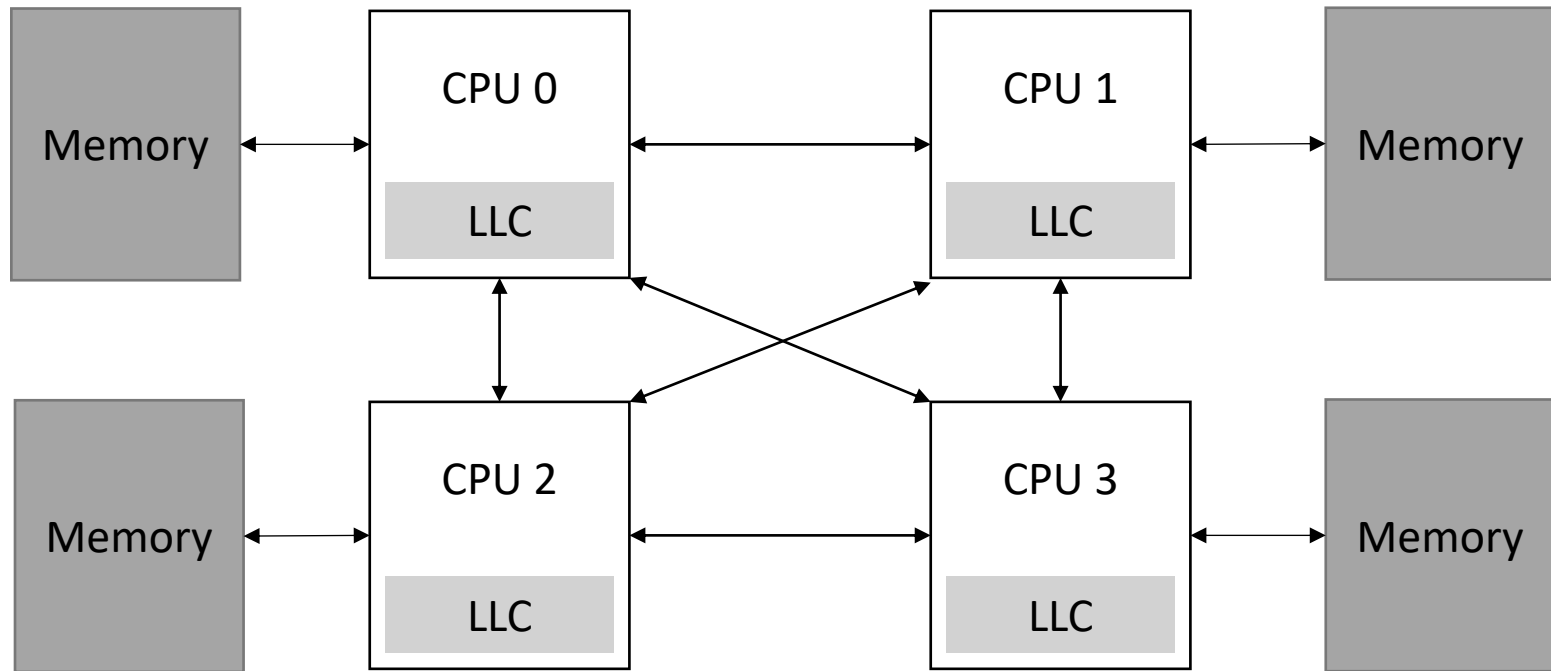
- An (implicit) linked list maintains the order between waiting threads
- An empty list represent an uncontended lock
- An arriving thread swaps the node with its private node
- Spin on the previous node
- Release on the new node



# CLH queue lock

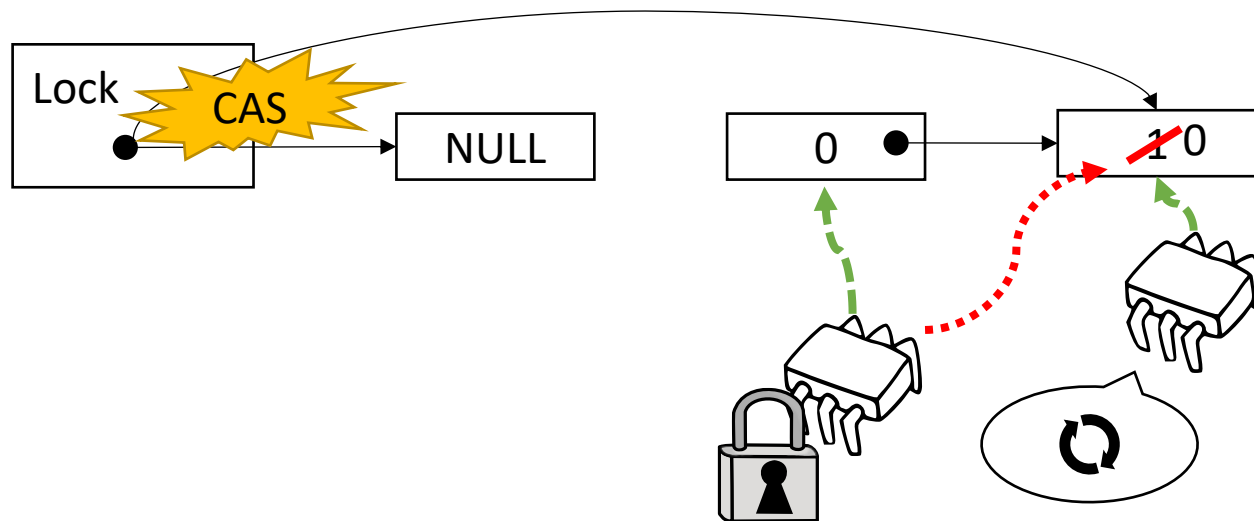
- Pros:
  - One variable updated once at each acquisition (like Ticket lock)
  - Write-1-Read-1 variable updated once per release (better than (T)TAS and Ticket)
- Cons:
  - Slightly increased memory footprint
- Let:
  - T be the number of threads
  - L be the number of locks
- Space Usage
  - CLH =  $O(L+T)$
  - Anderson =  $O(LT)$
  - TAS, TTAS, Ticket =  $O(L)$

# NUMA



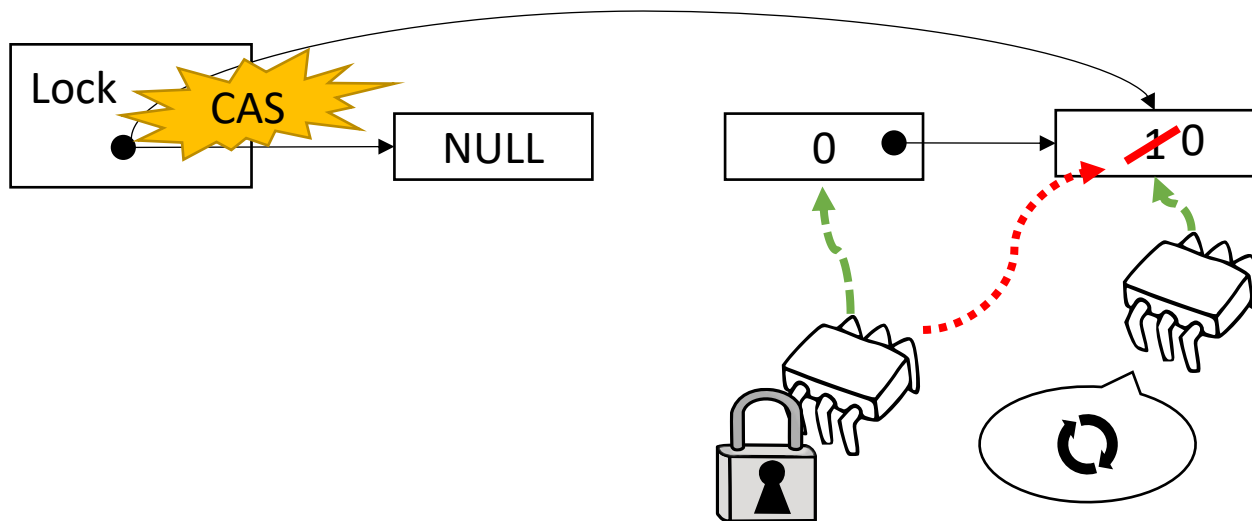
# MCS lock

- An explicit linked list maintains the order between waiting threads
- An empty list represent an uncontended lock
- An arriving thread swaps the node with its private node
- Spin on the just inserted node
- Release on the new node



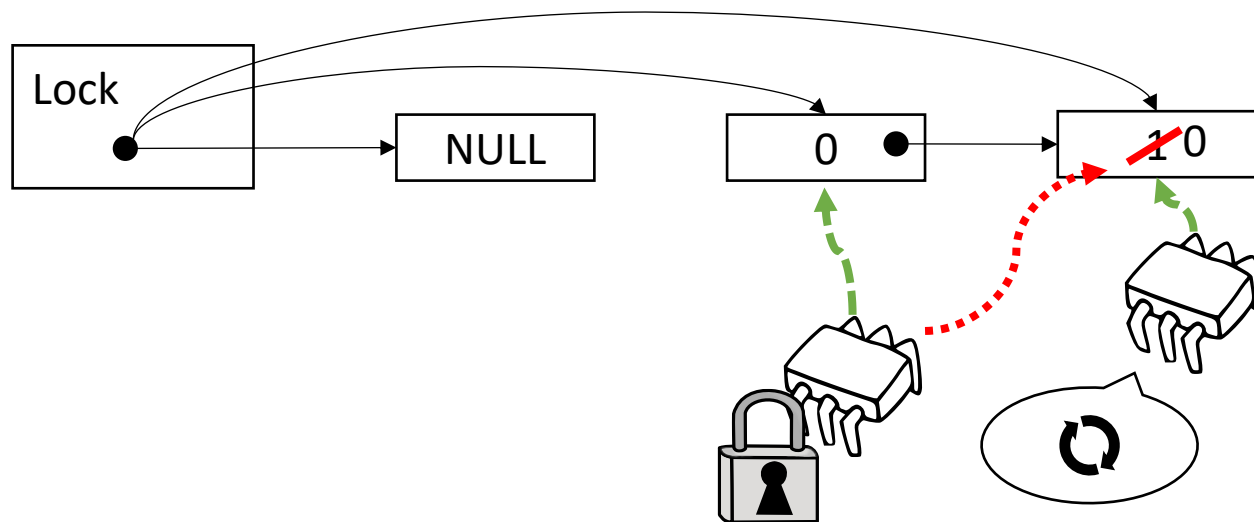
# MCS lock

- An explicit linked list maintains the order between waiting threads
- An empty list represent an uncontended lock
- An arriving thread swaps the node with its private node
- Spin on the just inserted node
- Release on the new node



# MCS lock

- An explicit linked list maintains the order between waiting threads
- An empty list represent an uncontended lock
- An arriving thread swaps the node with its private node
- Spin on the just inserted node
- Release on the new node



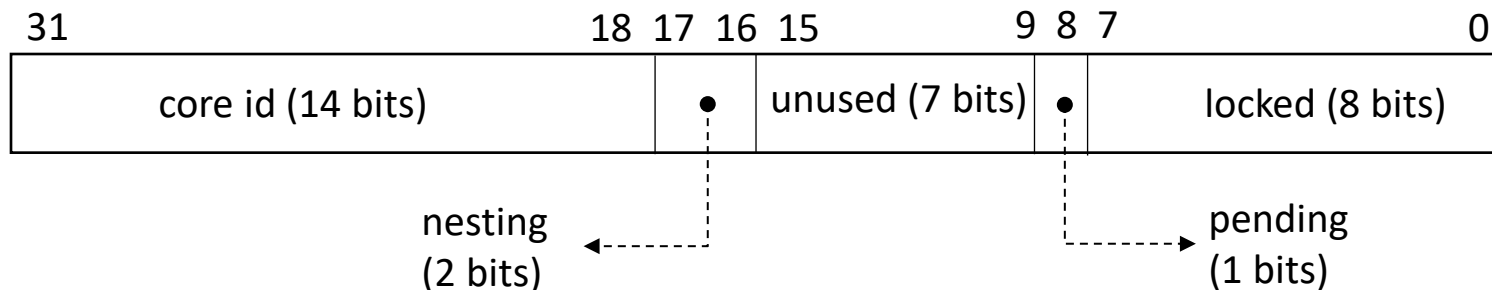


# MCS queue lock

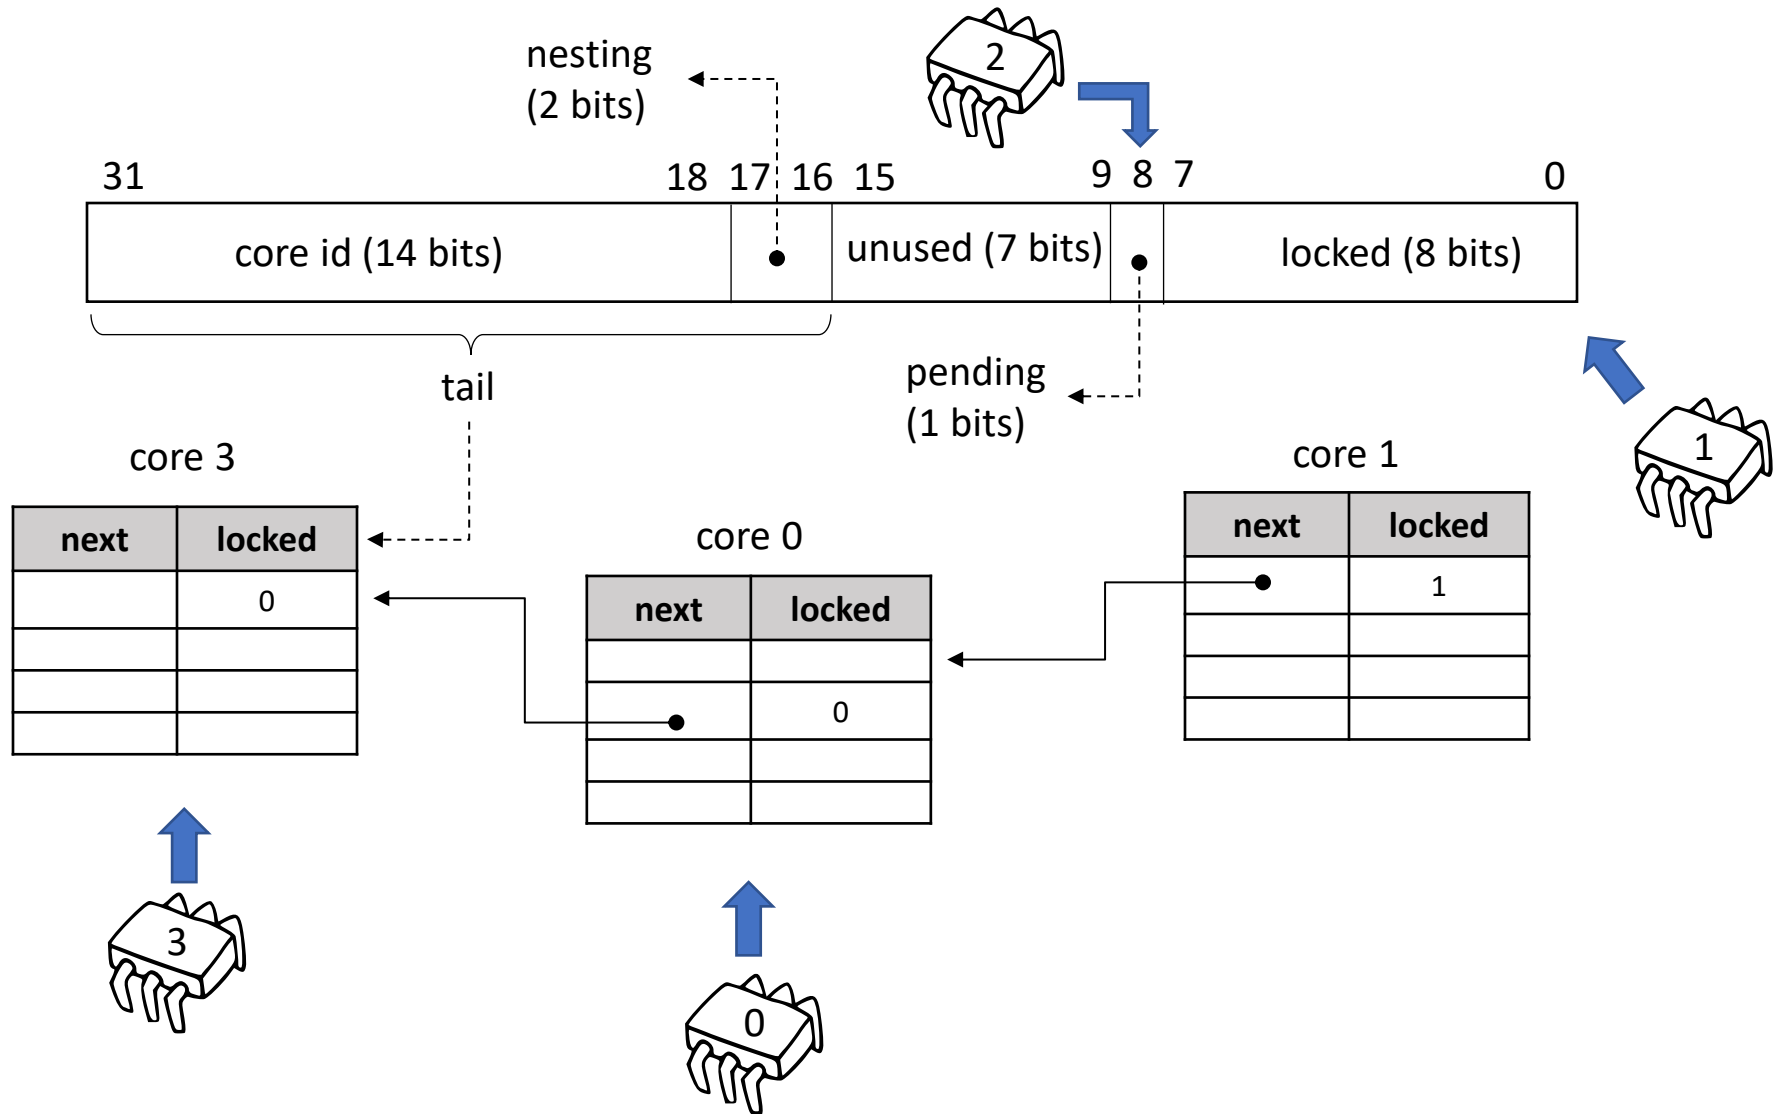
- Pros:
  - One variable updated once at each acquisition (like Ticket lock)
  - Write-1-Read-1 variable updated once per release (better than (T)TAS and Ticket)
  - No-remote spinning
- Cons:
  - Slightly increased memory footprint
- Let:
  - T be the number of threads
  - L be the number of locks
- Space Usage
  - MCS, CLH =  $O(L+T)$
  - Anderson =  $O(LT)$
  - TAS, TTAS, Ticket =  $O(L)$

# MCS in practice: the Linux kernel case

- The Linux kernel uses a particular implementation of a MCS lock: Qspinlock
- Additional challenge:
  - Maintain compatibility with classical 32-bit locks
  - MCS uses pointers (64-bit)
- Compact data:
  1. No recursion of same context in critical sections
  2. 4 different contexts (task, softirq, hardirq, nmi)
  3. Finite number of cores
- Use an additional bit for fast lock handover



# MCS in practice: the Linux kernel case



# A small benchmark

- We have an array of integers
- Each thread reverse the array



- This is done within a critical section

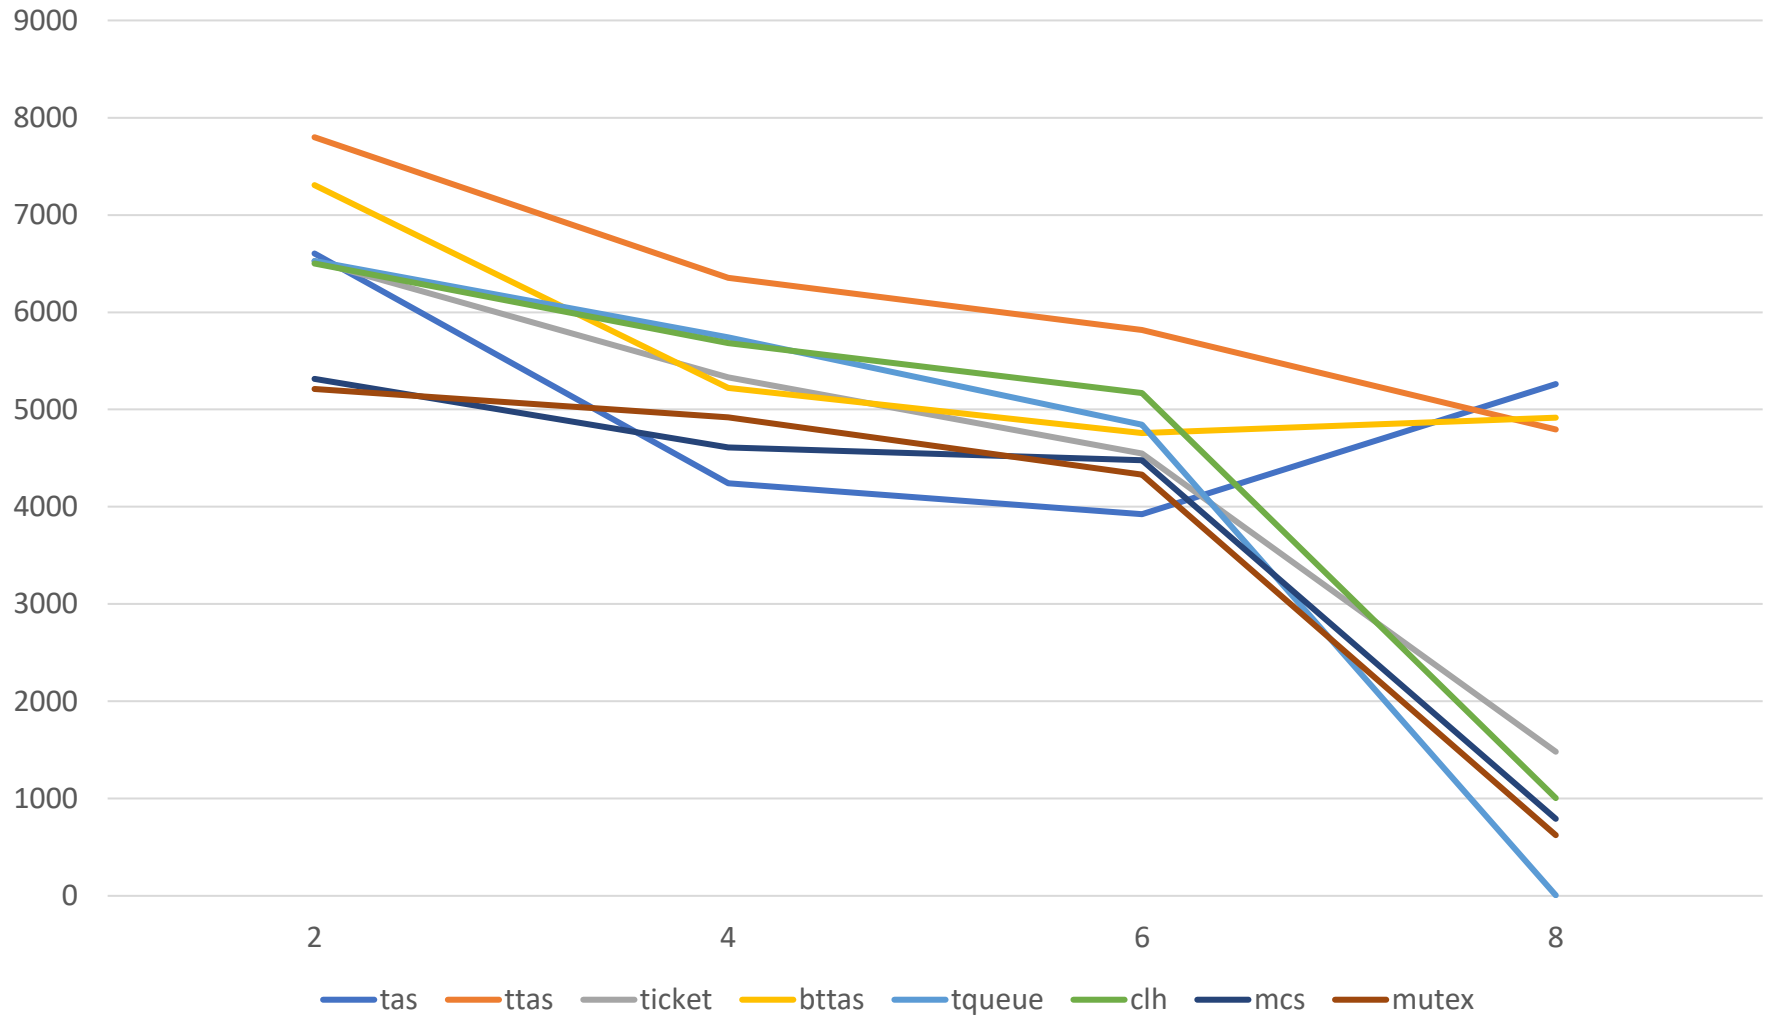
```
while(!stop){  
    acquire(&lock);  
    flip_array();  
    release(&lock);  
}
```

- Performance Metric:
  - Throughput = #Flips per second

**One lock  
to rule them all...**


# Performance

Intel i7-7700HQ – 8 cores



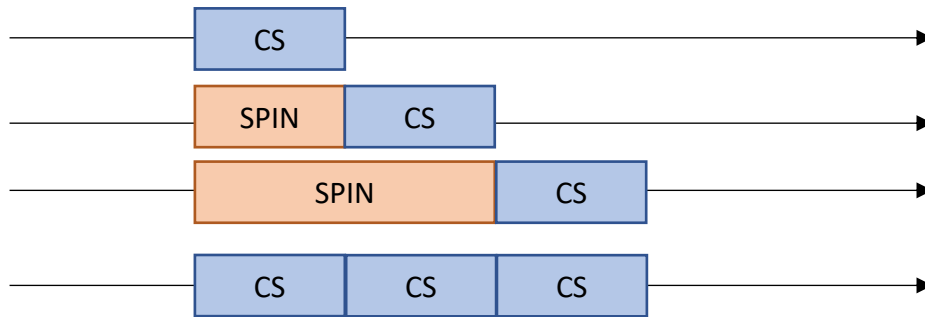
# Performance

AMD Opteron 6168 - 48 cores



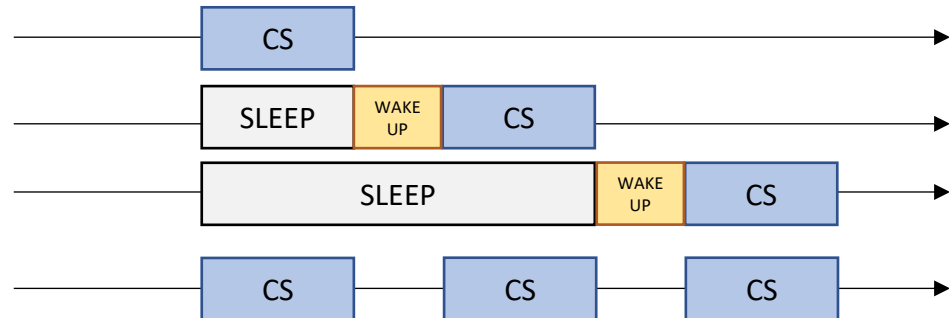
# At the beginning was... Spin vs Sleep

| Benefits                | Waiting Policy |          |
|-------------------------|----------------|----------|
|                         | Spinning       | Sleeping |
| Guaranteed low latency  | ✓              | ✗        |
| Computing power savings | ✗              | ✓        |



SPIN:  
++Waste of CPU Cycles  
--Latency

Sleep:  
--Waste of CPU Cycles  
++Latency





# How to avoid costs for sleeping?

A general approach exists:

- Reducing the frequency of sleep/wake-up pairs
- How?

→ Trading Fairness in favor of Throughput

- Make some thread sleep longer than others
- If the lock is highly contented, some thread willing to access the critical section will arrive soon
- If the lock is scarcely contented, we pay lower latency as TTAS locks

# An example - MutexEE

- MutexEE is a pthread\_mutex optimized for throughput and energy efficiency

**lock()**

| MUTEX                  | MUTEXEE |
|------------------------|---------|
| For up to 100 attempts |         |
| spin with pause        |         |
| if still busy, sleep   |         |


**unlock()**

| MUTEX                                    | MUTEXEE |
|------------------------------------------|---------|
| release in user space (lock->locked = 0) |         |
|                                          |         |
| wake up a thread                         |         |

Credits: Falsafi et al. "Unlocking energy"

# An example - MutexEE

- MutexEE is a pthread\_mutex optimized for throughput and energy efficiency



- Global lock
- 1000 cycles CS
- 40 cores

MUTEX x

TAS +

TTAS o

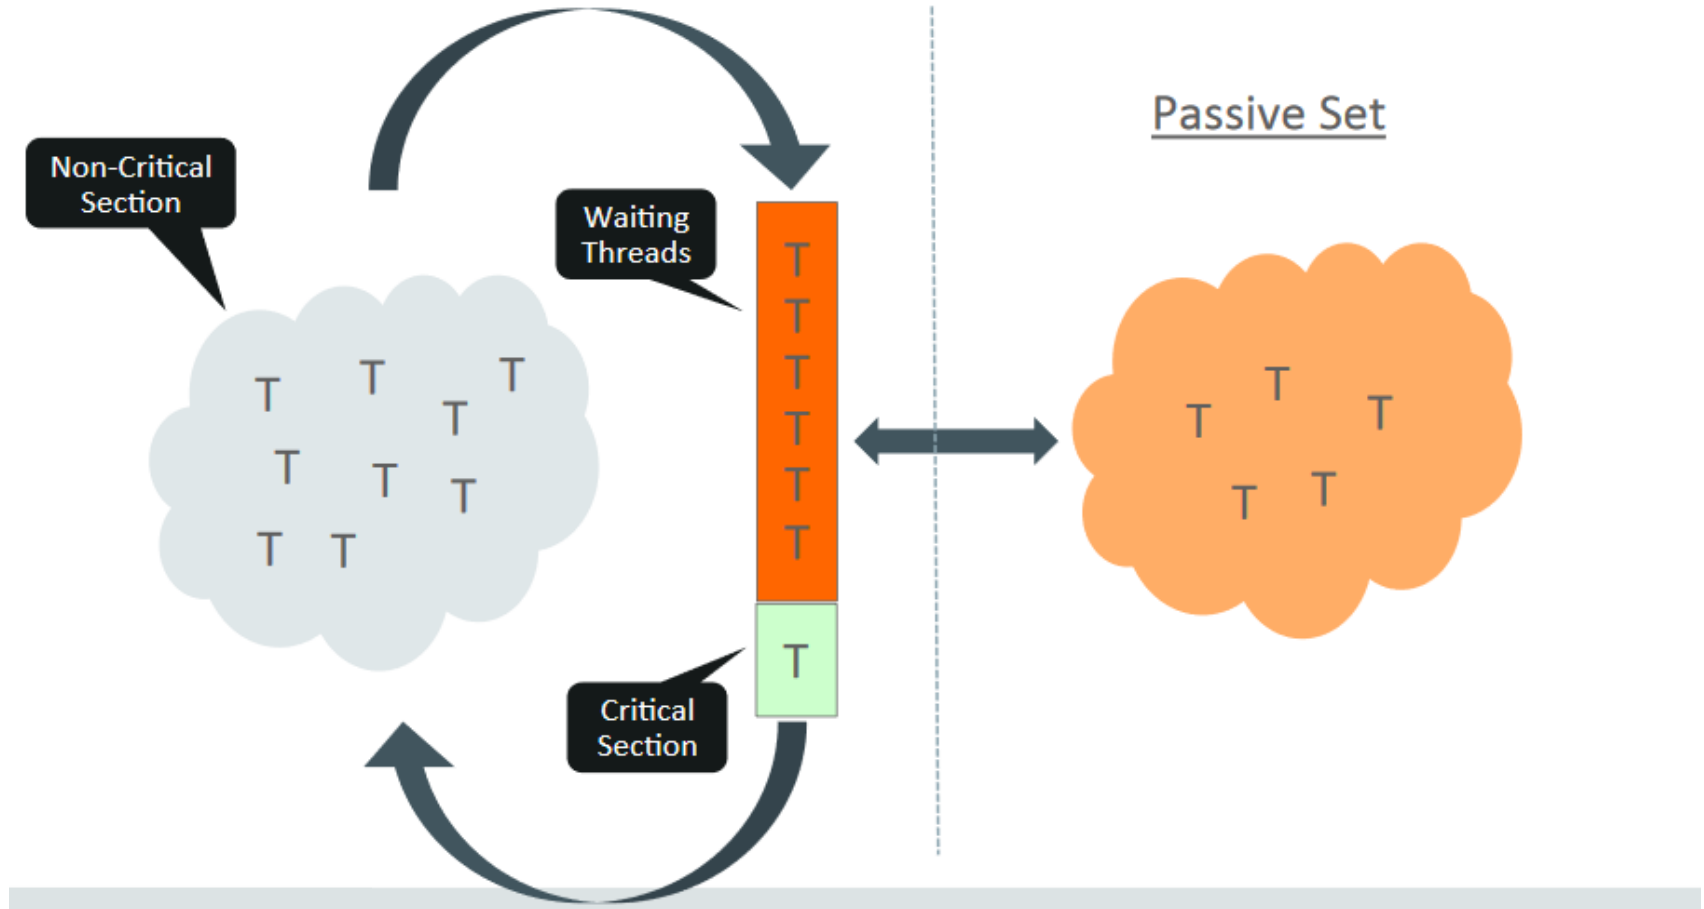
TICKET ■

MCS \*

MutexEE □


Credits: Falsafi et al. "Unlocking energy"

# An example 2 – Malthusian locks



Credits: Dave Dice "Malthusian locks"

# An example 2 – Malthusian locks



Credits: Dave Dice "Malthusian locks"