

Programmazione concorrente

Laurea Magistrale in Ingegneria Informatica

Università Tor Vergata

Docente: Romolo Marotta

Transactional Memory

Synchronization approaches:

- **Non-blocking data structures**
- **Locks**
- **Transactional Memory**

Transactional Memory

- Why?
 - Fine grain locking (or non-blocking synchronization) can scale but it is hard
 - Locks do not scale in general, but they are hard too:
 - Deadlocks
 - Races (forgotten locks)
 - Do not compose
- Transactions: **Begin_transaction**
 - `x.op()`
 - `y.op2(k)`
 - `z.op(j)`
 - End_transaction**
 - They compose (e.g. nested transactions)
 - Simpler to reason about

Transactions

- Well known in the context of databases
- Conceived integration of transaction in hardware (1993)
- Software implementations (1995-2005)
- Commercial hardware support (2013)




Transactions

Transaction on top of
DBMS



Transaction on top of
Transactional Memory



```
x = 2; y = 1
```

Begin:

```
d = x
```

```
n = y
```

```
write(z, 1/(n-d))
```

Abort

Begin:

```
y++
```

```
x++
```



Commit

Transactions

Transaction on top of
DBMS



Transaction on top of
Transactional Memory



```
x = 2; y = 1
```

Begin:

```
d = x
```

```
n = y
```

```
write(z, 1/(n-d))
```

Abort

Begin:

```
y++
```

```
x++
```



Commit

Transactions

Transaction on top of
DBMS



Transaction on top of
Transactional Memory



$x = 2; v = 1$

B

Managed by DBMS instead
of developers

Begin:

$y++$

$x++$

Commit

$n = y$

$\text{write}(z, 1/(n-d))$

Abort

Concurrent and parallel programming

Float Exceptions are not
transparent to developers

Transactions

Transaction on top of
DBMS



Transaction on top of
Transactional Memory

(view) serializability:
Committed transactions see
consistent values

Opacity:
Both committed and aborted
transactions see
consistent values

$x = 2; v = 1$

B

Managed by DBMS instead
of developers

Begin:

$y++$

$x++$

Commit

$n = y$

$\text{write}(z, 1/(n-d))$

Abort

Float Exceptions are not
transparent to developers

Histories

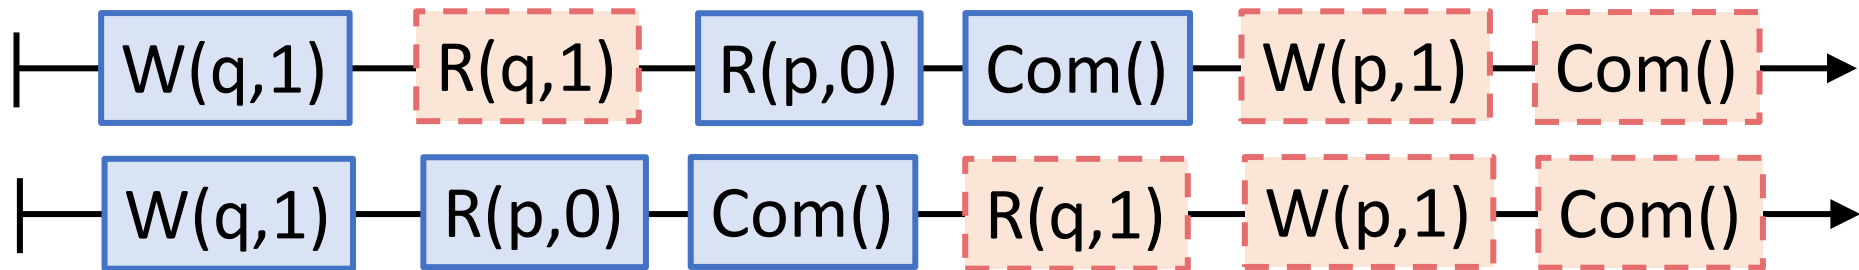
- The execution of transaction on a set of objects is modeled by a history
- A history is a sequence of:
 - Operations (e.g., read, write, push, pop ...)
 - Commits
 - Aborts
- Two transactions are:
 - **sequential** if one invokes its first operations after the other one commits or aborts
 - **concurrent** otherwise
- A history is:
 - **sequential** if has only sequential transactions
 - **concurrent** otherwise
- Two histories are equivalent if they have the same transactions

Correctness conditions (recall)

- A **concurrent execution** is correct if it is equivalent to a correct **sequential execution**
- ⇒ A **history** is correct if it is **equivalent** to a ~~correct~~ **sequential history** which satisfies a given correctness condition
- A correctness condition specifies the set of histories to be considered as reference
- ⇒ In order to implement correctly a concurrent object wrt a correctness condition, we must guarantee that every possible history on our implementation satisfies the correctness condition

(View) Serializability [Papadimitriou1979]


- A history H of committed transactions is serializable if
 - It is equivalent to a sequential history H'
 - H' is sequential
 - H' is legal, aka every read returns the last written value
- Serializable?



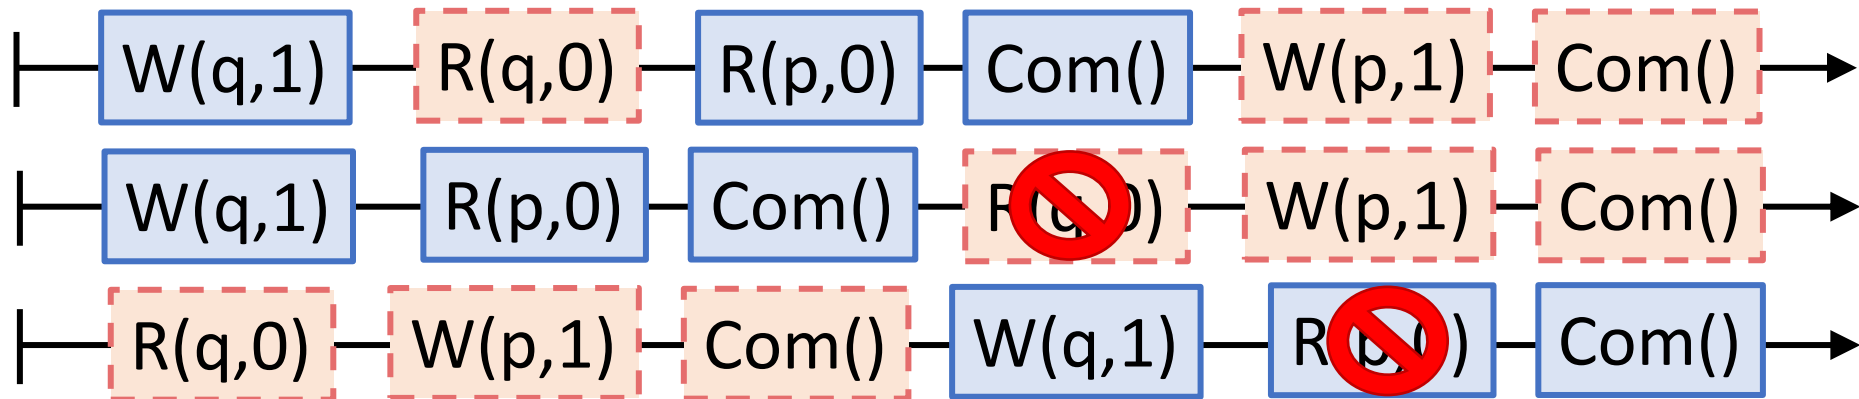
(View) Serializability [Papadimitriou1979]

- A history H of committed transactions is serializable if
 - It is equivalent to a sequential history H'
 - H' is sequential
 - H' is legal, aka every read returns the last written value

- Serializable? Yes!




- Serializable?




(View) Serializability [Papadimitriou1979]

- A history H of committed transactions is serializable if
 - It is equivalent to a sequential history H'
 - H' is sequential
 - H' is legal, aka every read returns the last written value

- Serializable? Yes!




- Serializable? No!




(View) Serializability [Papadimitriou1979]

- A history H of committed transactions is serializable if
 - It is equivalent to a sequential history H'
 - H' is sequential
 - H' is legal, aka every read returns the last written value


- Serializable? Yes!



- Serializable? No!

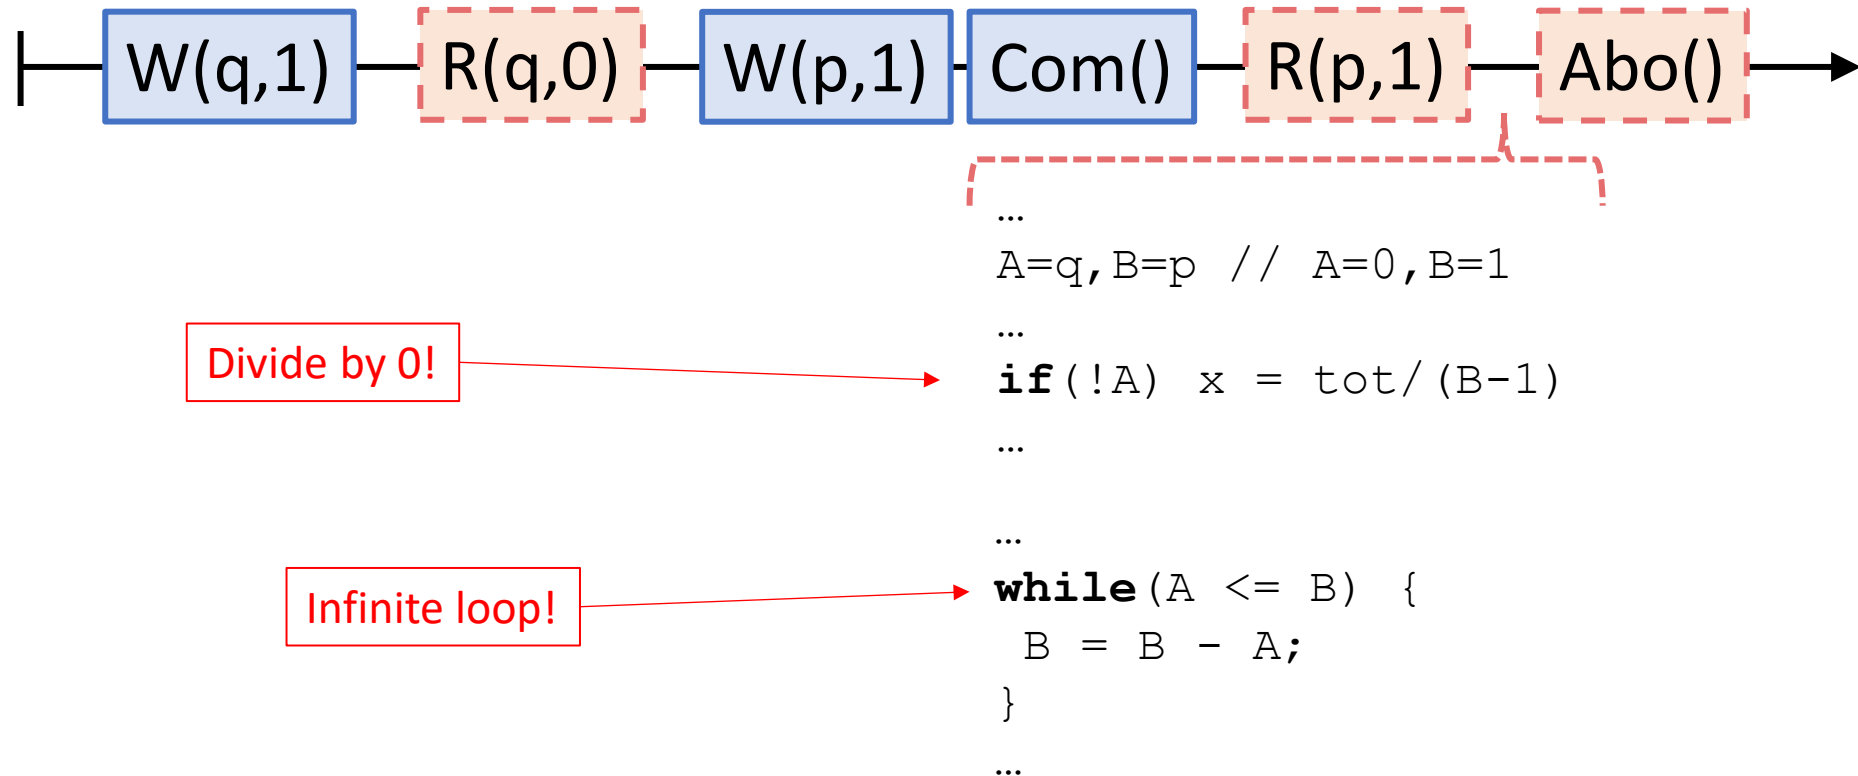


- Serializable? Yes!



(View) Serializability [Papadimitriou1979]

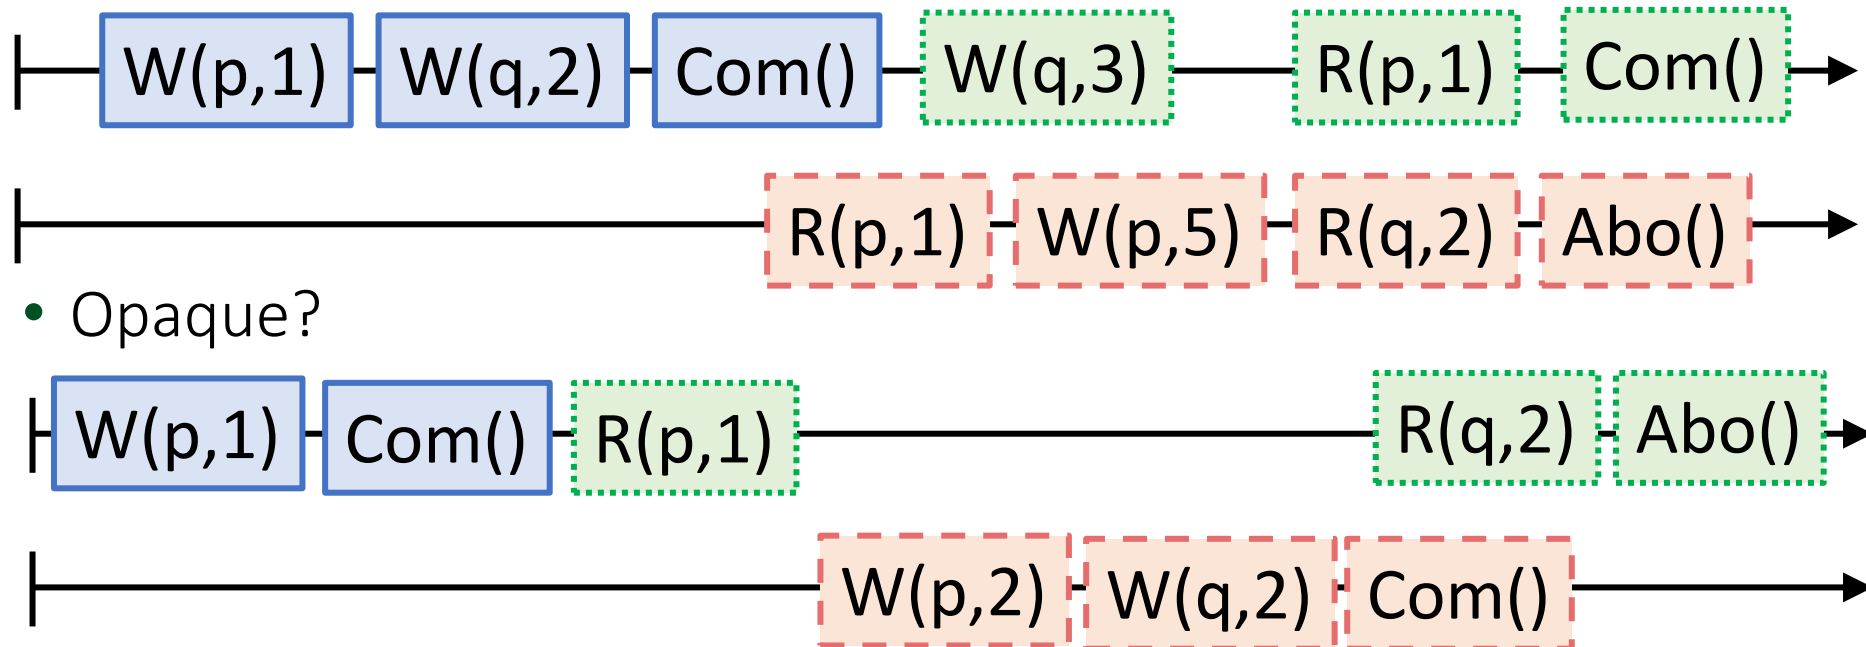
- Serializable? Yes! But, what happens in the case of TM?



- Could strict serializability be of any help?
 - Serializability + Real-time order
 - It predicates only on committed transactions

Opacity [Guerraoui2008]

- A history H is opaque if
 - It is equivalent to a sequential history H'
 - H' is sequential
 - H' preserves transactions' real-time order
 - H' is legal
- Opaque?



Transactions

Transaction on top of
DBMS



Transaction on top of
Transactional Memory

(view) serializability:
Committed transactions see
consistent values

Opacity:
Both committed and aborted
transactions see
consistent values

$x = 2; v = 1$

B

Managed by DBMS instead
of developers

Begin:

$y++$

$x++$

Commit

$n = y$

$\text{write}(z, 1/(n-d))$

Abort

Float Exceptions are not
transparent to developers

Transactions

Deadlock or starvation
freedom



Obstruction
freedom

(view) serializability:
Committed transactions see
consistent values

Opacity:
Both committed and aborted
transactions see
consistent values

$x = 2; v = 1$

B

Managed by DBMS instead
of developers

Begin:

$y++$

$x++$

Commit

$n = y$


$\text{write}(z, 1/(n-d))$

Abort

Float Exceptions are not
transparent to developers


Wait freedom

- Every correct transaction eventually commits
- Finite number of aborts




Wait freedom

- Every correct transaction eventually commits
- Finite number of aborts



Wait freedom

- Every correct transaction eventually commits
- Finite number of aborts



IMPOSSIBLE IN AN ASYNCHRONOUS SYSTEM

Obstruction freedom

- Every correct transaction that runs in isolation (without contention) eventually commits
- Abort is unavoidable
- Contention manager can help with contention scenarios
- When a new transaction A creates a conflict with B
 - Aggressive
 - always abort B
 - Backoff
 - B waits an exp. back-off time, then abort A if still conflicting
 - Karma
 - Assign priority to A and B, abort lowest priority, increase priority after abort
 - Greedy
 - Use start time as priority, if $P_b < P_a$ and A is not waiting then B wait, otherwise abort A

Transactions

Deadlock or starvation
freedom



Obstruction
freedom

(view) serializability:
Committed transactions see
consistent values

Opacity:
Both committed and aborted
transactions see
consistent values

`x = 2; y = 1`

Begin:

`d = x`

`n = y`

`write(z, 1/(n-d))`

Abort

Begin:

`y++`

`x++`

Commit

`x = 2; y = 1`

Begin:

`d = x`

Abort

Begin:

`y++`

`x++`

Commit

Software Transactional Memory

DSTM

JVSTM

RSTM

TL2

TinySTM

SwissTM

McRT-STM

Bartok-STM

NOrec

LSA


E-STM

SXM

ASTM

WSTM

PhTM

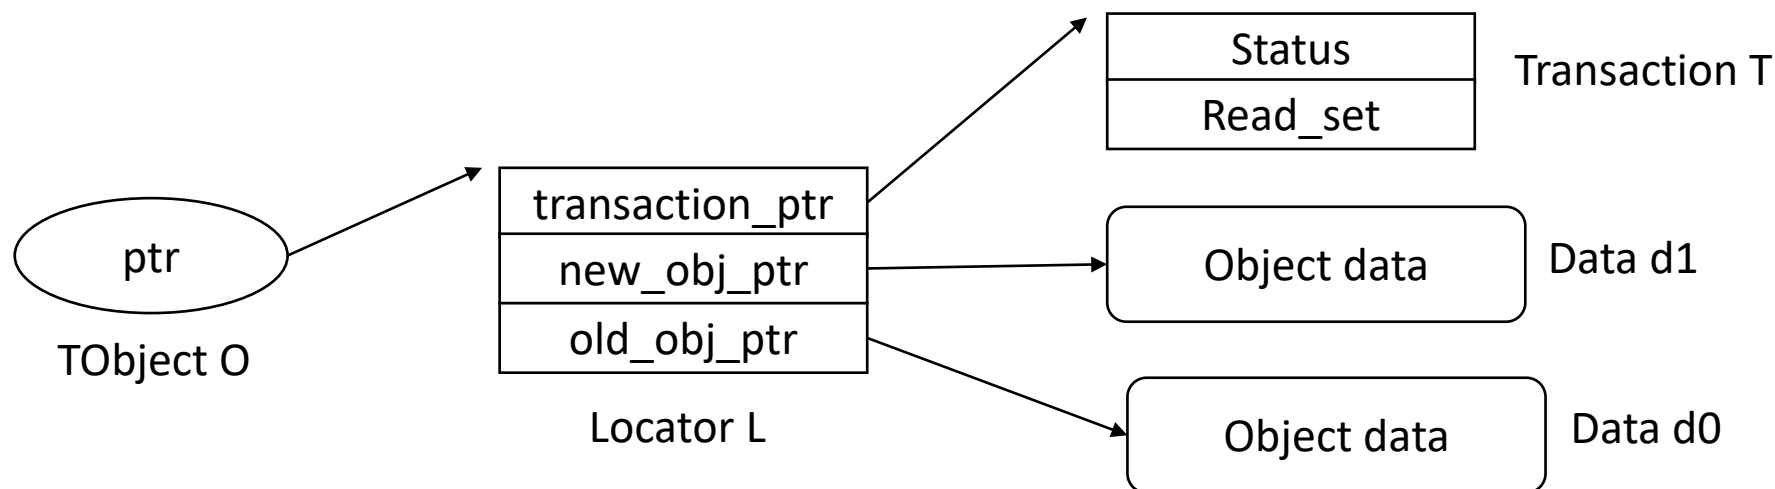


DSTM [Hearlihy2003]

- Obstruction freedom + contention manager
- It works at object granularity
 - Transactions open objects in READ/WRITE mode to apply an operation
 - Conflicts are detected when opening objects
- A conflicting write makes one of the two conflicting transaction abort via contention manager (**killer write**)
- A read requires that all already-read objects are still the most recently committed version (**careful read**)
- Validate all objects read upon commit

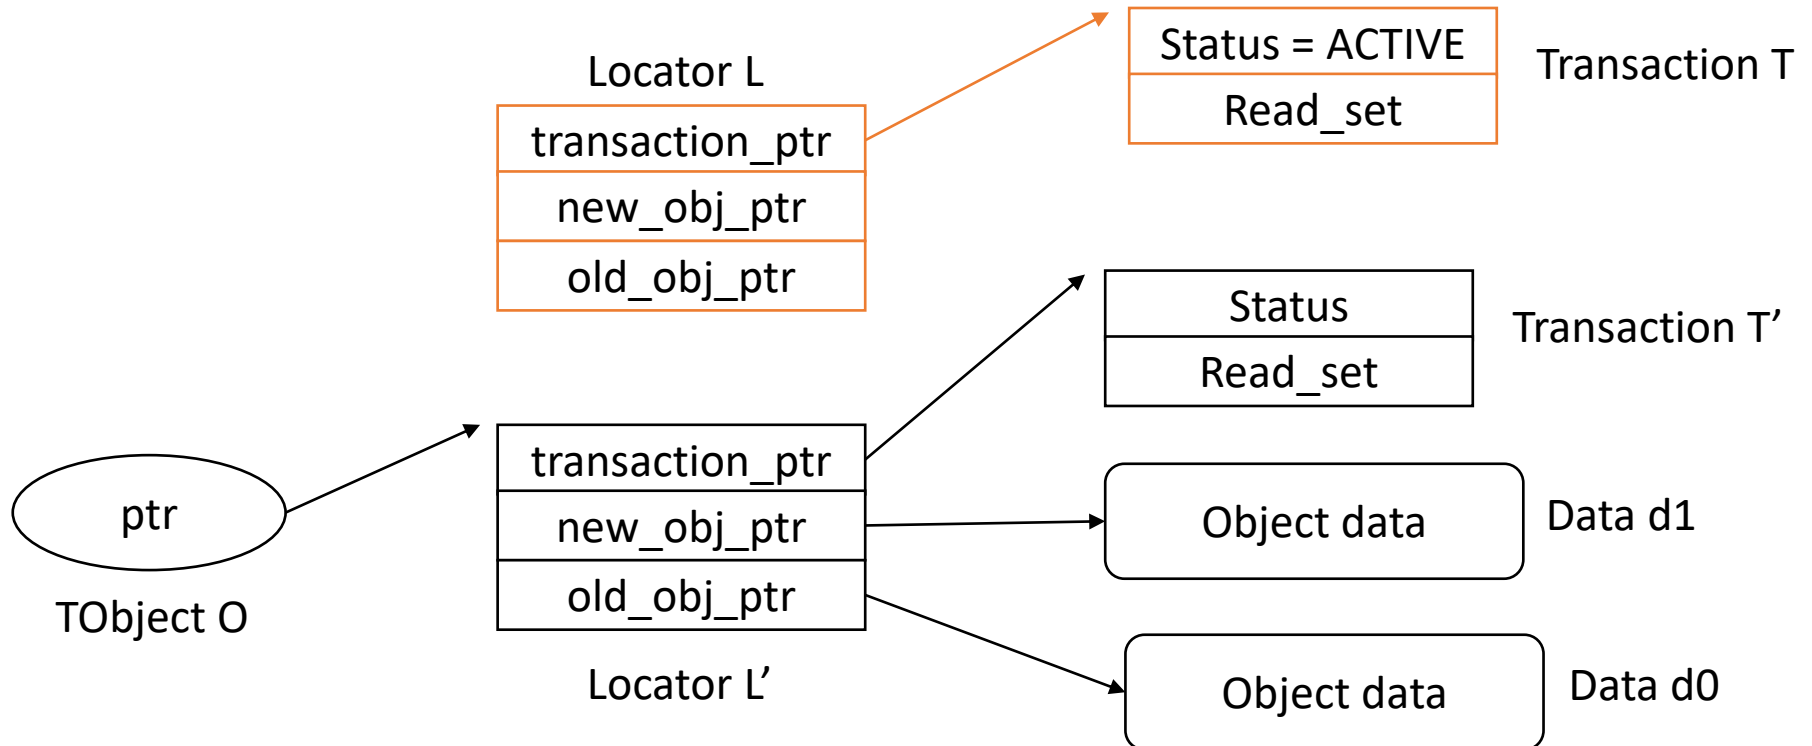
DSTM [Hearlihy2003]

- Transactions have:
 - A status
 - Committed
 - Active
 - Aborted
 - Collection of objects opened in READ mode
 - Objects are encapsulated within a Transactional Object which keeps references to
 - Transaction currently manipulating the object in WRITE mode
 - Current and tentative versions of the object
- with an intermediate object called Locator



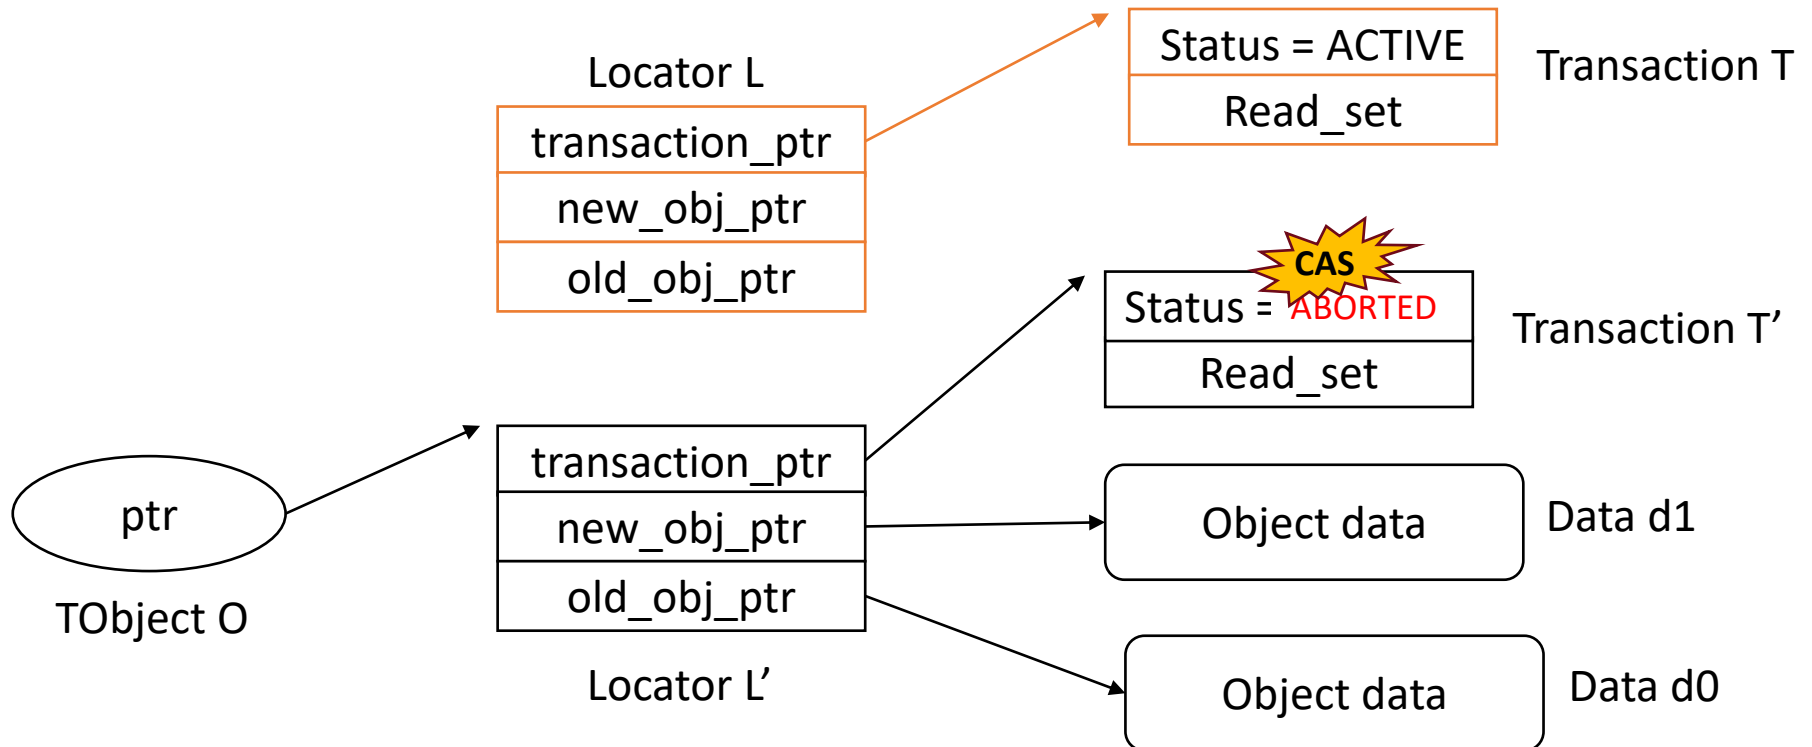
DSTM – First open in WRITE mode [Hearlihy2003]

- **T** is the current transaction, whose status is **ACTIVE**
- **T** allocates a new Locator **L**
- **T** accesses to current locator **L'** of **O** to retrieve last transaction **T'** that executed the last open in WRITE mode



DSTM – First open in WRITE mode [Hearlihy2003]

- **T** behaves accordingly to **T'** status
 - **ACTIVE:** **T** calls the contention manager
 - **T** waits a back-off time
 - **T** makes **T'** abort via **Compare&Swap**

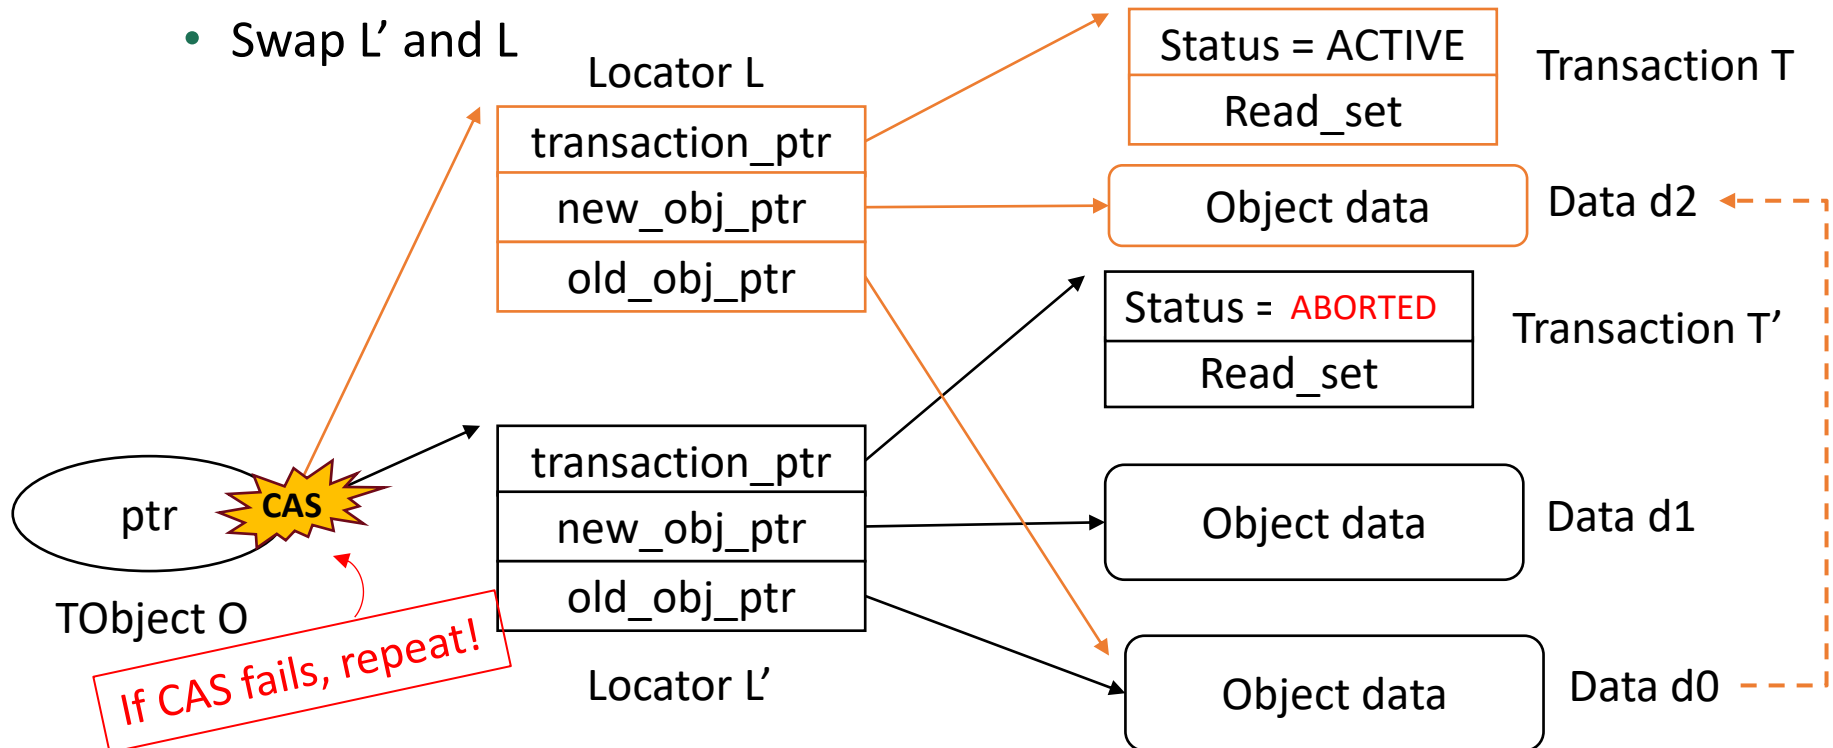


DSTM – First open in WRITE mode [Hearlihy2003]

- T behaves accordingly to T' status

- **ABORTED:**

- T use L'.old_obj_ptr to get current version of O
- L.old_obj_ptr = L'.old_obj_ptr
- L.new_obj_ptr = CLONE(L'.old_obj_ptr)
- Swap L' and L

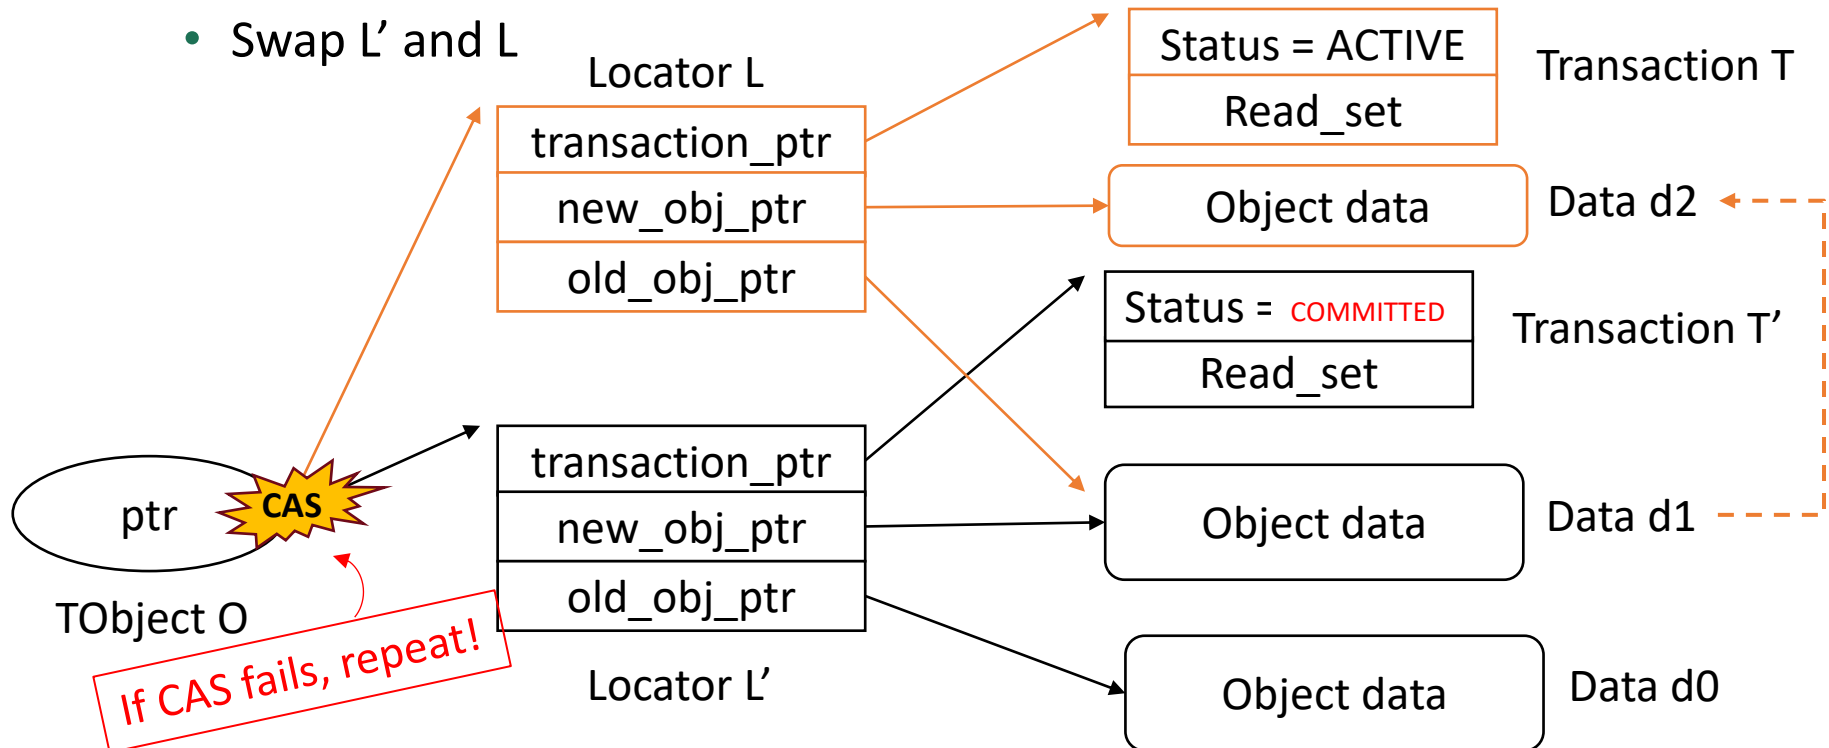


DSTM – First open in WRITE mode [Hearlihy2003]

- T behaves accordingly to T' status

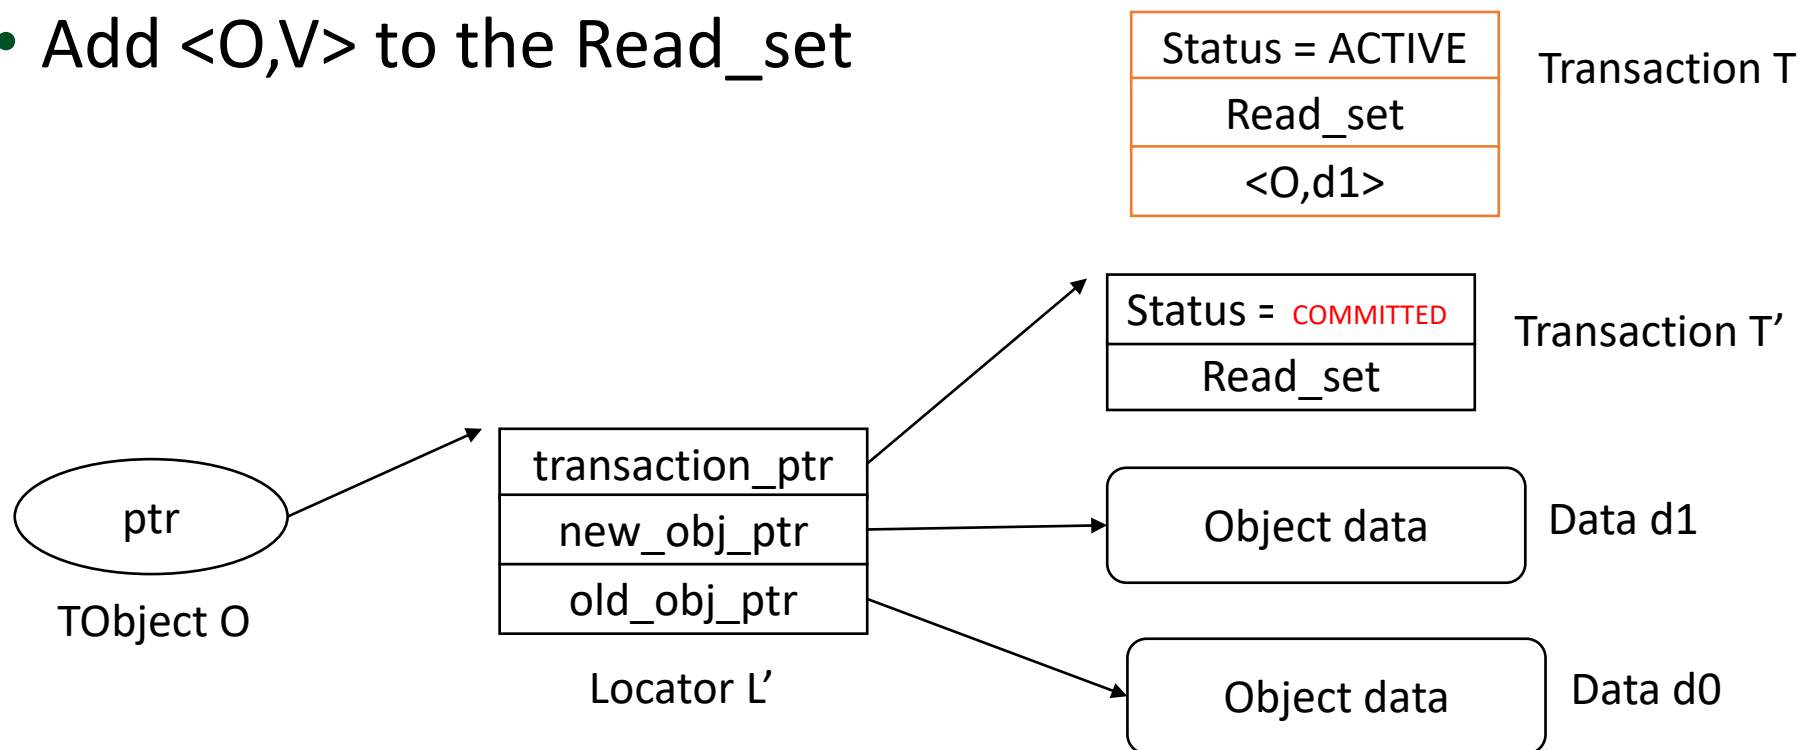
- **COMMITTED:**

- T use L'.old_obj_ptr to get current version of O
- L.old_obj_ptr = L'.new_obj_ptr
- L.new_obj_ptr = CLONE(L'.new_obj_ptr)
- Swap L' and L



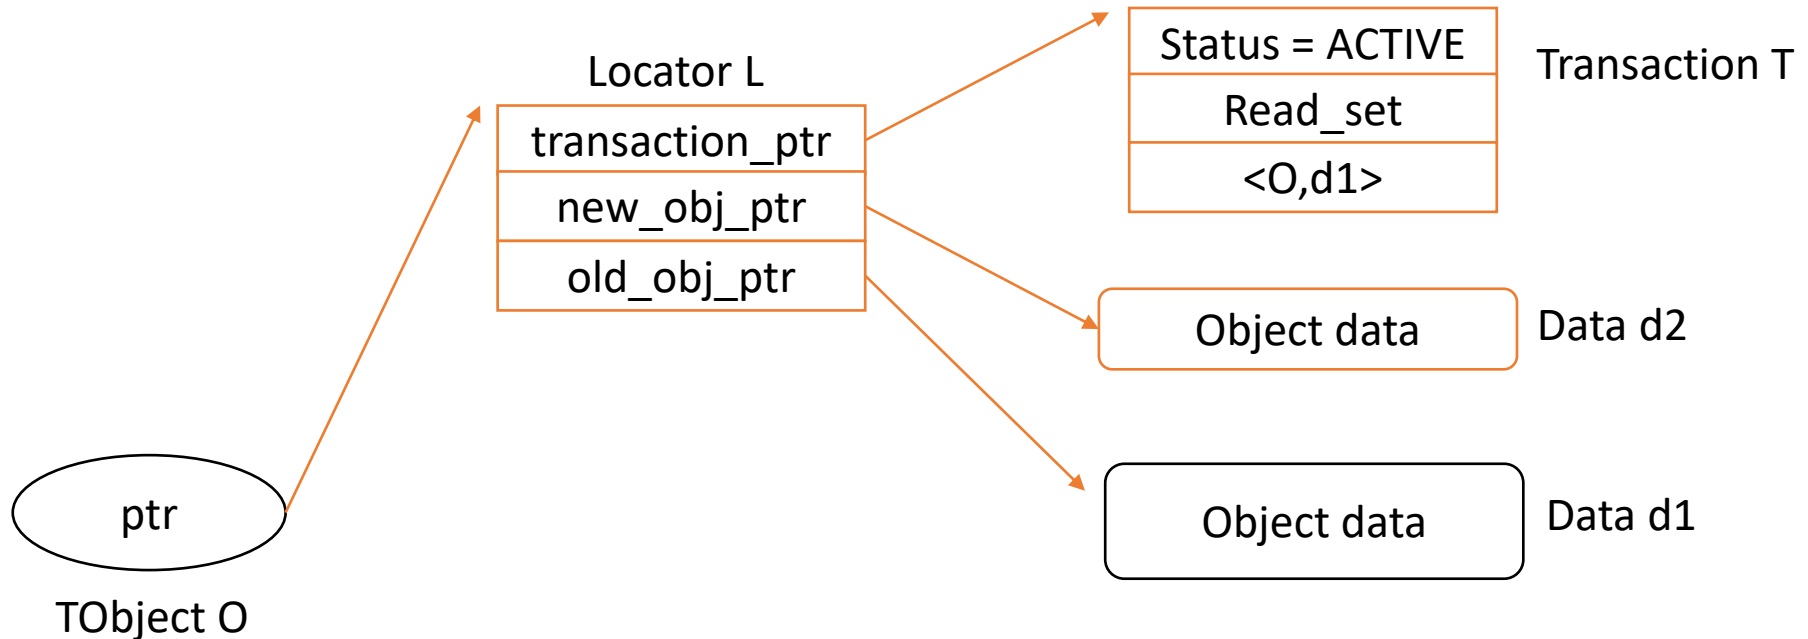
DSTM – First open in READ mode [Hearlihy2003]

- Validate Read_set (see later)
- Fetch current **committed** version **V** via current locator
 - New_obj_ptr if T' is committed
 - Old_obj_ptr otherwise
- Add $\langle O, V \rangle$ to the Read_set



DSTM – Already opened objects [Hearlihy2003]

- Already opened in READ mode:
 - Retrieve **V** from the Read_set
- Already opened in WRITE mode:
 - Retrieve **V** from the current locator



DSTM – Commit [Hearlihy2003]

1. Validate the transaction

- Transaction aborts on WRITE/WRITE conflicts
 - No need to validate WRITE upon commit
- Validate Read_set
 - For each pair $\langle O, V \rangle$ check that V is still the most recent committed version
- Read_set validation is non atomic
 - Check the status is still ACTIVE

2. If OK Change status then from **ACTIVE** to **COMMITTED** else from **ACTIVE** to **ABORTED**

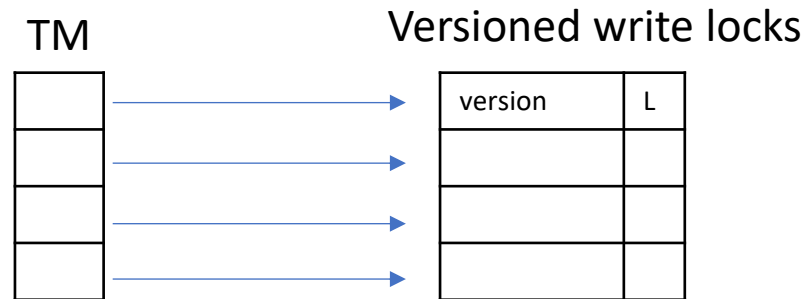
- Individual CAS

DSTM – Final remarks [Hearlihy2003]

- Read-only transactions do not need any ATOMIC instruction for each read
- Committed transactions appear to take effect when the transition ACTIVE->COMMITTED occurs
 - Linearizable/Strict serializable
- Why careful read (validation at each read)?
- Obstruction freedom
 - Transactions abort iff conflicts occur

Transactional Locking 2 [Dice2006]

- Word-based STM
 - Each transactional memory location is associated with a **versioned write lock** <version,is_locked>



- Exploits a Global Version Clock (GVC) to quickly detect updates (it increases before a write-transaction commits)
- Transactions keep track of
 - GVC
 - Read set
 - Write set

Transactional Locking 2 [Dice2006]

- BEGIN:
 - Sample GVC and store it in a transaction(thread)-local variable RV
- WRITE(m,v) operation:
 - Add <m,v> to the write set
- READ(m)(v) operation:
 - **IF** m in write set **THEN** return the associated v
 - **ELSE**
 - Load the versioned lock <version,locked> associated to m
 - IF locked or version > RV abort
 - Load v from m
 - IF locked or version > RV abort
 - Add <m> to the readset

Transactional Locking 2 [Dice2006]

- COMMIT:
 - For each m in the write set acquire the related versioned lock
 - If acquisition fails abort
 - Increment GVC via Add&Fetch obtaining WV
 - **IF WV != RV+1**
 - Validate the read set (abort if locked or version > RV)
 - Store each value in the read set
 - Release each versioned lock by using WV as version

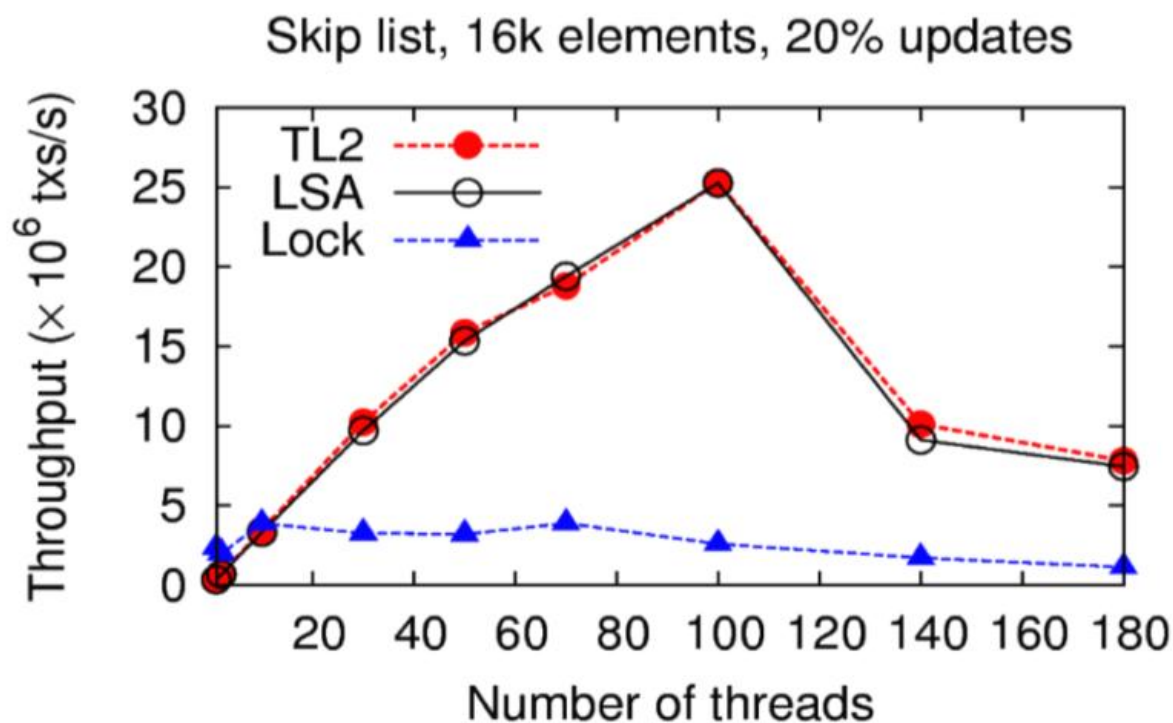
Transactional Locking 2 [Dice2006]

- REMARKS:
 - Re-validating the read set before applying updates is required due to possible concurrent updates during write-set locking and GVC increment
- Read-only transactions
 - Do not need to increase GVC
 - Do not need to acquire any lock
 - Do not need to revalidate the read set
 - Do not need the read set

Transactional Locking 2 [Dice2006]

- REMARKS:
 - Re-validating the read set before applying updates is required due to possible concurrent updates during write-set locking and GVC increment
- Read-only transactions
 - Do not need to increase GVC
 - Do not need to acquire any lock
 - Do not need to revalidate the read set
 - Do not need the read set

What about Software Transactional Memory



What about Software Transactional Memory

- Scale as (or better than) fine-grain locking
- Overheads hamper scalability
 - Due to instrumented access (overhead for each read/write)
 - Read set validation
- Hot topic in 2000s
 - A plethora of implementations for several programming languages
 - C/C++: TinySTM, G++ v4.7 (still experimental)
 - C#: SXM by Microsoft (discontinued)
 - Haskell: STM is part of the Haskell platform
 - Scala: Akka framework
- Large debate on its practical impact
 - Software Transactional Memory: Why Is It Only a Research Toy?: The promise of STM may likely be undermined by its overheads and workload applicabilities. [Cascaval2008]
 - Transactional Memory Should Be an Implementation Technique, Not a Programming Interface [Boehm2009]
 - Why STM can be more than a Research Toy [Dragojević2011]