

Programmazione concorrente

Laurea Magistrale in Ingegneria Informatica

Università Tor Vergata

Docente: Romolo Marotta

Safe Memory Reclamation

Safe Memory Reclamation

WHAT is SMR?

- Approach to release/reuse memory only when it is safe
- A bit different to Garbage Collectors

WHY SMR?

- In non-blocking algorithms, shared memory accesses are NOT protected by critical section via locks
- A thread can dereference a pointer whose target memory might be released (dangling pointers)
- Undefined behaviors, crash, memory corruptions

Terminologia SMR

- Allocation: creating the object
- Publish: making it visible to other threads
- Detach: removing it from shared structures
- Retiring: marking it for future reclamation
- Reclaim: actual deallocation

Life cycle of a shared object

Allocate

Locally by
a thread

Publish

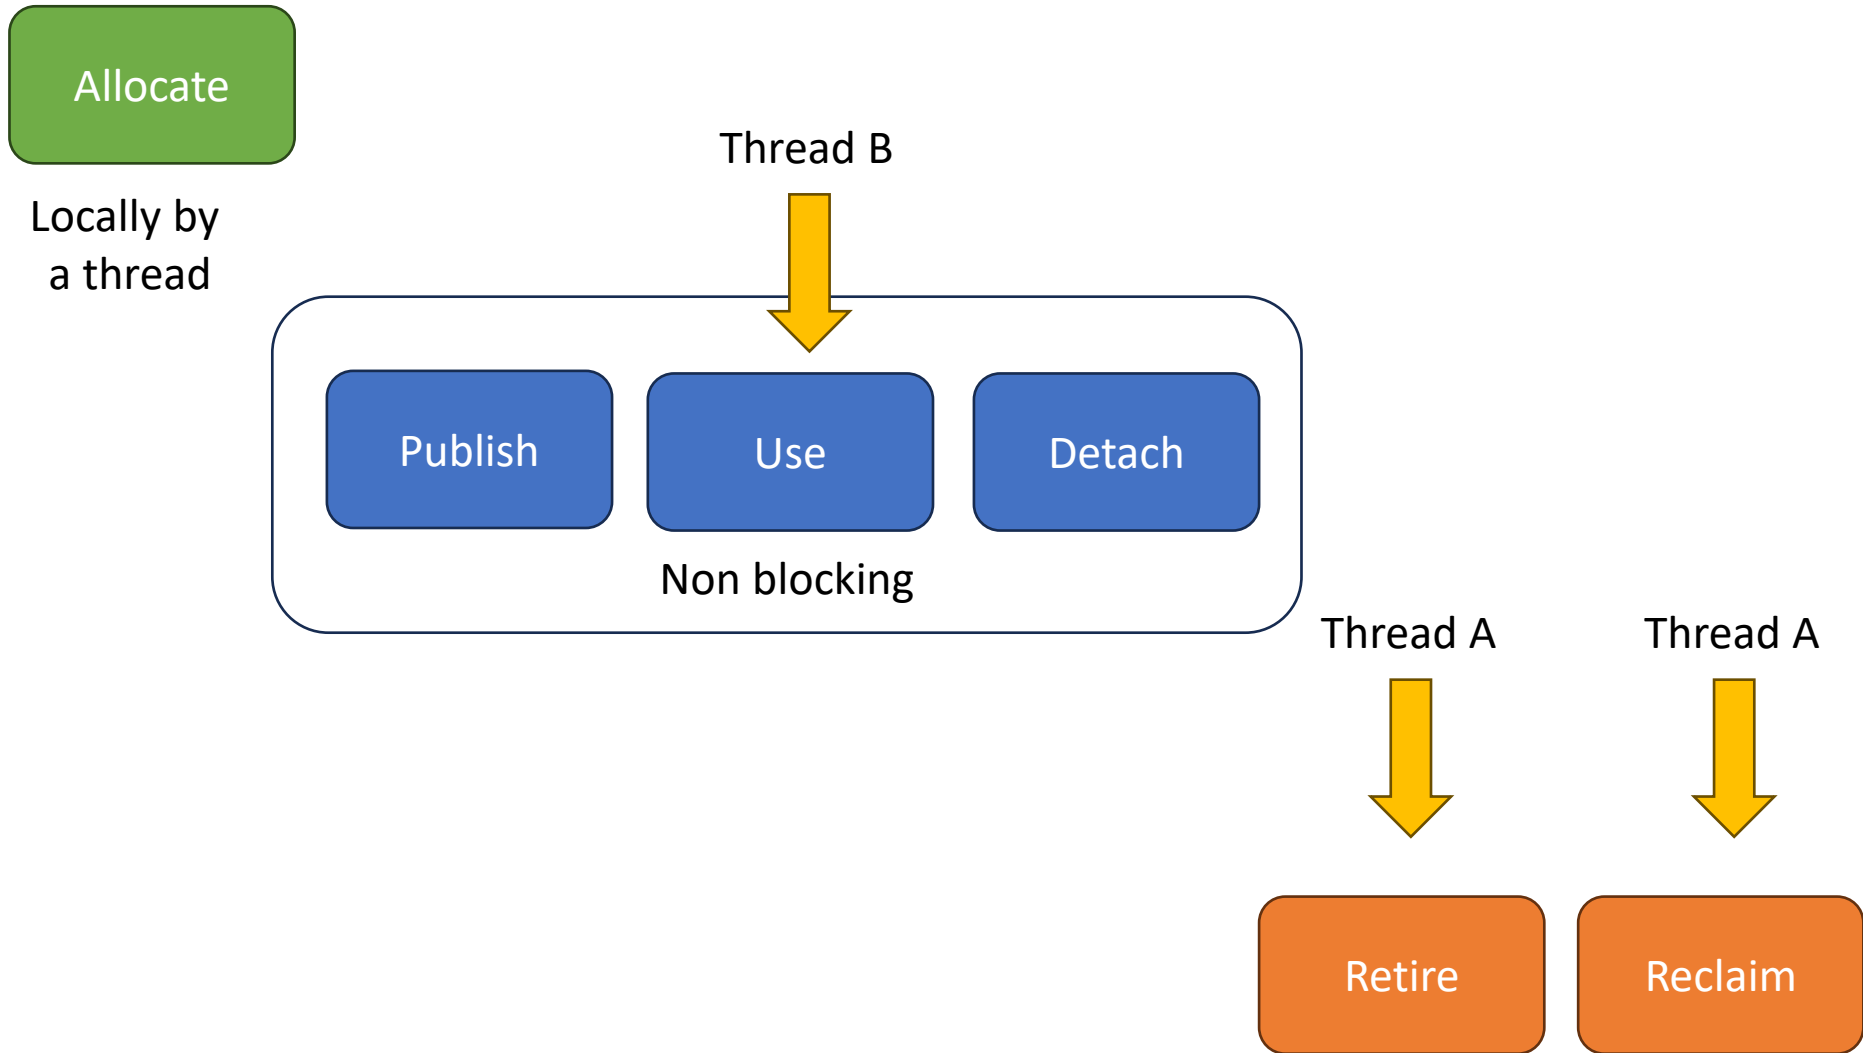
Use

Detach

Critical section

Reclaim

Life cycle of a shared object



SMR design

- Avoid stop-the-world approaches
- **Robustness**: a stalled thread should prevent only a **bounded** number of blocks from being reclaimed
- Threads needs to:
 - Maintain a list of memory buffers for later reclamation (retired list)
 - Identify grace periods for buffer reuse

Hazard Pointers

- Thread:
 - publish hazard pointers that they are going to use
 - unpublish pointers when unused by replacing it with NULL
- Memory referenced by a non-null HP cannot be reclaimed
- Let be:
 - N the number of threads
 - K the maximum number of HPs per thread
- We need:
 - $N \times K$ array of HPs publicly available to all threads
- The scheme can be generalized to support dynamic number of threads/HPs

Hazard Pointers: how to reclaim?

- After detaching a memory buffer, put references into a retired list
- Periodically scan the HP array
- While scanning HPs build a fast map for later searches (e.g. hashmap or binary search tree)
- For each retired buffer check the map
 - If present re-put in a new retired list
 - If not reclaim memory
- Swap the old retired list with the new one

Hazard Pointers: why works?

- Identifies the end of per-pointer grace periods as
 - The buffer is detached (and hence retired)
 - There is no HP pointing to it

Hazard Pointers: example

```
structure NodeType { Data:DataType; Next:*NodeType; }  
// Shared variables  
Head,Tail:*NodeType;  
// Initially both Head and Tail point to a dummy node  
  
Enqueue(data:DataType) {  
  1: node ← NewNode();  
  2: node^.Data ← data;  
  3: node^.Next ← null;  
    while true {  
  4:   t ← Tail;  
  5:   next ← t^.Next;  
  6:   if (Tail ≠ t) continue;  
  7:   if (next ≠ null) { CAS(&Tail,t,next); continue; }  
  8:   if CAS(&t^.Next,null,node) break;  
    }  
  9: CAS(&Tail,t,node);  
}
```

```
Dequeue() : DataType {  
  while true {  
  11:  h ← Head;  
  12:  t ← Tail;  
  13:  next ← h^.Next;  
  14:  if (Head ≠ h) continue;  
  15:  if (next = null) return EMPTY;  
  16:  if (h = t) { CAS(&Tail,t,next); continue;}  
  17:  data ← next^.Data;  
  18:  if CAS(&Head,h,next) break;  
    }  
  19: return data;  
}
```

Hazard Pointers: why works?

- Identifies the end of per-pointer grace periods as
 - The buffer is detached (and hence retired)
 - There is no HP pointing to it
- Threads double check that just published HP are still usable